

Improving Scheduling Algorithms via a Mathematical Perspective

Sami Davies

A dissertation
submitted in partial fulfillment of the
requirements for the degree of

Doctor of Philosophy

University of Washington

2021

Reading Committee:

Thomas Rothvoss, Chair

Anna Karlin

Isabella Novik

Program Authorized to Offer Degree:
Department of Mathematics

© Copyright 2021

Sami Davies

University of Washington

Abstract

Improving Scheduling Algorithms via a Mathematical Perspective

Sami Davies

Chair of the Supervisory Committee:

Thomas Rothvoss

Department of Mathematics

This document will discuss three problems that I worked on during my Ph.D. Chapter 2 contains my work on the Santa Claus problem, and Chapters 3 and 4 contain my work on scheduling with precedence constraints and communication delays.

In the *Santa Claus problem*, Santa has a set of gifts, and he wants to distribute them among a set of children so that the least happy child is as happy as possible. Child i has value p_{ij} for present j , where p_{ij} is in $\{0, p_j\}$. A modification of Haxell's hypergraph matching argument by Annamalai, Kalaitzis, and Svensson gives a 12.33-approximation algorithm for the problem. In joint work with Thomas Rothvoss and Yihao Zhang, we introduce a matroid version of the Santa Claus problem. While our algorithm is also based on the augmenting tree by Haxell, the introduction of the matroid structure allows us to solve a more general problem with cleaner methods. Using our result from the matroid version of the problem, we obtain a $(4 + \varepsilon)$ -approximation algorithm for Santa Claus.

In scheduling theory, one of the most poorly understood, yet practically interesting, models is scheduling in the presence of *communication delays*. Here, if two jobs are dependent and scheduled on different machines, then at least c units of time must pass between their executions. Even

for the special case where an unlimited number of identical machines are available, the best known approximation ratio for minimizing makespan is $O(c)$. An outstanding open problem in the top-10 list by Schuurman and Woeginger (and its recent update by Bansal) asks whether there exists a constant-factor approximation algorithm. In joint work with Janardhan Kulkarni, Thomas Rothvoss, Jakub Tarnawski, and Yihao Zhang, we prove a $O(\log c \cdot \log m)$ -approximation algorithm for the problem of minimizing makespan on m identical machines; this work is presented in Chapter 3. Our approach is based on a Sherali-Adams lift of a linear programming relaxation and a randomized clustering of the semimetric space induced by this lift. We extend our work to the related machines setting and study the objectives of minimizing makespan and minimizing the weighted sum of completion times. Here, we also obtain polylogarithmic approximation algorithms, and these results are presented in Chapter 4.

Contents

1	Introduction and Background	1
1.1	The Santa Claus problem	4
1.1.1	Previous work	5
1.1.2	Our contribution	7
1.2	Scheduling with Communication Delays	10
1.2.1	Previous work	11
1.2.2	Our contribution	13
2	Santa Claus	19
2.1	Introduction	19
2.1.1	Our contributions	22
2.2	An algorithm for Matroid Max-Min Allocation	23
2.2.1	Intuition for the algorithm	24
2.2.2	A detailed procedure	26
2.2.3	Correctness of the algorithm	28
2.2.4	Termination and runtime	34
2.3	Application to Santa Claus	36
2.4	Conclusion	40
3	Scheduling with Communication Delays on Identical Machines	41

3.1	Introduction	41
3.1.1	Our Contributions	44
3.1.2	Independent work of Maiti <i>et al</i>	45
3.1.3	Our Techniques	46
3.1.4	History of the Problem	48
3.2	Preliminaries	50
3.2.1	The Sherali-Adams Hierarchy for LPs with Assignment Constraints	50
3.2.2	Semimetric Spaces	53
3.3	Approximation for $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$	57
3.3.1	The Linear Program	58
3.3.2	Scheduling a Single Batch of Jobs	62
3.3.3	The Complete Algorithm for $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$	66
3.4	Reductions	67
3.5	Integrality Gap for the Compact LP	74
3.6	Minimizing Weighted Sum of Completion Times	76
3.6.1	The linear program	77
3.6.2	The Rounding Algorithm	79
4	Scheduling with Communication Delays on Related Machines	85
4.1	Introduction	85
4.1.1	Our Contributions	88
4.1.2	Technical Challenges	89
4.1.3	Our Techniques and Algorithm Overview	90
4.1.4	Outline	95
4.2	Preliminaries	95
4.2.1	The Sherali-Adams Hierarchy for LPs with Assignment Constraints	95
4.2.2	Semimetric Spaces	96

4.2.3	Concentration	96
4.3	The Linear Program and Its Properties	97
4.3.1	The Linear Program	97
4.3.2	Properties of the LP	99
4.4	The Rounding Algorithm for $Q \mid \text{prec}, p_j = 1, c \mid \sum w_j C_j$	104
4.4.1	First Batch Scheduling	105
4.4.2	Intermediate Batch Scheduling	112
4.5	Proof of Theorem 56	113
4.5.1	Main Scheduling Subroutine	113
4.5.2	The Complete Algorithm	116
4.6	Reductions	118
4.7	Makespan Minimization on Related Machines	123
4.8	Conclusion	126

Bibliography	127
---------------------	------------

Acknowledgments

It is hard for me to imagine a more capable and supportive advisor than Thomas. He taught me a lot about how to approach problems, what to do when I get stuck, and when to move on. In doing so, he always showed an incredible amount of patience, kindness, and humility in our work.

I received a great deal of mentorship from Janardhan Kulkarni during my Ph.D, and I thank Jana for his continued honesty, guidance, and support. Along with Jana and Thomas, several works in this thesis are joint with Yihao Zhang and Jakub Tarnawski, all of whom I was fortunate to work with and learn from. Jana and Jakub also hosted me at MSR during the summer of 2020, an experience for which I am grateful. While our work is not featured in this thesis, I am grateful to my coauthors Cyrus Rashtchian and Miki Racz for their friendship and mentorship over the past several years.

I'm indebted to our union, UAW 4121, and its organizers for the constant efforts in fighting for our rights and for serving as a gateway for me into local organizing spaces.

I made many friends in the Math department and CS theory group (and PLSE group), and I thank the students and faculty of both for allowing me to fluidly move between the two at my whim. People who really helped me to go to grad school and then study TCS were Bill Hrusa, Brendan Sullivan, Lev Reyzin, and Jeremy Kun. I thank my family—especially my sister, Addie—for supporting me as I live far away and do something they don't really understand. Lastly, to Max, thank you for helping me find balance in life as we did this thing.

DEDICATION

to Mr. Raymond Kreiser

Chapter 1

Introduction and Background

Many optimization problems can be modeled as linear programs of the form $\max_x \{c^T x \mid Ax \leq b\}$. There are methods to solve linear programs efficiently and obtain a fractional solution, but often any valuable x is integral. Finding optimal integral solutions is generally theoretically intractable, so instead we often take a fractional solution and devise some algorithm that *rounds* the fractional solution to an integral one. The rounding algorithm is seen as decent if not too much is lost in the objective value from moving from the fractional solution to the integral one.

In designing linear programming relaxations, we define that *integrality gap* of an LP to be the maximum ratio over all input instances I between the optimal objective value of the integral solution of I and the optimal objective value of the relaxation of I . The integrality gap is the standard way of measuring how well a linear program actually captures the problem.

We accept that it is extremely likely that there exists a separation between the set of “easy” and the set of “hard” problems, i.e. those for which we can always find an optimal solution in time polynomial in the size of the problem, and those for which we cannot. Skipping over the complexity formalities that can be found in many textbooks, the latter class is referred to as the set of **NP-hard** problems. Instead, we work towards obtaining approximate solutions efficiently.

An α -*approximation algorithm* for an optimization problem is an algorithm which given any input instance to the problem is guaranteed to return a solution whose value is within an α factor of the value of the optimal solution, and further the algorithm runs in time polynomial in the instance's size¹.

Designing LPs with small integrality gap—where small depends on the target approximation factor—can be difficult. One normally begins by writing down the basic constraints that arise from the problem statement; let us call this the basic LP for now. To judge the quality of the basic LP, the next step is to try and think of instances that might lead to a larger-than-desired integrality gap. If indeed there is an instance certifying that the integrality gap is large, then we have to come up with additional constraints to add the LP that will eliminate the bad instances. Sometimes this ad hoc approach works relatively well, in that coming up with constraints to eliminate bad instances is not too difficult given previous work, intuition, or luck².

But other times, this cycle of dreaming up new integrality gap instances and then nixing those instances with new constraints can go for a while. Even if there is an LP for which there are no obvious integrality gap instances, it might difficult to come up with a way to round an LP solution or analyze a rounding procedure.

Appealing though these options may be, hierarchies can help avoid them. LP (or SDP) hierarchies provide a systematic way to add new constraints and variables to an LP (or SDP) in order to strengthen them. Additionally, there are nice properties of solutions to lifted LPs and SDPs that make analyzing rounding procedures simpler. Sometimes after analyzing the problem using a lifted solution, it is possible to reverse-engineer a simpler LP—one that is not necessarily a lifted LP— by only using the constraints and variables enforced by the lift that are needed in the analysis

¹In some of the literature, the runtime is not specified in the definition, but in this thesis everything will be in polynomial time, so we build it into the definition.

²In the Santa Claus problem, this ad hoc approach works well enough for our purposes!

of the rounding procedure³. Other times, the analysis uses properties of the lifted solution that are global in the sense that it is not obvious how to extract a compact LP⁴.

For the purposes of this thesis, we will only give some background on the Sherali-Adams hierarchy, but the interested reader should see the chapter by Chlamtac and Tulsiani in the first reference or the work by Laurent for more details[AL12, Lau03]. From the basic LP relaxation, we strengthen as specified by the *rank* of the lift, where higher lifts are stronger. As one should expect, the increased strength with higher rank comes with higher complexity; an integer program on n variables, x_1, \dots, x_n , lifted to rank r can be solved in time $n^{O(r)}$. At a high level, a Sherali-Adams lift of rank r enforces a consistent probability distribution on sets of variables of size at most r . More specifically, observe that for K a polytope of a relaxation of a 0/1 integer program on n variables, we can view any point in $\text{conv}(K \cap \{0, 1\}^n)$ as a probability distribution X over points in the solution set $K \cap \{0, 1\}^n$. For $I \subseteq [n]$ with $|I| \leq r$, one could manually add variables y_I representing the probability over the distribution X that $X_i = 1$ for all $i \in I$ and constraints enforcing these probabilities agree according to inclusion-exclusion. We would ideally like to enforce that $y_I = \prod_{i \in I} y_i$; instead the Sherali-Adams lift requires the following:

Definition 1. Let $\mathcal{P}_r([n]) = \{I \subseteq [n] \mid |I| \leq r\}$ and $K = \{x \in \mathbb{R}^n \mid Ax \geq b\}$ for $A \in \mathbb{R}^{m \times n}$. Then $SA_r(K)$, the Sherali-Adams lift of K of rank r , is the set of vectors $y \in [0, 1]^n$ with $y_\emptyset = 1$ and for all $I, J \subseteq [n]$ with $|I| + |J| \leq r$,

$$\sum_{J' \subseteq J} (-1)^{|J'|} \cdot \left(\sum_{i=1}^n A_{\ell, i} y_{I \cup J' \cup \{i\}} - b_\ell y_{I \cup J'} \right) \geq 0 \quad \forall \ell \in [m].$$

More properties of the Sherali-Adams hierarchy are discussed in Chapters 3 and 4.

Problems for which we use approximation algorithms and that are considered in theoretical

³This is what happened in our work on scheduling with communication delays on identical machines. One can view the Sparsest Cut algorithm by Arora, Rao, and Vazirani this way, too.

⁴This is what happened in our work in scheduling with communication delays on related machines.

computer science are often motivated from practical situations. But it is worth noting that in the definition of an α -approximation algorithm, we require that the algorithm obtains a solution whose value is only α off for *any* input. Depending on your persuasion, this may seem like a very strong definition, and maybe impractical... perhaps the inputs that make algorithm design difficult are contrived or never really arise in practice. So we have an—at first glance—strange juxtaposition, where we motivate a problem by saying that it arises in practice, but we measure success with a potentially impractical metric. From my perspective, this view is orthogonal to the work I did during my Ph.D., as well as a lot of the most important work in TCS⁵. The design and analysis of approximation algorithms represented in this thesis is all about the math. Specifically, it is about discovering what mathematical structures we can superimpose onto a hard problem in order to make a little more sense of why it is indeed hard. We focus on two specific **NP**-hard problems in this thesis: Santa Claus, also known as Restricted Max-Min Fair Allocation, and scheduling with precedence constraints and communication delays.

1.1 The Santa Claus problem

In the *Santa Claus problem*, the objective is to allocate gifts, M , to children, J , in such a way that the minimum value of any child is maximized. We assume that child $i \in M$ has value $p_{i,j}$ for gift $j \in J$ and that a child's value is the sum of the values of the gifts they receive. Formally, we seek an assignment $\phi : J \rightarrow M$ such that $\min_{i \in M} \sum_{\phi^{-1}(i)} p_{i,j}$ is maximized. Let m denote the number of children and n the number of gifts. In Santa Claus, also known as restricted Max-Min Fair Allocation, we assume that $p_{i,j} \in \{0, p_j\}$, i.e if two children both want to receive a gift, then they both value it equally. We will refer to the unrestricted version of MMFA occasionally for

⁵Though there is a lot of meaningful work in approximation algorithms, TCS, and beyond that takes runtime, memory, input characteristics, and other practical concerns into consideration, and I very well might work in this space someday.

background, where in the unrestricted version $p_{i,j}$ does not necessarily belong to $\{0, p_j\}$ ⁶.

1.1.1 Previous work

It is **NP**-hard to approximate Santa Claus within a factor less than 2 [BD05]. Though interestingly, the known approximation factors for upper bounds wildly differ between the restricted and unrestricted versions. Before our work, constructive methods for both the restricted and unrestricted versions of Max-Min Fair Allocation in some way use the existence of a solution to an exponentially large *configuration linear program* [BS06, AKS15, AS10, CCK09, CM18].

For a CLP with objective value T , the variables in the CLP take the form $x_{i,C}$, which represents whether child i receives the subset of gifts $C \in \mathcal{C}(i, T)$, those sets of gifts bringing child i at least value T :

$$\begin{aligned}
 (CLP) \quad & \sum_{C \in \mathcal{C}(i, T)} x_{i,C} \geq 1 \quad \forall i \in M \\
 & \sum_{C: j \in C} \sum_i x_{i,C} \leq 1 \quad \forall j \in J \\
 & x \geq 0.
 \end{aligned}$$

The CLP was introduced by Bansal and Sviridenko when studying the Santa Claus problem. They use it to give an $O(\log \log m / \log \log \log m)$ -approximation for the Santa Claus problem [BS06]. Since then, progress has been made towards better constant factor approximations for the Santa Claus problem, with Feige proving an (unspecified constant) $O(1)$ -approximation and Annamalai, Kalaitzis, and Svensson proving a 12.3-approximation [Fei08, AKS15]. The CLP for the Santa Claus problem has an integrality gap of 3.808 [CM19]. In contrast, the CLP for

⁶In Max-Min Fair Allocation, we normally use the terms “agents” and “items” instead of “children” and “gifts”, but since we only briefly discuss the unrestricted version, we will stick with the latter lingo.

unrestricted Max-Min Fair Allocation has an integrality gap of $\Omega(\sqrt{m})$, but Chakrabarty, Chuzhoy, and Khanna cleverly beat the \sqrt{m} gap to obtain an $\tilde{O}(n^\varepsilon)$ -approximation in $n^{1/\varepsilon}$ time for $\varepsilon = \Omega(\log \log n / \log n)$ by iteratively using solutions to the LP, whose gaps become smaller in each iteration⁷. From these works, it is helpful to note that the algorithms rely on assigning gifts with large versus small $p_{i,j}$ differently [AKS15, CCK09]. Similar ideas have been used in scheduling to minimize makespan [Sve11, LST87].

More generally, studying Max-Min Fair Allocation can be viewed from the lens of investigating sufficient conditions for perfect matchings in *bipartite hypergraphs*. A hypergraph $\mathcal{H} = (V \dot{\cup} W, \mathcal{E})$ is called bipartite if for all $e \in \mathcal{E}$, $|e \cap V| = 1$, and hyperedges $F \subseteq \mathcal{E}$ form a (left) perfect matching if they are disjoint and every node in V is contained in exactly one edge in F . In particular, a solution to the CLP is a fractional perfect matching on the hypergraph $(M \dot{\cup} J, \bigcup_{i \in M} \mathcal{C}(i, T))$. Finding perfect matchings in bipartite hypergraphs is in general **NP**-hard, though *Haxell's condition* defines a sufficient condition:

Theorem 1. [Hax95] *Let $\mathcal{H} = (V \dot{\cup} W, \mathcal{E})$ be a bipartite hypergraph with $|e| \leq r$ for all $e \in \mathcal{E}$. Then either \mathcal{H} contains a (left) perfect matching or there are subsets $V' \subseteq V$ and $W' \subseteq W$ so that all hyperedges incident to V' intersect W' and $|W'| \leq (2r - 3)(|V'| - 1)$.*

While the proof of Haxell's condition relies on a potentially exponential time argument, Annamalai made the argument polynomial by restricting edge sizes to be uniform and allowing some slack in the inequality [Ann16]. His argument, in addition to the one made in the 12.33-approximation algorithm by Annamalai, Kalaitzis, and Svensson, takes Haxell's argument and makes it polynomial by introducing an augmenting tree. This generalizes the notion of an augmenting path used to find maximum matchings in bipartite graphs. Instead of swapping every

⁷Note that some results are in terms of m , and some in terms of n . Though in general we think of n as much larger than m , and in fact, the results of Chakrabarty, Chuzhoy, and Khanna hold replacing n with m for constant ε .

other edge in an augmenting path, as is the case for a bipartite graph, the augmenting tree swaps sets of alternating hyperedges to find more space in the hypergraph.

1.1.2 Our contribution

In my work on the Santa Claus problem with Thomas Rothvoss and Yihao Zhang, we exploit an underlying matroid to design a clean bipartite hypergraph matching algorithm [DRZ20]. This leads to an improved approximation factor (from 12.3 to $4 + \varepsilon$) compared to the solution of a new linear program relaxation⁸.

We build off the idea of separating gifts by their values by designating a gift j as either large or small, depending on its value p_j . Our linear program has only $O(n^2)$ many variables and constraints and uses this distinction between large and small gifts. In particular, we introduce an expansion constraint, forcing children who do not entirely receive a large gift to fractionally receive at least T small gifts, for T the optimal objective value. Let $J_L := \{j \in J \mid p_j > \delta_2 T\}$ denote the large gifts and $J_S := \{j \in J \mid p_j \leq \delta_1 T\}$ denote the small gifts for constants $0 < \delta_1 \leq \delta_2 < 1$ such that all gifts have values in $[0, \delta_1 T] \cup (\delta_2 T, T]$. M_j denotes the set of children who wish to receive gift j , in other words the set of $i \in M$ such that $p_{i,j} > 0$. Let $P(T, \delta_1, \delta_2)$ be the set of vectors $z \in \mathbb{R}_{\geq 0}^{J \times M}$ satisfying

$$\begin{aligned} \sum_{j \in J_S: i \in M_j} p_j z_{ij} &\geq T \cdot \left(1 - \sum_{j \in J_L: i \in M_j} z_{ij}\right) && \forall i \in M \\ \sum_{i \in M_j} z_{ij} &\leq 1 && \forall j \in J \\ z_{ij} &\leq 1 - \sum_{j' \in J_L: i \in M_{j'}} z_{ij'} && \forall j \in J_S \forall i \in M_j. \end{aligned}$$

⁸After obtaining our results, we learned that Cheng and Mao simultaneously and independently obtained a $6 + \varepsilon$ approximation to the Santa Claus problem by altering the algorithm by Annamalai, Kalaitzis and Svensson [CM18]. Using ideas established by Asadpour, Feige, and Saberi, they further improved this to a $4 + \varepsilon$ approximation [CM19, AFS08].

The separation between large and small gifts was critical for our approach as it allowed us to define an underlying matroid. Recall a matroid is a structure generalizing the notion of linear independence in vector spaces, and in particular their independent sets have useful exchange properties. Formally, $\mathcal{M} = (X, \mathcal{I})$ is a matroid with groundset X and independent sets $\mathcal{I} \subseteq 2^X$ if and only if the following hold:

- *Non-emptiness*: $\emptyset \in \mathcal{I}$;
- *Monotonicity*: For $Y \in \mathcal{I}$ and $Z \subseteq Y$ one has $Z \in \mathcal{I}$;
- *Exchange property*: For all $Y, Z \in \mathcal{I}$ with $|Y| < |Z|$ there is an element $z \in Z \setminus Y$ so that $Y \cup \{z\} \in \mathcal{I}$.

The bases of a matroid, denoted by $\mathcal{B}(\mathcal{M})$, are all inclusion-wise maximal independent sets, and the convex hull of all bases is called the base polytope, denoted by $P_{\mathcal{B}(\mathcal{M})} := \text{conv}\{\chi(S) \in \{0, 1\}^X \mid S \text{ is basis}\}$, where $\chi(S)$ is the characteristic vector of S .

Now consider a bipartite graph $G = (V \dot{\cup} W, E)$ with ground set V on one side and a set of resources W on the other side; each resource $w \in W$ has value $p_w \geq 0$. We define the problem *Matroid Max-Min Allocation*, where the goal is to find a basis $S \in \mathcal{B}(\mathcal{M})$ and an assignment $\sigma : W \rightarrow S$ with $(\sigma(w), w) \in E$ so that $\min_{v \in S} \sum_{w \in \sigma^{-1}(v)} p_w$ is maximized. For $T \geq 0$ the target objective function value, we define a linear programming relaxation $Q(T)$ as the set of vectors $(x, y) \in \mathbb{R}_{\geq 0}^X \times \mathbb{R}_{\geq 0}^E$ satisfying the constraints:

$$\begin{aligned}
 x &\in P_{\mathcal{B}(\mathcal{M})} \\
 \sum_{w \in N(v)} p_w y_{vw} &\geq T \cdot x_v && \forall v \in V \\
 y(\delta(w)) &\leq 1 && \forall w \in W \\
 y_{vw} &\leq x_v && \forall (v, w) \in E.
 \end{aligned}$$

Here, the decision variable x_v expresses whether element v should be part of the basis, and y_{vw} expresses whether resource w should be assigned to element v . We abbreviate $N(v)$ for the neighborhood of v and $y(\delta(w))$ for $\sum_{v:(v,w) \in E} y_{vw}$. Our main technical result is the following:

Theorem 2. [DRZ20] *Suppose $Q(T) \neq \emptyset$ and membership in the matroid can be tested in time polynomial in n . Then for any $\varepsilon > 0$ one can find*

$$(x, y) \in Q \left(\left(\frac{1}{3} - \varepsilon \right) \cdot T - \frac{1}{3} \cdot \max_{w \in W} p_w \right)$$

with both x and y integral in time $n^{\Theta_\varepsilon(1)}$, where $n := |V| + |W|$.

Viewing the set of children and gifts as parts of a bipartite graph, the sets of children that can be given a large gift form a matchable set matroid, \mathcal{M} . Thus children who do not receive a large gift form bases of \mathcal{M}^* , the co-matroid of \mathcal{M} . Set δ_1 so that $\delta_1 \cdot T$ is the largest gift value that is at most $\frac{1}{4} \cdot T$, and set δ_2 so that $\delta_2 \cdot T$ is the smallest gift value that is at least $\frac{1}{4} \cdot T$. In an instance of the Santa Claus problem where the optimal integral solution has value T , we transform a solution in $P(T, \delta_1, \delta_2)$ into a solution of $Q(T)$ with $\max_{w \in W} p_w = \delta_1 \cdot T$.

By using properties of matroids, we find an assignment of small gifts to a basis of \mathcal{M}^* by simplifying a bipartite hypergraph matching algorithm of Annamalai, Kalaitzis, and Svensson that builds an augmenting tree [AKS15]. Using Theorem 2, each child receiving only small gifts has value at least $(\frac{1}{3} - \varepsilon) \cdot T - \frac{1}{3} \cdot \delta_1 \cdot T \geq (\frac{1}{4} - \varepsilon) \cdot T$. The remaining children who received no small gifts get a single large gift with value at least $\delta_2 \cdot T \geq \frac{1}{4} \cdot T$, guaranteeing every child at least $\frac{1}{4} - \varepsilon$ of the optimal value. Interestingly, in the case where the distribution of gift values is bimodal and the gap between δ_1 and δ_2 is large, our approximation factor improves. In the most extreme case when gift values are either 1 or T , our algorithm returns a $3 + \varepsilon$ approximation.

1.2 Scheduling with Communication Delays

Distributing tasks onto processors occurs in many areas of computing. Frequently, some tasks use as their inputs the outputs of other tasks. The cost of transferring data between processors can be expensive, e.g. in data center scheduling and Deep Neural Network training [CS12, ZCB⁺15]. This dependence and latency in scheduling is modelled with precedence constraints and communication delays.

Given a set of n jobs, and a set of m machines on which the jobs can be performed, scheduling algorithms find an assignment of jobs to machines and time slots. The inclusion of *precedence constraints* model the flow of jobs as a directed acyclic graph, where if there is an edge in the DAG from j_1 to j_2 , denoted $j_1 \prec j_2$, then j_1 must be completed before any machine can start to process j_2 . The inclusion of *communication delays* model the latency between dependent jobs, where if $j_1 \prec j_2$ and j_2 is performed on a different machine than j_1 , then at least c_{j_1, j_2} time units must pass after the completion of j_1 before any machine can start to process j_2 .

One objective of interest is minimizing *makespan*, which is the time that the last machine processes its final job. In the setting where each jobs is additionally assigned a weight, w_j for job $j \in J$, and C_j is the time at which job j is completed, another objective is to minimize the *weighted sum of completion times*, $\sum_{j \in J} w_j C_j$. In the *identical* machines setting, all machines process all jobs equally quickly. In the *related* machines setting, each machine is equipped with some speed s_i for $i \in [m]$ and job j is processed in p_j/s_i time units by machine i . We are interested in *non-preemptive* scheduling, where a job must be processed in its entirety on a single machine in a consecutive period of time.

1.2.1 Previous work

Overall, there was very little known about scheduling with communication delays for non-constant c before our work and the concurrent work of Maiti, Rajaraman, Stalfa, Svitkina, and Vijayaraghavan[MRS⁺20], despite much practical and theoretical intrigue. In fact, obtaining constant factor approximation algorithms to the problem is in a list of the top 10 open problems in scheduling theory that was proposed by Schuurman and Woeginger [SW99]. More recently in a 2017 survey talk, Bansal described scheduling with precedence constraints and communication delays as “not understood at all” and “almost completely open” [Ban17].

Simple arguments obtained $O(c)$ -approximations for the objective of minimizing makespan [Gra66, GKMP08]. In a *list scheduling* algorithm, a heuristic orders the jobs and then jobs are assigned according to the order to the least full machine. List scheduling already gives a $(c+2)$ -approximation for the problem, and the best known for general c was a $2/3(c+1)$ approximation. While there are tighter bounds when $c = 1$, little else was known for when c is large, which is the parameter setting arising in practice [MK97, HM01].

Without communication delays, but with precedence constraints, Graham showed that with any ordering, list-scheduling gives a $(2 - 1/m)$ -approximation. And this is roughly tight, as assuming a variant of the Unique Games Conjecture, in 2010 Svensson showed that it is **NP**-hard to obtain better than a $(2 - \epsilon)$ -approximation for constant $\epsilon > 0$. It should also be noted that one can instead consider the problem with *duplication*; here a job can be performed on multiple machines in order to forego the communication cost incurred for placing dependent jobs on different machines.

One hypothesis of why this problem eluded researchers for so many decades is the lack of good LP relaxations. To see why, let us consider the LP of Munier and König, which—together with a simple rounding procedure—was used to obtain a $4/3$ -approximation for scheduling with

communication delays in the special case when $p_j = 1$ for all $j \in J$ and $c = 1$. In the LP below, we call $\Gamma^-(j) = \{j' \in J \mid j' \prec j\}$ the *predecessors* of j and $\Gamma^+(j) = \{j' \in J \mid j \prec j'\}$ the *successors* of j . Further, C_j is the completion time of job j and $x_{j,j'}$ is the decision variable for whether jobs $j \prec j'$ are performed on the same machine.

$$\begin{aligned}
C_{j'} + 2 - x_{j,j'} &\leq C_j && \forall j' \prec j \\
\sum_{j' \in \Gamma^-(j)} x_{j',j} &\leq 1 && \forall j \in J \\
\sum_{j' \in \Gamma^+(j)} x_{j,j'} &\leq 1 && \forall j \in J \\
0 &\leq C_j && \forall j \in J \\
0 &\leq x_{j,j'} \leq 1 && \forall j \prec j'
\end{aligned}$$

Now, we could try to extend the LP above to the case when c is large. Assume the time horizon is partitioned into intervals of length c and the completion time variables now refer to which c -length interval to schedule in, e.g. the first c length interval is from time 0 to c and $C_j = 0$ for job j if and only if job j is performed by the LP in the first interval. Also, we slightly change the variables $y_{j,j'}$ to be the decision variable for whether jobs j and j' are performed on the same machine and in the same c -length interval. Then we end up with the following LP:

$$\begin{aligned}
C_{j'} + 1 - y_{j,j'} &\leq C_j && \forall j' \prec j \\
\sum_{j' \in J} y_{j,j'} &\leq c && \forall j \in J \\
0 &\leq C_j && \forall j \in J \\
0 &\leq y_{j,j'} \leq 1 && \forall j \prec j'
\end{aligned}$$

Unfortunately, this LP has a non-constant integrality gap; for a DAG with degree at most c , the LP can simply set schedule all jobs in the first interval, even though an integral solution will need a superconstant number of intervals.

See Chapters 3 and 4 for a comparison on our work versus that of Maiti et. al.

1.2.2 Our contribution

In joint work with Janardhan Kulkarni, Thomas Rothvoss, Jakub Tarnawski, and Yihao Zhang, we studied the problem of scheduling with precedence constraints and communication delay c on identical machines. In the identical machine setting, all machines have the same speed and are able to process every job.

As outlined above, it was not clear how to construct a useful LP using an ad hoc approach. Therefore, we used a *Sherali-Adams* lift. The LP is constructed in two steps; the first step constructs a basic assignment LP and the second step uses the variables produced from lifting the basic LP. Consider the assignment variables

$$x_{j,i,s} = \begin{cases} 1 & \text{if } j \text{ is scheduled on machine } i \text{ in interval } I_s \\ 0 & \text{otherwise} \end{cases} \quad \forall j \in J, i \in [m], s \in \{0, \dots, S-1\},$$

and let K be the set of fractional solutions to

$$\begin{aligned} \sum_{i \in [m]} \sum_{s \geq 0} x_{j,i,s} &= 1 \quad \forall j \in J \\ \sum_{j \in J} x_{j,i,s} &\leq c \quad \forall i \in [m] \quad \forall s \in \{0, \dots, S-1\} \\ 0 \leq x_{j,i,s} &\leq 1 \quad \forall j \in J, i \in [m], s \in \{0, \dots, S-1\}. \end{aligned}$$

A lift $x \in \text{SA}_r(K)$ introduces variables $x_{(j_1, i_1, s_1), (j_2, i_2, s_2)}$ that are the probability that job j_1 is scheduled in interval s_1 on machine i_1 and job j_2 is scheduled in interval s_2 on machine i_2 . Variables C_j and $y_{j, j'}$ are those of our previously attempted LP, and S is the number of c -length intervals on the time horizon. We let $Q(r)$ denote the set of vectors (x, y, C) satisfying the following:

$$\begin{aligned} y_{j_1, j_2} &= \sum_{s \in \{0, \dots, S-1\}} \sum_{i \in [m]} x_{(j_1, i, s), (j_2, i, s)} \\ C_{j_2} &\geq C_{j_1} + (1 - y_{j_1, j_2}) \quad \forall j_1 \prec j_2 \\ C_j &\geq 0 \quad \forall j \in J \\ x &\in \text{SA}_r(K) \end{aligned}$$

Note that for our purposes, it suffices to take $r = 5$, so we can solve the LP in polynomial time. Our main technical theorem is the following:

Theorem 3. *Consider an instance with unit-length jobs J , a partial order \prec , and parameters $c, S, m \in \mathbb{N}$ such that $Q(r)$ is feasible for $r := 5$. Then there is a randomized algorithm with expected polynomial running time that finds a schedule for the instance using at most $O(\log m \cdot \log c) \cdot S$ intervals.*

The Sherali-Adams lift guarantees that the triangle inequality holds on the function $d(j, j') = 1 - y_{j, j'}$. Therefore the set of jobs equipped with d is a *semimetric space*, where jobs close with respect to the semimetric are likely to be scheduled on the same machine and in the same c -length interval. Equipped with a semimetric, we design an algorithm that clusters jobs with a clustering algorithm by Calinescu, Karloff, and Rabani [CKR04].

Since we are willing to pay a $O(\log c)$ factor in the approximation, it suffices to partition the set of jobs into band of width $\Theta(1/\log c)$ and just schedule each band, paying c between scheduling

each band. For a fixed band $J^* = \{j \mid T \leq C_j \leq T + \frac{1}{\Theta(\log n)}\}$, jobs $j, j' \in J^*$ are very likely to be scheduled on the same machine and in the same interval according to the LP. Using this fact, we can show that the probability that a node $j \in J^*$ is assigned by the CKR algorithm to a cluster different from any of its predecessors in J^* is at most $1/2$. Therefore the downward closed subsets of each cluster can be scheduled on an unused machine in an interval of size at most $2c$. Repeating this iteratively, and inserting c idle time slots between each iteration to account for the communication delay between jobs in different iterations, we show that with high probability this procedure schedules all the jobs⁹ after $\Theta(\log m)$ iterations.

We can reduce from the general setting to this special unit-length jobs setting with the next theorem.

Theorem 4. *Given an instance with unit-length jobs, suppose that there is a polynomial time algorithm that takes a solution for the LP $Q(r)$ with parameters $m, c, S \in \mathbb{N}$ and $r \geq 5$ and transforms it into a schedule using at most $\alpha \cdot S$ intervals. Then there is a polynomial time $O(\alpha)$ -approximation for scheduling with communication delays on identical machines for arbitrary job lengths.*

Overall, this implies a $O(\log m \cdot \log c)$ -approximation algorithm for scheduling with communication delays on identical machines.

Sherali-Adams gives the insight that the following compact LP actually suffices for our analysis. Our compact LP is the following:

$$\begin{aligned} y_{j_1, j_1} &= 1 & \forall j_1 \in J \\ y_{j_1, j_2} &= y_{j_2, j_1} & \forall j_1, j_2 \in J \\ 1 - y_{j_1, j_2} &\leq 1 - y_{j_1, j_3} + 1 - y_{j_2, j_3} & \forall j_1, j_2, j_3 \in J \end{aligned}$$

⁹Potentially a set of jobs is left after $O(\log m)$ iterations, but at that point, they can be scheduled together with high probability.

$$\begin{aligned}
\sum_{j \in J} y_{j_1, j} &\leq c & \forall j_1 \in J \\
C_{j_2} &\geq C_{j_1} + (1 - y_{j_1, j_2}) & \forall j_1 \prec j_2 \\
C_j &\geq 0 & \forall j \in J \\
0 \leq y_{j, j'} &\leq 1 & \forall j, j' \in J
\end{aligned}$$

With the same team of authors, we also obtained a polylogarithmic approximation algorithms for the more general setting of related machines with the objective of minimizing the weighted sum of completion times [DKR⁺21]. This work generalizes the Sherali-Adams framework established in their previous paper, though there are new technical challenges. For the objective of minimizing makespan, it was sufficient to determine in which c -length time interval a job is scheduled. However, for the weighted sum of completion times objective, this is not careful enough; jobs that have very small fractional completion time by the LP must be scheduled very early. Additionally, as opposed to the identical machines setting, the problem now involves assigning jobs to speed classes. We introduce more constraints into our LP and prove stronger structural insights about the Sherali-Adams solution. Our main result is an $O(\log^4 n)$ -approximation algorithm for the problem. As a byproduct of our result, we also obtain an $O(\log^3 n)$ -approximation algorithm for the problem of minimizing makespan $Q \mid \text{prec}, c \mid C_{\max}$, which improves upon the $O(\log^5 n / \log \log n)$ -approximation algorithm due to a recent work of Maiti et al. [MRS⁺20].

We consider the problem of scheduling jobs with communication delays on related machines to minimize the weighted sum of completion times. Let M denote the number of speed classes, i.e. the number of distinct s_i . While only losing a constant factor in the approximation factor, we can assume that $M = \log(s_{\max}/s_{\min}) \leq O(\log m)$, where s_{\max} and s_{\min} are the fastest and slowest speeds of machines. Our main result in the related machines setting is the following theorem.

Theorem 5. *There is a randomized $O(M^2 \cdot \log^2 n)$ -approximation algorithm for scheduling with communication delays on related machines for minimizing the weighed sum of completion times with expected polynomial running time. When jobs have unit processing lengths, the approximation factor of the algorithm improves to $O(M \cdot \log^2 n)$.*

Due to the bound on M , the theorem implies an $O(\log^4 n)$ -approximation algorithm to the general problem and an $O(\log^3 n)$ -approximation algorithm when all jobs are unit-length. We also obtain an improved approximation factor for the objective of minimizing makespan.

Theorem 6. *There is a randomized $O(M \cdot \log m \cdot \log n)$ -approximation algorithm for scheduling with communication delays on related machines for minimizing makespan with expected polynomial running time. When jobs have unit processing lengths the approximation factor of the algorithm improves to $O(\log m \cdot \log n)$.*

Chapter 2

Santa Claus

2.1 Introduction

Formally, the *Santa Claus* problem takes as input a set M of children, a set J of gifts, and values $p_{ij} \in \{0, p_j\}$ for all $i \in M$ and $j \in J$. In other words, a child is only interested in a particular subset of the gifts, but then its value only depends on the gift itself. The goal is to find an assignment $\sigma : J \rightarrow M$ of gifts to children so that $\min_{i \in M} \sum_{j \in \sigma^{-1}(i)} p_{ij}$ is maximized.

The first major progress on this problem is due to Bansal and Sviridenko [BS06], who showed a $O(\log \log n / \log \log \log n)$ -approximation based on rounding a *configuration LP*. The authors of [BS06] also realized that in order to obtain a $O(1)$ -approximation, it suffices to answer a purely combinatorial problem: show that in a uniform bipartite hypergraph with equal degrees on all sides, there is a left-perfect matching that selects a constant fraction of nodes from original edges. This question was affirmatively answered by Feige [Fei08] who proved a large unspecified constant using the Lovász Local Lemma repeatedly. Then Asadpour, Feige and Saberi [AFS08] showed that one can answer the question of [BS06] by using a beautiful theorem on hypergraph

matchings due to Haxell [Hax95]; their bound¹ of 4 has been slightly improved to 3.84 by Jansen and Rohwedder [JR18c] and then to 3.808 by Cheng and Mao [CM19]. Recently, Jansen and Rohwedder [JR18a] also showed (still non-constructively) that it suffices to compare to a linear program with as few as $O(n^3)$ many variables and constraints, in contrast to the exponential size configuration LP.

A hypergraph $\mathcal{H} = (X \dot{\cup} W, \mathcal{E})$ is called *bipartite* if $|e \cap X| = 1$ for all hyperedges $e \in \mathcal{E}$. A (left-) *perfect matching* is a set of hyperedges $F \subseteq \mathcal{E}$ that are disjoint but cover each node in X . In general, finding perfect matchings in even bipartite hypergraphs is **NP**-hard, but there is an intriguing sufficient condition:

Theorem 7 (Haxell [Hax95]). *Let $\mathcal{H} = (X \dot{\cup} W, \mathcal{E})$ be a bipartite hypergraph with $|e| \leq r$ for all $e \in \mathcal{E}$. Then either \mathcal{H} contains a left-perfect matching or there is a subset $C \subseteq X$ and a subset $U \subseteq W$ so that all hyperedges incident to C intersect U and $|U| \leq (2r - 3) \cdot (|C| - 1)$.*

It is instructive to consider a “standard” bipartite graph with $r = 2$. In this case, if there is no perfect matching, then there is a set $C \subseteq X$ with at most $|C| - 1$ many neighbors — so Haxell’s condition generalizes *Hall’s Theorem*. Unlike Hall’s Theorem, Haxell’s proof is non-constructive and based on a possibly exponential time augmentation argument. Only very recently and with a lot of care, Annamalai [Ann16] managed to make the argument polynomial. This was accomplished by introducing some slack into the condition and assuming the parameter r is a constant. Preceding [Ann16], Annamalai, Kalaitzis and Svensson [AKS15] gave a non-trivially modified version of Haxell’s argument for Santa Claus, which runs in polynomial time and gives a 12.33-approximation². Recently, Cheng and Mao altered their algorithm to improve the approximation to $6 + \varepsilon$, for any constant $\varepsilon > 0$ [CM18]. Our algorithm will also borrow a lot

¹The conference version [AFS08] provides a factor of 5, and it was improved to 4 in the journal version [AFS12].

²To be precise they obtain a $(6 + 2\sqrt{10} + \varepsilon)$ -approximation in time $n^{O(\frac{1}{\varepsilon^2} \log(\frac{1}{\varepsilon}))}$.

from [AKS15]. However, through a much cleaner argument we obtain a result that works in a more general matroid setting and implies a better approximation of $4 + \varepsilon$ for Santa Claus.

It should not go without mention that the version of the Santa Claus problem with arbitrary p_{ij} has also been studied before under the name *Max-Min Fair Allocation*. Interestingly, the integrality gap of the configuration LP is at least $\Omega(\sqrt{n})$ [BS06]. Still, Chakrabarty, Chuzhoy and Khanna [CCK09] found a (rather complicated) $O(\log^{10}(n))$ -approximation algorithm in $n^{O(\log n)}$ time³.

Santa Claus has a very well studied “dual” minmax problem. Usually it is phrased as *Makespan Scheduling* with *machines* $i \in M$ and *jobs* $j \in J$. Then we have a running time p_{ij} of job j on machine i , and the goal is to assign jobs to machines so that the maximum load of any machine is minimized. In this general setting, the seminal algorithm of Lenstra, Shmoys and Tardos [LST87] gives a 2-approximation — with no further improvement since then. In fact, a $(\frac{3}{2} - \varepsilon)$ -approximation is NP-hard [LST87], and the configuration LP has an integrality gap of 2 [VW11]. In the restricted assignment setting with $p_{ij} \in \{p_j, \infty\}$, the breakthrough of Svensson [Sve11] provides a non-constructive 1.942-bound on the integrality gap of the configuration LP using a custom-tailored Haxell-type search method. Recently, this was improved by Jansen and Rohwedder [JR17] to 1.834. In an even more restricted variant called *Graph Balancing*, each job is admissible on exactly 2 machines. In this setting Ebenlendr, Krcál and Sgall [EKS08] gave a 1.75-approximation based on an LP-rounding approach, which has again been improved by Jansen and Rohwedder [JR18b] to 1.749 using a local search argument.

³The factor is n^ε if only polynomial time is allowed, where $\varepsilon > 0$ is arbitrary but fixed.

2.1.1 Our contributions

Let $\mathcal{M} = (X, \mathcal{I})$ be a *matroid* with *groundset* X and a family of *independent sets* $\mathcal{I} \subseteq 2^X$. Recall that a matroid is characterized by three properties:

- (i) *Non-emptiness*: $\emptyset \in \mathcal{I}$;
- (ii) *Monotonicity*: For $Y \in \mathcal{I}$ and $Z \subseteq Y$ one has $Z \in \mathcal{I}$;
- (iii) *Exchange property*: For all $Y, Z \in \mathcal{I}$ with $|Y| < |Z|$ there is an element $z \in Z \setminus Y$ so that $Y \cup \{z\} \in \mathcal{I}$.

The *bases* $\mathcal{B}(\mathcal{M})$ of the matroid are all inclusion-wise maximal independent sets. The cardinalities of all bases are identical, with size denoted as $\text{rank}(\mathcal{M})$. The convex hull of all bases is called the *base polytope*, that is $P_{\mathcal{B}(\mathcal{M})} := \text{conv}\{\chi(S) \in \{0, 1\}^X \mid S \text{ is basis}\}$, where $\chi(S)$ is the *characteristic vector* of S .

Now consider a bipartite graph $G = (X \dot{\cup} W, E)$ with the ground set X on one side and a set of *resources* W on the other side; each resource $w \in W$ has a *size* $p_w \geq 0$. In a problem that we call *Matroid Max-Min Allocation*, the goal is to find a basis $S \in \mathcal{B}(\mathcal{M})$ and an assignment $\sigma : W \rightarrow S$ with $(\sigma(w), w) \in E$ so that $\min_{i \in S} \sum_{w \in \sigma^{-1}(i)} p_w$ is maximized. To the best of our knowledge, this problem has not been studied before. In particular if $T \geq 0$ is the target objective function value, then we can define a linear programming relaxation $Q(T)$ as the set of vectors $(x, y) \in \mathbb{R}_{\geq 0}^X \times \mathbb{R}_{\geq 0}^E$ satisfying the constraints

$$x \in P_{\mathcal{B}(\mathcal{M})}; \quad \sum_{w \in N(i)} p_w y_{iw} \geq T \cdot x_i \quad \forall i \in X; \quad y(\delta(w)) \leq 1 \quad \forall w \in W; \quad y_{iw} \leq x_i \quad \forall (i, w) \in E.$$

Here, the decision variable x_i expresses whether element i should be part of the basis, and y_{iw} expresses whether resource w should be assigned to element i . We abbreviate $N(i)$ as the neighborhood of i and $y(\delta(w))$ is shorthand for $\sum_{i: (i, w) \in E} y_{iw}$. Then our main technical result is:

Theorem 8. *Suppose $Q(T) \neq \emptyset$. Then for any $\varepsilon > 0$ one can find*

$$(x, y) \in Q \left(\left(\frac{1}{3} - \varepsilon \right) \cdot T - \frac{1}{3} \cdot \max_{w \in W} p_w \right)$$

with both x and y integral in time $n^{\Theta_\varepsilon(1)}$, where $n := |X| + |W|$. This assumes that membership in the matroid can be tested in time polynomial in n .

Previously this result was not even known with non-constructive methods. We see that Matroid Max-Min Allocation is a useful framework by applying it to the Santa Claus problem:

Theorem 9. *The Santa Claus problem admits a $(4 + \varepsilon)$ -approximation algorithm in time $n^{\Theta_\varepsilon(1)}$.*

For a suitable threshold $0 < \delta < 1$, call a gift j *small* if $p_j \leq \delta \cdot OPT$ and *large* otherwise. Then the family of sets of children that can get assigned large gifts forms a *matchable set matroid*. We apply Theorem 8 to the *co-matroid* of the matchable set matroid. Then we obtain a basis $S := \{i \in M \mid x_i = 1\}$, which contains the children *not* receiving a large gift. These children can receive small gifts of total value $(\frac{1}{3} - \frac{\delta}{3} - \varepsilon) \cdot OPT$. The remaining children receive a large gift with value at least $\delta \cdot OPT$. Setting $\delta := \frac{1}{4}$ implies the claim. Note the approximation factor $4 + \varepsilon$ will be with respect to a natural, compact linear program with $O(n^2)$ many variables and constraints. The smallest LP that was previously known to have a constant integrality gap was the $O(n^3)$ -size LP of [JR18a].

2.2 An algorithm for Matroid Max-Min Allocation

In this section we provide an algorithm that proves Theorem 8.

2.2.1 Intuition for the algorithm

We provide some insight by starting with an informal overview of our algorithm. Let $G = (X \cup W, E)$ be the bipartite graph defined in Section 2.1.1. If $U \subset W$ and $i \in X$ with $(i, j) \in E$ for all $j \in U$, we can consider the pair (i, U) to be a hyperedge. Then for $0 < \nu < 1$ and $\text{val}(\cdot)$ the function summing the value in a hyperedge's resources, we say that (i, U) is a ν -edge if it is a hyperedge with minimal (inclusion wise) resources such that $\text{val}(U) := \sum_{w \in U} p_w \geq \nu T$. By $\mathcal{E}_{\nu T}$ we denote the set of ν -edges.

Fix constants $0 < \beta < \alpha < 1$ and $0 < \delta < 1$, to be chosen later. The goal of the algorithm is to find a basis $S \in \mathcal{B}(\mathcal{M})$ and a hypergraph matching $M \subseteq \mathcal{E}_{\beta T}$ covering S . The algorithm is initialized with $S := \{i_0\}$, for any node $i_0 \in X$, and $M := \emptyset$. We perform $\text{rank}(\mathcal{M})$ many phases, where in each phase we find a larger matching, and the set it covers in X is independent with respect to the matroid. In an intermediate phase, we begin with $S \in \mathcal{I}$ and $M \subseteq \mathcal{E}_{\beta T}$ a hypergraph matching covering $S \setminus \{i_0\}$ with one exposed node $i_0 \in X$. At the end of a phase, the algorithm produces an updated matching covering an independent set S' , with $|S'| = |S|$. For $|S'| < \text{rank}(\mathcal{M})$, there exists $i'_0 \in X \setminus S'$ such that $S' \cup \{i'_0\} \in \mathcal{I}$. Repeating this $\text{rank}(\mathcal{M})$ times, we end with a basis which is well-covered by β -edges.

The algorithm generalizes the notion of an augmenting path used to find a maximum matchings in bipartite graphs to an *augmenting tree*. Though instead of swapping every other edge in an augmenting path, as is the case for a bipartite graph, the algorithm swaps sets of edges in the augmenting tree to find more space in the hypergraph. During a phase, the edges are swapped in such a way that the underlying set in X covered by the matching is always independent with respect to the matroid. The edges which are candidates for being swapped into the matching are called *adding edges* and denoted by A , while those which are candidates for being swapped out of the matching are called *blocking edges* and denoted by B . It is helpful to discuss the nodes covered

by adding and blocking edges in each part, and so for hyperedges $H \subseteq \mathcal{E}_{\nu T}$ we define H_X and H_W as the nodes covered by H in X and W , respectively. The algorithm gives some slack by allowing the adding edges to be slightly larger than the blocking edges.

The parameters α and β determine the value of the adding and blocking edges, respectively, so the adding edges are a subset of $\mathcal{E}_{\alpha T}$ while the blocking edges are a subset of $\mathcal{E}_{\beta T}$. Set $\delta := \max_w p_w / T$, so that all elements in the basis receive resources with value at most δT . The following observations follow from minimality of the hyperedges:

1. A ν -edge has value less than $(\nu + \delta)T$. This implies that an add edge has value less than $(\alpha + \delta)T$ and a blocking edge has value less than $(\beta + \delta)T$.
2. Every blocking edge has value at most $\beta \cdot T$ not covered by an add edge.

To build the augmenting tree, the algorithm starts from the node in S uncovered by M , i_0 , and chooses an edge $e \in \mathcal{E}_{\alpha T}$ covering i_0 which is added to A . If there is a large enough hyperedge $e' \in \mathcal{E}_{\beta T}$ such that $e' \subset e$ and e' is disjoint from M , then there is enough available resources that we simply update M by adding e' to it. Otherwise, e does not contain a set of resources with total value βT free from M . The edges of M intersecting e are added to the set of blocking edges, B . Nodes in $C = \{i_0\} \cup B_X$ are called *discovered* nodes, as they are the nodes covered by the hypermatching M which appear in the augmenting tree.

Continuing to build the augmenting tree in later iterations, the algorithm uses an *Expansion Lemma* to find a large set of disjoint hyperedges, $H \subset \mathcal{E}_{\alpha T}$, that cover a subset which can be swapped into S in place of some subset of C while maintaining independence in the matroid. The set of hyperedges H either (i) intersects many edges of M or (ii) has a constant fraction of edges which contain a hyperedge from $\mathcal{E}_{\beta T}$ that is disjoint from M .

In the first case, a subset of H which intersects M , denoted $A_{\ell+1}$, is added to A , and the edges of M intersecting $A_{\ell+1}$, denoted $B_{\ell+1}$, are added to B , for ℓ the index of the iteration. Note we

naturally obtain *layers* which partition the adding and blocking edges in our augmenting tree. The layers for the adding and blocking edges respectively are denoted as A_ℓ and B_ℓ , with

$$A_{\leq \ell} := \bigcup_{i=0}^{\ell} A_i \quad \text{and} \quad B_{\leq \ell} := \bigcup_{i=0}^{\ell} B_i.$$

The layer indices are tracked because they are useful in proving the algorithm's runtime. In the second case, for the set of edges $H' \subset \mathcal{E}_{\alpha T}$ that have a hyperedge from $\mathcal{E}_{\beta T}$ disjoint from M , the algorithm finds a layer which has a large number of discovered nodes that can be swapped out for a subset of nodes which H' covers.

2.2.2 A detailed procedure

Recall, we fixed $\delta = \max_{w \in W} p_w / T$. Then, we set $\beta = \frac{1}{3} - \frac{\delta}{3} - \varepsilon$ and $\alpha = \frac{1}{3} - \frac{\delta}{3} - \frac{\varepsilon}{2}$, for $0 < \varepsilon < (1 - \delta)/3$. Here lies the subtle but crucial difference to previous work. In [AKS15] the authors have to use adding edges that are a large constant factor bigger than blocking edges. In our setup we can allow adding edges that are only marginally larger than the blocking edges. This results in an improved approximation factor of $4 + \varepsilon$ for Santa Claus compared to the 12.33 factor by [AKS15].

The algorithm is described in Figure 2.1. For later reference, the constant from Lemma 13 is $\frac{1-2\alpha-\beta-\delta}{1+\delta} = \frac{2\varepsilon}{1+\delta} \geq \varepsilon$, and the constant from Lemma 14 is $\frac{\alpha-\beta}{\delta+\alpha} = \frac{\varepsilon}{2(1+\delta)} \geq \varepsilon/4$. We use Lemma 15, with constant $c = \frac{1-2\alpha-\beta-\delta}{1+\delta} \cdot \frac{\alpha-\beta}{\delta+\alpha}$. Our bounds for constants do not use a specific choice of δ , and instead they only use the fact that $0 < \delta < 1$. Both cases in the algorithm are visualized in Figure 2.2 and Figure 2.3.

Input: Node i_0 and set $S \in \mathcal{I}$ with $i_0 \in S$. Matching $M \subseteq \mathcal{E}_{\beta T}$ with $M_X = S \setminus \{i_0\}$.

Initialize: $A = A_0 = \emptyset$, $B = B_0 = \emptyset$, $C = \{i_0\}$, $\ell = 0$.

while TRUE **do**

Find disjoint $H \subseteq \mathcal{E}_{\alpha T}$ covering $D \subseteq (X \setminus S) \cup C$, s.t. $|D| \geq \varepsilon \cdot |C|$, $(S \setminus C) \cup D \in \mathcal{I}$, *and H_W is disjoint from $A_W \cup B_W$.

// Build the next layer in the augmenting tree

if H intersects at least $\frac{\varepsilon}{4} \cdot |H| \geq \frac{\varepsilon^2}{4} \cdot |C|$ many edges M on W -side **then**

$B \leftarrow B \cup B_{\ell+1}$, $B_{\ell+1} = \{e \in M : e \cap H \neq \emptyset\}$

// Find subset of H to add to A

for $b \in B_{\ell+1}$ **do**

Choose one edge $h_b \in H$ such that $h_b \cap b \neq \emptyset$

$A_{\ell+1} \leftarrow A_{\ell+1} \cup \{h_b\}$

end for

$C \leftarrow B_X \cup \{i_0\}$

$\ell \leftarrow \ell + 1$

// Swap sets and collapse layers

else $H' = \{e \in H : \text{val}(e_W \setminus M_W) \geq \beta T\}$ has size at least $\frac{\varepsilon}{4} \cdot |H|$

For all $e \in H'$, choose one $e' \subset e$ with $e' \in \mathcal{E}_{\beta T}$ and $e'_W \cap M_W = \emptyset$. Replace e for e' in H' .

// Find a set to swap in, \tilde{D} , and a set to swap out, \tilde{C}

$D' \subseteq D$ are the nodes covered by H'

$C' \subseteq C$ is such that $|C'| = |D'|$ and $S \setminus C' \cup D' \in \mathcal{I}$

if $i_0 \in C'$ **then**

Let $i_1 \in D'$ so that $S' := S \setminus \{i_0\} \cup \{i_1\} \in \mathcal{I}$ and let $e_1 \in H'$ be edge covering i_1 .

Return $M' := M \cup \{e_1\}$ covering all of S' and **terminate**.

end if

Layer $\tilde{\ell} \leq \ell$ contains $\tilde{C} \subset C' \cap (B_{\tilde{\ell}})_X$, with $|\tilde{C}| \geq \gamma|C'|$. **

Let $\tilde{D} \subset D'$ be such that $|\tilde{C}| = |\tilde{D}|$ and $S' := S \setminus \tilde{C} \cup \tilde{D} \in \mathcal{I}$.

$\tilde{M} \subset M$ covers \tilde{C} and $\tilde{H} \subset H'$ covers \tilde{D} .

$M \leftarrow M \setminus \tilde{M} \cup \tilde{H}$, and $S \leftarrow S'$

$A \leftarrow A_{\leq \tilde{\ell}}$, $B \leftarrow B_{\leq \tilde{\ell}} \setminus \tilde{M}$, $C \leftarrow B_X \cup \{i_0\}$

$\ell \leftarrow \tilde{\ell}$

end if***

end while

* Possible by Lemma 13 with $W' := A_W \cup B_W$.

** By Lemma 15, such a \tilde{C} exists.

*** One of the conditionals occurs by Lemma 14.

Figure 2.1: Main algorithm

2.2.3 Correctness of the algorithm

Here, we prove several lemmas used in the algorithm which implies Theorem 8. We begin by building up to our *Expansion Lemma*, Lemma 13. Our algorithm takes a fixed independent set, S , and swaps $C \subset S$ out of S for a set of nodes D in order to construct a new independent set of the same size. This is possible by Lemma 13.

Recall a variant of the so-called *Exchange Lemma*. For independent sets $Y, Z \in \mathcal{I}$, let $H_{\mathcal{M}}(Y, Z)$ denote the bipartite graph on parts Y and Z (if $Y \cap Z \neq \emptyset$, then have one copy of the intersection on the left and one on the right). For $i \in Y \setminus Z$ and $j \in Z \setminus Y$ we insert an edge (i, j) in $H_{\mathcal{M}}(Y, Z)$ if $Y \setminus \{i\} \cup \{j\} \in \mathcal{I}$. Otherwise, for $i \in Y \cap Z$, there is an edge between the left and right copies of i , and this is the only edge for both copies of i .

Lemma 10 (Exchange Lemma). *For any matroid $\mathcal{M} = (X, \mathcal{I})$ and independent set $Y, Z \in \mathcal{I}$ with $|Y| \leq |Z|$, the exchange graph $H_{\mathcal{M}}(Y, Z)$ contains a left perfect matching.*

Next, we prove several lemmas about vectors in the base polytope with respect to sets containing swappable elements. Lemma 13 relies on a *Swapping Lemma*, Lemma 12, for which the next lemma serves as a helper function.

Lemma 11 (Weak Swapping Lemma). *Let $\mathcal{M} = (X, \mathcal{I})$ be a matroid with an independent set $S \in \mathcal{I}$. For $C \subseteq S$, define*

$$U := \{i \in (X \setminus S) \cup C \mid (S \setminus C) \cup \{i\} \in \mathcal{I}\}.$$

Then for any vector $x \in P_{B(\mathcal{M})}$ in the base polytope one has $\sum_{i \in U} x_i \geq |C|$.

Proof. Note that in particular $C \subseteq U$. Moreover, an equivalent definition of U is

$$U = \{i \in (X \setminus S) \cup C \mid \exists j \in C : (S \setminus \{j\}) \cup \{i\} \in \mathcal{I}\}.$$

Due to the integrality of the base polytope, there is a basis $B \in \mathcal{I}$ with $\sum_{i \in U} x_i \geq \sum_{i \in U} (\chi(B))_i = |U \cap B|$, where $\chi(B) \in \{0, 1\}^X$ is the characteristic vector of B . As S and B are independent sets with $|S| \leq |B|$, from Lemma 10 there is a left-perfect matching in the exchange graph $H_{\mathcal{M}}(S, B)$. The neighborhood of C in $H_{\mathcal{M}}(S, B)$ is $U \cap B$. As there is a left-perfect matching, $|B \cap U|$ is least $|C|$ and hence $\sum_{i \in U} x_i \geq |U \cap B| \geq |C|$. \square

Next, we derive a more general form of the Swapping Lemma (which coincides with the previous Lemma 11 if $D = \emptyset$):

Lemma 12 (Strong Swapping Lemma). *Let $\mathcal{M} = (X, \mathcal{I})$ be a matroid with an independent set $S \in \mathcal{I}$. Let $C \subseteq S$ and $D \subseteq (X \setminus S) \cup C$ with $|D| \leq |C|$ and $S \setminus C \cup D \in \mathcal{I}$. Define*

$$U := \{i \in ((X \setminus S) \cup C) \setminus D \mid S \setminus C \cup D \cup \{i\} \in \mathcal{I}\}.$$

Then for any vector $x \in P_{B(\mathcal{M})}$ in the base polytope one has $\sum_{i \in U} x_i \geq |C| - |D|$.

Proof. Partition $C = C_1 \dot{\cup} C_2$ so that $C \cap D \subseteq C_1$, $|C_1| = |D|$ and $S' := S \setminus C_1 \cup D \in \mathcal{I}$. Then note that

$$\begin{aligned} U &= \left\{ i \in X \setminus \underbrace{(S \setminus C \cup D)}_{=S' \setminus C_2} \mid \underbrace{S \setminus C \cup D}_{=S' \setminus C_2} \cup \{i\} \in \mathcal{I} \right\} \\ &= \{i \in (X \setminus S') \cup C_2 \mid S' \setminus C_2 \cup \{i\} \in \mathcal{I}\}. \end{aligned}$$

Then applying Lemma 11 gives

$$\sum_{i \in U} x_i \geq |C_2| = |C| - |D|.$$

□

Having proved our swapping lemma, we are equipped to prove the Expansion Lemma. Note that in our algorithm, layers are built to ensure that $|A_{\ell+1}| \leq |B_{\ell+1}|$. Due to this and the minimality of the edges in $\mathcal{E}_{\alpha T}$ and $\mathcal{E}_{\beta T}$, $W' := A_W \cup B_W$ has $\text{val}(W') \leq (\alpha + \beta + \delta)T \cdot |C|$.

Lemma 13 (Expansion Lemma). *Let $C \subseteq S \in \mathcal{I}$, $W' \subseteq W$ with $\text{val}(W') \leq (\alpha + \beta + \delta)T \cdot |C|$. Further, let $\mu := \frac{1-2\alpha-\beta-\delta}{1+\delta} > 0$ and assume that there exists $(x, y) \in Q(T)$. Then there is a set $D \subseteq (X \setminus S) \cup C$ of size $|D| \geq \lceil \mu \cdot |C| \rceil$ covered by a matching $H \subseteq \mathcal{E}_{\alpha T}$ so that $H_W \cap W' = \emptyset$ and $(S \setminus C) \cup D \in \mathcal{I}$.*

Proof. Note that D may contain elements from C . Greedily choose D and the matching H with $|D| = |H|$ one node/edge after the other. Suppose the greedy procedure gets stuck — no edge can be added without intersecting $W' \cup H_W$. For the sake of contradiction assume this happens when $|D| < \mu|C|$. First, let

$$U := \{i \in ((X \setminus S) \cup C) \setminus D \mid (S \setminus C) \cup D \cup \{i\} \in \mathcal{I}\}$$

be the nodes which could be added to D while preserving independence. Then for our fixed $x \in P_{\mathcal{B}(\mathcal{M})}$, by Lemma 12 one has

$$\sum_{i \in U} x_i \geq |C| - |D| > (1 - \mu) \cdot |C|.$$

Let $W'' := W' \cup H_W$ be the right hand side resources that are being covered by the augmenting tree. Here, we let $W' = A_W \cup B_W$. Using the minimality of the adding and blocking edges,

$$\text{val}(W'') \leq \mu|C|(\alpha + \delta)T + |C|(\beta + \alpha + \delta)T = |C|T(\mu(\alpha + \delta) + \beta + \alpha + \delta).$$

By the assumption that the greedy procedure is stuck, there is no edge $e \in \mathcal{E}_{\alpha T}$ with $e_X \in U$ and $e \cap W'' = \emptyset$. If $N(i)$ denotes the neighborhood of $i \in X$ in the bipartite graph G , then this means that $\text{val}(N(i) \setminus W'') < \alpha T$ for all $i \in U$. For every fixed $i \in U$ we can then lower bound the y -weight going into W'' as

$$\begin{aligned} \sum_{(i,w) \in E: w \in W''} p_w y_{i,w} &= \underbrace{\sum_{w \in \delta(i)} p_w y_{i,w}}_{\geq T x_i} - \sum_{(i,w) \in E: w \notin W''} p_w \underbrace{y_{i,w}}_{\leq x_i} \\ &\geq T x_i - x_i \left(\underbrace{\sum_{(i,w) \in E: w \notin W''} p_w}_{< \alpha T} \right) \geq T \cdot x_i \cdot (1 - \alpha) \quad (*) \end{aligned}$$

Then double counting the y -weight running between U and W'' with a lower and upper bound shows that

$$(1 - \alpha)T \underbrace{\sum_{i \in U} x_i}_{\geq (1-\mu)|C|} \leq \sum_{(i,w) \in E: i \in U, w \in W''} p_w y_{i,w} \leq \sum_{w \in W''} p_w \underbrace{y(\delta(w))}_{\leq 1} \leq \text{val}(W'').$$

Simplifying the above,

$$(1 - \alpha) \cdot (1 - \mu) \cdot T|C| < (\mu(\alpha + \delta) + \beta + \alpha + \delta) \cdot T|C| \quad \Rightarrow \quad \frac{1 - 2\alpha - \beta - \delta}{1 + \delta} < \mu.$$

Thus we reach a contradiction for our choice of μ . □

The algorithm relies on the fact that from the set of hyperedges, H , guaranteed by the Expansion Lemma, there is either some constant fraction of H to swap into the matching, or a constant fraction of H is blocked by edges in the current matching. In the former, significant space is found in W for S . In the latter, enough edges of the matching are intersected to guarantee the next layer in the augmenting tree is large. The following lemma proves at least one of these conditions occurs.

Lemma 14. Set $\mu := \frac{\alpha-\beta}{\delta+\alpha} > 0$. Let $M \subseteq \mathcal{E}_{\beta T}$ and $F \subseteq \mathcal{E}_{\alpha T}$ both be hypergraph matchings. Further, let

$$H := \{e \in F \mid \text{val}(e_W \setminus M_W) \geq \beta T\}$$

be the edges in F that still have value βT after overlap with M is removed. Then either (i) $|H| \geq \mu|F|$ or (ii) F intersects at least $\mu|F|$ edges of M .

Proof. Let $W' := M_W \cap F_W$ be the right hand side nodes where the hypermatchings overlap and suppose for the sake of contradiction that neither of the two cases occur. Then double counting the value of W' gives

$$\mu \cdot (\beta + \delta) \cdot T \cdot |F| > (\beta + \delta) T \cdot \underbrace{(\# \text{edges in } M \text{ intersecting } W')}_{\mu|F|} \geq \text{val}(W') \geq \underbrace{|F \setminus H|}_{\geq (1-\mu) \cdot |F|} \cdot (\alpha - \beta) \cdot T.$$

Rearranging and simplifying, the above implies $\mu > \frac{\alpha-\beta}{\delta+\alpha}$. Thus we contradict our choice of μ . \square

Our last lemma will show that a constant fraction of the nodes which could be swapped out of the augmenting tree come from the same *layer* in the tree. This allows us to swap out enough nodes from the same layer to make substantial progress with each iteration. Here C' and \tilde{C} are labelled the same as in the algorithm.




Figure 2.2: Case 1 of the algorithm, where a set $A_{\ell+1} \subseteq \mathcal{E}_{\alpha T}$ of hyperedges is found that intersects many new edges $B_{\ell+1} \subseteq (M \setminus B_{\leq \ell})$. In particular $|B_{\ell+1}| \geq \Omega_\varepsilon(|C|)$. Note that D might contain nodes from C .

Lemma 15. *Let sets C' and $\{B_i\}_{i=0}^\ell$ be such that $C' \subset (B_{\leq \ell})_X$. Further, suppose there exists constant $c > 0$ such that $|C'| \geq c \cdot |B_{\leq \ell}|$ and $|B_{i+1}| \geq c \cdot |B_{\leq i}|$ for $i = 0, \dots, \ell - 1$. Then, there exists a layer $0 \leq \tilde{\ell} \leq \ell$ and constant $\gamma := \gamma(c) > 0$, such that $\tilde{C}' := C' \cap (B_{\tilde{\ell}})_X$ has size $|\tilde{C}'| \geq \gamma \cdot |C'|$.*

Proof. By induction, $|B_{\leq \ell}|$ can be written in terms of lower indexed sets as

$$|B_{\leq \ell}| \geq (1 + c)^k \cdot |B_{\leq \ell-k}|,$$

for $k = 0, \dots, \ell$. Therefore, the size of C' can be written as $|C'| \geq c(1 + c)^k \cdot |B_{\leq \ell-k}|$. As c is a constant, take k large enough so $c(1 + c)^k \geq 2$, namely $k \geq \frac{\log(\frac{2}{c})}{\log(1+c)}$. Then the collection of sets $(B_{\ell-i})_X$ for $i = 0, \dots, k$ contain at least half of C' , so one of them must contain at least $\gamma = \frac{1}{2(k+1)}$ of C' . \square




Figure 2.3: Case 2 of the algorithm, where $\tilde{H} \subseteq \mathcal{E}_{\beta T}$ of size $|\tilde{H}| \geq \Omega_\epsilon(|C|)$ is found so that (i) \tilde{H} is disjoint on the W -side to the matching M and the adding edges in the augmenting tree, (ii) \tilde{H} covers a set \tilde{D} with $S \setminus \tilde{C} \cup \tilde{D} \in \mathcal{I}$, and (iii) \tilde{C} is from one layer of the augmenting tree. Here \tilde{D} and \tilde{C} do not have to be disjoint.

2.2.4 Termination and runtime

As seen in Lemma 15,

$$|X| \geq |B_{\leq \ell}| \geq \left(1 + \frac{\epsilon^2}{4}\right)^\ell |B_0|,$$

and solving for ℓ shows $\frac{\log(|X|)}{\log\left(1 + \frac{\epsilon^2}{4}\right)} \geq \ell$. Thus the total number of layers at any step in the algorithm is $O(\log |X|)$. Note after each collapse of the layers, the matching M and possibly the independent set S are updated. However, the fixed exposed node i_0 will remain in S until the very last iteration in which the algorithm finds an edge e_1 that augments the matching. Before we begin discussing the proof guaranteeing our algorithm terminates, we need a lemma to compare the number of blocking edges after a layer is collapsed to the number of blocking edges at the beginning of the iteration.

Lemma 16. *Let $\tilde{\ell}$ be the index of the collapsed layer and let B' be the updated blocking edges after a collapse step. Then, $|B'_{\leq \tilde{\ell}}| \leq |B_{\leq \tilde{\ell}}| \cdot \left(1 - \frac{\epsilon^2}{4} \cdot \gamma\right)$.*

Proof. Recall $B'_\ell = B_\ell \setminus F$ for F the edges of M covering \tilde{C} . Further, the blocking edges in layers

indexed less than $\tilde{\ell}$ are not effected in the iteration. Hence

$$|B'_{\leq \tilde{\ell}}| = |B'_{\leq \tilde{\ell}-1}| + |B'_\ell| = |B_{\leq \tilde{\ell}-1}| + |B'_\ell|$$

From Lemmas 13 and 14, $|B_{\ell+1}| \geq \frac{\varepsilon^2}{4}|B_{\leq \ell}|$. Then examining the collapsed layer by itself, we see

$$|B'_\ell| = |B_\ell| - |F| \leq |B_\ell| - \frac{\varepsilon^2}{4} \cdot \gamma |B_{\leq \ell}|.$$

Substituting back into $|B'_{\leq \tilde{\ell}}|$, we find that

$$\begin{aligned} |B'_{\leq \tilde{\ell}}| &\leq |B_{\leq \tilde{\ell}-1}| + |B_\ell| - \frac{\varepsilon^2}{4} \cdot \gamma |B_{\leq \ell}| \\ &= |B_{\leq \tilde{\ell}}| - \frac{\varepsilon^2}{4} \cdot \gamma |B_{\leq \ell}| = |B_{\leq \tilde{\ell}}| \cdot \left(1 - \frac{\varepsilon^2}{4} \cdot \gamma\right). \end{aligned}$$

□

To prove the algorithm terminates in polynomial time, we consider a signature vector $s = (s_0, s_1, \dots, s_\ell, \infty)$, where $s_j = \lfloor \log_c |B_{\leq j}| \rfloor$ for $c = \frac{1}{1 - \frac{\varepsilon^2}{4} \cdot \gamma}$. The signature vector and proof that the algorithm terminates is inspired by [AKS15], but it is subtly different.

Lemma 17. *The signature vector decreases lexicographically after each iterative loop in the algorithm.*

Proof. Let $s = (s_0, \dots, s_\ell, \infty)$ be a signature vector at the beginning of a step in the algorithm, and let s' be the result of s through one iteration of the algorithm. For $\ell + 1$ denoting the newest built layer in the algorithm, if the newest set of hyperedges found intersects at least $\frac{\varepsilon^2}{4}|C|$ many edges of M , then another layer in the augmenting tree is built and no layer is collapsed. Then $s' = (s_0, \dots, s_\ell, s_{\ell+1}, \infty)$ is lexicographically smaller than s .

Otherwise, layer $0 \leq \tilde{\ell} \leq \ell$ is collapsed. All finite coordinates above $s_{\tilde{\ell}}$ are deleted from the signature vector, and all coordinates before $s_{\tilde{\ell}}$ are unaffected. So it suffices to check that $s'_{\tilde{\ell}} < s_{\tilde{\ell}}$. Again, let B' be the updated blocking edges after a collapse step. As $B_{\tilde{\ell}}$ is the only set of blocking edges in $B_{\leq \tilde{\ell}}$ affected by the collapse, by Lemma 16 one has $|B'_{\leq \tilde{\ell}}| \leq |B_{\leq \tilde{\ell}}|(1 - \frac{\varepsilon^2}{4} \cdot \gamma)$. Taking a log we compare the coordinates

$$s'_{\tilde{\ell}} = \left\lfloor \log_c \left(|B'_{\leq \tilde{\ell}}| \right) \right\rfloor \leq \left\lfloor \log_c \left(|B_{\leq \tilde{\ell}}| \left(1 - \frac{\varepsilon^2}{4} \cdot \gamma \right) \right) \right\rfloor = \left\lfloor \log_c (|B_{\leq \tilde{\ell}}|) \right\rfloor - 1 = s_{\tilde{\ell}} - 1.$$

□

Choose the infinite coordinate to be some integer larger than $\log |X|$. Since for every layer ℓ , we have $|B_{\leq \ell}| \leq |X|$, then every coordinate of the signature vector is upper bounded by $U = O(\log |X|)$. Recall the number of layers, and thus the number of coordinates in the signature vector, is also upper bounded by U . Together, these imply that the sum of the coordinates of the signature vector is at most U^2 .

As the signature vector has non-decreasing order, each signature vector corresponds to a partition of an integer $z \leq U^2$. On the other hand, every partition of some $z \leq U^2$ has a corresponding signature vector. Thus we apply a result of Hardy and Ramanujan to find the total number of signature vectors is $\sum_{k \leq U^2} e^{O(\sqrt{k})} = |X|^{O(1)}$. Since each iteration of the algorithm can be done in polynomial time and the signature vector decreases lexicographically after each iteration, the algorithm terminates after a total time of $n^{\Theta_\varepsilon(1)}$.

2.3 Application to Santa Claus

In this section, we show a polynomial time $(4 + \varepsilon)$ -approximation algorithm for the Santa Claus problem. Recall that for a given set of children M , and a set of presents J , the Santa Claus

problem asks how Santa should distribute presents to children in order to maximize the minimum happiness of any child⁴. Here, present j is only wanted by some subset of children that we denote by $A_j \subseteq M$, and present j has value p_j to child $i \in A_j$. The happiness of child i is the sum of all p_j for presents j assigned to child i . We assume w.l.o.g. to know the integral objective function value T of the optimum solution, otherwise T can be found by binary search.

We partition gifts into two sets: *large* gifts $J_L := \{j \in J \mid p_j > \delta_2 T\}$ and *small* gifts $J_S := \{j \in J \mid p_j \leq \delta_1 T\}$, for parameters $0 < \delta_1 \leq \delta_2 < 1$ such that all gifts have values in $[0, \delta_1 T] \cup (\delta_2 T, T]$. Let $P(T, \delta_1, \delta_2)$ be the set of vectors $z \in \mathbb{R}_{\geq 0}^{J \times M}$ satisfying

$$\begin{aligned} \sum_{j \in J_S: i \in A_j} p_j z_{ij} &\geq T \cdot \left(1 - \sum_{j \in J_L: i \in A_j} z_{ij}\right) && \forall i \in M \\ \sum_{i \in A_j} z_{ij} &\leq 1 && \forall j \in J \\ z_{ij} &\leq 1 - \sum_{j' \in J_L: i \in A_{j'}} z_{ij'} && \forall j \in J_S \forall i \in A_j \end{aligned}$$

If $n = |J| + |M|$, then this LP has $O(n^2)$ many variables and $O(n^2)$ many constraints. To see that this is indeed a relaxation, take any feasible assignment $\sigma : J \rightarrow M$ with $\sum_{j \in \sigma^{-1}(i)} p_j \geq T$ for all $i \in M$. Now let $\sigma : J \rightarrow M \cup \{\emptyset\}$ be a modified assignment where we set $\sigma(j) = \emptyset$ for gifts that we decide to drop. For each child $i \in M$ that receives at least one large gift we drop all small gifts and all but one large gift. Then a feasible solution $z \in P(T, \delta_1, \delta_2)$ is obtained by letting

$$z_{ij} := \begin{cases} 1 & \text{if } \sigma(j) = i \\ 0 & \text{otherwise.} \end{cases}$$

We will show that given a feasible solution $z \in P(T, \delta_1, \delta_2)$, there exists a feasible solution

⁴We assume Santa to be an equitable man– not one influenced by bribery, social status, etc.

(x^*, y^*) to $Q(T)$. To do this, we will exploit two underlying matroids in the Santa Claus problem, allowing us to apply Theorem 8. Let

$$\mathcal{I} = \{M_L \subseteq M \mid \exists \text{ left-perfect matching between } M_L \text{ and } J_L \text{ using edges } (i, j) : i \in A_j\},$$

be a family of independent sets. Then $\mathcal{M} = (M, \mathcal{I})$ constitutes a *matchable set matroid*. We denote the *co-matroid* of \mathcal{M} by $\mathcal{M}^* = (M, \mathcal{I}^*)$. Recall that the independent sets of the co-matroid are given by

$$\mathcal{I}^* = \{M_S \subseteq M \mid \exists M_L \in \mathcal{B}(\mathcal{M}) : M_S \cap M_L = \emptyset\}.$$

We can define a vector $x \in \mathbb{R}^M$ with $x_i = \sum_{j \in J_L : i \in A_j} z_{ij}$ that lies in the matroid polytope of \mathcal{M} . This fact follows easily from the integrality of the fractional matching polytope in bipartite graphs. It is instructive to think of x_i as the decision variable telling whether child $i \in M$ should receive a large present.

Unfortunately, x does not have to lie in the base polytope — in fact the sum $\sum_{i \in M} x_i$ might not even be integral. However, there always exists a vector x' in the base polytope that covers every child just as well with large presents as x does. This observation can be stated for general matroids:

Lemma 18. *Let $\mathcal{M} = (X, \mathcal{I})$ be any matroid and let x be a point in its matroid polytope. Then in polynomial time one can find a point x' in the base polytope so that $x' \geq x$ coordinate-wise.*

In fact the algorithm behind this claim is rather trivial: as long as $x \in P_{\mathcal{M}}$ is not in the base polytope, there is always a coordinate i and a $\mu > 0$ so that $x + \mu e_i \in P_{\mathcal{M}}$.

With the new vector $x' \in P_{\mathcal{B}(\mathcal{M})}$ at hand, we can redefine the z -assignments by letting

$$z'_{ij} = \begin{cases} z_{ij} & x_i = 1 \\ \frac{1-x'_i}{1-x_i} z_{ij} & x_i \neq 1. \end{cases}$$

for $j \in J_S$; the new values z'_{ij} for $j \in J_L$ can be obtained from the fractional matching that corresponds to x'_i . Note that $0 \leq z'_{ij} \leq z_{ij}$ for $j \in J_S$. The reader should be convinced that still $z' \in P(T, \delta_1, \delta_2)$, just that the corresponding vector x' now lies in $P_{\mathcal{B}(\mathcal{M})}$ ⁵.

It is well known in matroid theory that the complementary vector $x^* := \mathbb{K} - x'$ lies in $P_{\mathcal{B}(\mathcal{M}^*)}$. Again, it is instructive to think of x^*_i as the decision variable whether child i has to be satisfied with small gifts. Finally, the assignments y^* are simply the restriction of z' on the coordinates $(i, j) \in M \times J_S$. The obtained pair (x^*, y^*) lies in $Q(T)$, where the matroid in the definition of $Q(T)$ is \mathcal{M}^* .

As $Q(T) \neq \emptyset$, we can apply Theorem 8 which results in a subset $M_S \in \mathcal{B}(\mathcal{M}^*)$ of the children and an assignment $\sigma : J_S \rightarrow M_S$, where each child in M_S receives happiness at least $(\frac{1}{3} - \frac{\delta_1}{3} - \varepsilon) \cdot T$ from the assignment of small gifts. Implicitly due to the choice of the matroid \mathcal{M}^* , we know that the remaining children $M \setminus M_S = M_L$ can all receive one large gift and this assignment can be computed in polynomial time using a matching algorithm. Overall, each child receives either one large present of value at least $\delta_2 \cdot T$ or small presents of total value at least $(\frac{1}{3} - \frac{\delta_1}{3} - \varepsilon) \cdot T$. Therefore each child receives value at least

$$\min \left\{ \left(\frac{1}{3} - \frac{\delta_1}{3} - \varepsilon \right) \cdot T, \delta_2 \cdot T \right\} \geq \left(\frac{1}{4} - \varepsilon \right) \cdot T \quad (2.1)$$

⁵There is an alternative proof without the need to replace x by x' . Add the constraint $\sum_{j \in J_L, i \in A_j} z_{ij} = \text{rank}(\mathcal{M})$ to $P(T, \delta_1, \delta_2)$. There is always a feasible integral solution satisfying this constraint. Then for any fractional solution $z \in P(T, \delta_1, \delta_2)$, the corresponding vector x will immediately lie in the base polytope.

for a choice of $\delta_2 = \delta_1 = \frac{1}{4}$. In some instances of Santa Claus, we can do better. Set δ_1 so that $\delta_1 \cdot T$ is the largest gift value that is at most $\frac{1}{4}T$, and set δ_2 so that $\delta_2 \cdot T$ is the smallest gift value that is at $\frac{1}{4}T$. Then the algorithm guarantees that each child receives value at least as in the left hand side of Equation 2.1. When δ_1 and δ_2 are bounded away from $1/4$, then the approximation improves. For example, when $\delta_2 \geq 1/3$ and $\delta_1 T$ is close to 0, such as in the case where all gifts have value either T or 1, we approach a $(3 + \varepsilon)$ -approximation.

2.4 Conclusion

In this work, we introduced Matroid Max-Min Fair Allocation. for this problem, we construct a new, compact LP to model it, and we prove a $(3 + \epsilon)$ -approximation, modulo an additive term of $\frac{1}{3}$ of the most valuable item's value $\frac{1}{3} \max_{w \in W} p_w$. Our algorithm is a local search based on Haxell's augmenting tree argument. Overall, we can use our result on Matroid Max-Min Fair Allocation as a blackbox to obtain a $(4 + \epsilon)$ -approximation for Santa Claus.

One obvious, immediate open question is to resolve the gap between the lower bound of 2 and the upper bound of 4 for approximation Santa Claus. We do not know the integrality gap for our LP, so it is unclear whether the LP can be used to obtain approximations better than 4. More broadly, this problem of Matroid Max-Min Fair Allocation might be useful to study other scheduling problems.

Chapter 3

Scheduling with Communication Delays on Identical Machines

3.1 Introduction

Scheduling jobs with precedence constraints is a fundamental problem in approximation algorithms and combinatorial optimization. In this problem we are given m identical machines and a set J of n jobs, where each job j has a processing length $p_j \in \mathbb{Z}_+$. The jobs have precedence constraints, which are given by a partial order \prec . A constraint $j \prec j'$ encodes that job j' can only start after job j is completed. The goal is to find a schedule of jobs that minimizes *makespan*, which is the completion time of the last job. This problem is denoted¹ by $P \mid \text{prec} \mid C_{\max}$. In a seminal result from 1966, Graham [Gra66] showed that the greedy list scheduling algorithm achieves a $(2 - \frac{1}{m})$ -approximation. By now, our understanding of the approximability of this

¹Throughout the paper we use the standard scheduling three-field notation [GLLK79, VLL90]. The respective fields denote: **(1) number of identical machines:** P_∞ : unlimited; P : number m of machines given as input; Pm : constant number m of machines, **(2) job properties:** prec : precedence constraints; $p_j = 1$: unit-size jobs; c : communication delays of length c (can be c_{jk} if dependent on jobs $j \prec k$); c -intervals: see Section 3.3; dup : allowed duplication of jobs, **(3) objective:** C_{\max} : minimize makespan; $\sum w_j C_j$: minimize weighted sum of completion times.

basic problem is almost complete: it had been known since the late ‘70s, due to a result by Lenstra and Rinnooy Kan [LRK78], that it is NP-hard to obtain better than $4/3$ -approximation, and in 2010 Svensson [Sve10] showed that, assuming a variant of the Unique Games Conjecture [BK10], it is NP-hard to get a $(2 - \varepsilon)$ -approximation for any $\varepsilon > 0$.

The above precedence-constrained scheduling problem models the task of distributing workloads onto multiple processors or servers, which is ubiquitous in computing. This basic setting takes the dependencies between work units into account, but not the data transfer costs between machines, which is critical in applications. A precedence constraint $j \prec j'$ typically implies that the input to j' depends on the output of j . In many real-world scenarios, especially in the context of scheduling in data centers, if j and j' are executed on different machines, then the *communication delay* due to transferring this output to the other machine cannot be ignored. This is an active area of research in applied data center scheduling literature, where several new abstractions have been proposed to deal with communication delays [CZM⁺11, GFC⁺12, HCG12, SZA⁺18, ZZC⁺12, ZCB⁺15, LYZ⁺16]. Another timely example is found in the parallelization of Deep Neural Network training (the machines being accelerator devices such as GPUs, TPUs, or FPGAs). There, when training the network on one sample/minibatch per device in parallel, the communication costs incurred by synchronizing the weight updates in fact dominate the overall running time [NHP⁺19]. Taking these costs into account, it turns out that it is better to split the network onto multiple devices, forming a “model-parallel” computation pipeline [HCB⁺19]. In the resulting *device placement* problem, the optimal split crucially depends on the communication costs between dependent layers/operators.

A classic model that captures the effect of data transfer latency on scheduling decisions is the problem of *scheduling jobs with precedence and communication delay constraints*, introduced by Rayward-Smith [RS87] and Papadimitriou and Yannakakis [PY90]. The setting, denoted by

$P \mid \text{prec}, c \mid C_{\max}$, is similar to the makespan minimization problem described earlier, except for one crucial difference. Here we are given a *communication delay parameter* $c \in \mathbb{Z}_{\geq 0}$, and the output schedule must satisfy the property that if $j \prec j'$ and j, j' are scheduled on different machines, then j' can only start executing at least c time units after j had finished. On the other hand, if j and j' are scheduled on the same machine, then j' can start executing immediately after j finishes. In a closely related problem, denoted by $P_{\infty} \mid \text{prec}, c \mid C_{\max}$, a schedule can use as many machines as desired. The goal is to schedule jobs *non-preemptively* so as to minimize the makespan. In a non-preemptive schedule, each job j needs to be assigned to a single machine and executed during p_j consecutive timeslots. The problems $P \mid \text{prec}, c \mid C_{\max}$ and $P_{\infty} \mid \text{prec}, c \mid C_{\max}$ are the focus of this paper.

Despite its theoretical significance and practical relevance, very little is known about the communication delay setting. A direct application of Graham's [Gra66] list scheduling algorithm yields a $(c + 2)$ -approximation, and no better algorithm is known for the problem. Over the years, the problem has attracted significant attention, but all known results, which we discuss below in Section 3.1.4, concern special settings, small communication delays, or hardness of approximation. To put this in perspective, we note that the current best algorithm for general c [GKMP08], which achieves an approximation factor of $2/3 \cdot (c + 1)$, only marginally improves on Graham's algorithm while requiring the additional assumptions that the number of machines is unbounded and $p_j = 1$. This is in sharp contrast to the basic problem $P \mid \text{prec} \mid C_{\max}$ (which would correspond to the case $c = 0$), where the approximability of the problem is completely settled under a variant of the Unique Games Conjecture. This situation hints that incorporating communication delays in scheduling decisions requires fundamentally new algorithmic ideas compared to the no-delay setting. Schuurman and Woeginger [SW99] placed the quest for getting better algorithms to the problem in their influential list of top-10 open problems in scheduling theory. In a recent MAPSP 2017 survey talk, Bansal [Ban17] highlighted the lack of progress on this model, describing it

as “not understood at all; almost completely open”, and suggested that this is due to the lack of promising LP/SDP relaxations.

3.1.1 Our Contributions

The main result of this paper is the following:

Theorem 19. *There is a randomized $O(\log c \cdot \log m)$ -approximation algorithm for $P \mid \text{prec}, c \mid C_{\max}$ with expected polynomial running time, where $c, p_j \in \mathbb{N}$.*

In any non-preemptive schedule the number m of machines is at most the number n of jobs, so for the easier P_{∞} version of the problem, the above theorem implies the following:

Corollary 20. *There is a randomized $O(\log c \cdot \log n)$ -approximation algorithm for $P_{\infty} \mid \text{prec}, c \mid C_{\max}$ with expected polynomial running time, where $c, p_j \in \mathbb{N}$.*

For both problems one can replace either c or m by n , yielding a $O(\log^2 n)$ -approximation algorithm. Our results make substantial progress towards resolving one of the questions in “Open Problem 3” in the survey of Schuurman and Woeginger [SW99], which asks whether a constant-factor approximation algorithm exists for $P_{\infty} \mid \text{prec}, c \mid C_{\max}$.

Our approach is based on a Sherali-Adams lift of a time-indexed linear programming relaxation for the problem, followed by a randomized clustering of the semimetric space induced by this lift. It is possible to write a more compact LP for the problem for which our algorithm and analysis still work. However, as we will show in Section 3.5, this compact LP has integrality gap $\Omega(\sqrt{\log n})$. To our knowledge, this is the first instance of a multiple-machine scheduling problem being viewed via the lens of metric space clustering. We believe that our framework is fairly general and should extend to other problems involving scheduling with communication delays. To demonstrate the

broader applicability of our approach, we also consider the objective of minimizing the weighted sum of completion times. Here each job j has a weight w_j , and the goal is to minimize $\sum_j w_j C_j$, where C_j is the completion time of j .

Theorem 21. *There is a randomized $O(\log c \cdot \log n)$ -approximation algorithm for $P_\infty \mid \text{prec}, p_j = 1, c \mid \sum_j w_j C_j$ with expected polynomial running time, where $c \in \mathbb{N}$.*

No non-trivial approximation ratio was known for this problem prior to our work.

3.1.2 Independent work of Maiti *et al*

In a parallel and independent work, Maiti et al. [MRS⁺20] developed an $O(\log^2 n \log^2 m \log c / \log \log c)$ -approximation algorithm for the makespan objective function on *related machines* ($Q \mid \text{prec}, c \mid C_{\max}$). Interestingly, they obtained the results using completely different techniques compared to ours. While our results are based on LP hierarchies and clustering, Maiti et al. [MRS⁺20] developed a novel framework based on *job duplication*. In their framework, they first construct a schedule where a single job can be scheduled on *multiple machines*, which is known to effectively “hide” the communication delay constraints [PY90]. Quite surprisingly, Maiti et al. [MRS⁺20] showed that one can convert a schedule with duplication to a feasible schedule without duplication, where every job is processed on a single machine, while increasing the makespan by at most an $O(\log^2 n \log m)$ factor. Maiti et al. also prove an integrality gap for the LP they consider, which is different than ours. However, after seeing their result, we found that their integrality gap instance works for a compact form of our LP as well, which we prove in Section 3.5.

3.1.3 Our Techniques

As we alluded earlier, there is a lack of combinatorial lower bounds for scheduling with communication delays. For example, consider Graham's list scheduling algorithm, which greedily processes jobs on m machines as soon as they become available. One can revisit the analysis of Graham [Gra66] and show that there exists a chain Q of dependent jobs such that the makespan achieved by list scheduling is bounded by

$$\frac{1}{m} \sum_{j \in J} p_j + \sum_{j \in Q} p_j + c \cdot (|Q| - 1).$$

The first two terms are each lower bounds on the optimum — the 3rd term is not. In particular, it is unclear how to certify that the optimal makespan is high because of the communication delays. However, this argument suffices for a $(c + 2)$ -approximation, since $p_j \geq 1$ for all $j \in J$.

As pointed out by Bansal [Ban17], there is no known promising LP relaxation. To understand the issue let us consider the special case $P_\infty \mid \text{prec}, p_j = 1, c \mid C_{\max}$. Extending, for example, the LP of Munier and König [MK97], one might choose variables C_j as completion times, as well as decision variables x_{j_1, j_2} denoting whether j_2 is executed in the time window $[C_{j_1}, C_{j_1} + c)$ on the same machine as j_1 . Then we can try to enforce communication delays by requiring that $C_{j_2} \geq C_{j_1} + 1 + (c - 1) \cdot (1 - x_{j_1, j_2})$ for $j_1 \prec j_2$. Further, we enforce load constraints $\sum_{j_1 \in J} x_{j_1, j_2} \leq c$ for $j_2 \in J$ and $\sum_{j_2 \in J} x_{j_1, j_2} \leq c$ for $j_1 \in J$. To see why this LP fails, note that in any instance where the maximum dependence degree is bounded by c , one could simply set $x_{j_1, j_2} = 1$ and completely avoid paying any communication delay. Moreover, this problem seems to persist when moving to more complicated LPs that incorporate indices for time and machines.

A convenient observation is that, in exchange for a constant-factor loss in the approximation guarantee, it suffices to find an assignment of jobs to length- c intervals such that dependent jobs scheduled in the same length- c interval must be assigned to the same machine. (The latter

condition will be enough to satisfy the communication delay constraints as, intuitively, between every two length- c intervals we will insert an empty one.) In order to obtain a stronger LP relaxation, we consider an $O(1)$ -round *Sherali-Adams lift* of an initial LP with indices for time and machines. From the lifted LP, we extract a *distance function* $d : J \times J \rightarrow [0, 1]$ which satisfies the following properties:

- (i) The function d is a *semimetric*.
- (ii) $C_{j_1} + d(j_1, j_2) \leq C_{j_2}$ for $j_1 \prec j_2$.
- (iii) Any set $U \subseteq J$ with a diameter of at most $\frac{1}{2}$ w.r.t. d , satisfies $|U| \leq 2c$.

Here we have changed the interpretation of C_j to the *index* of the length- c interval in which j will be processed. Intuitively, $d(j_1, j_2)$ can be understood as the probability that jobs j_1, j_2 are *not* being scheduled within the same length- c interval on the same machine. To see why a constant number of Sherali-Adams rounds are helpful, observe that the triangle inequality behind (i) is really a property depending only on *triples* $\{j_1, j_2, j_3\}$ of jobs and an $O(1)$ -round Sherali-Adams lift would be locally consistent for every triple of variables.

We will now outline how to round such an LP solution. For jobs whose LP completion times are sufficiently different, say $C_{j_1} + \Theta(\frac{1}{\log(n)}) \leq C_{j_2}$, we can afford to deterministically schedule j_1 and j_2 at least c time units apart while only paying a $O(\log n)$ -factor more than the LP. Hence the critical case is to sequence a set of jobs $J^* = \{j \in J \mid C^* \leq C_j \leq C^* + \Theta(\frac{1}{\log(n)})\}$ whose LP completion times are very close to each other. Note that by property (ii), we know that any dependent jobs $j_1, j_2 \in J^*$ must have $d(j_1, j_2) \leq \Theta(\frac{1}{\log(n)})$. As d is a semimetric, we can make use of the rich toolset from the theory of metric spaces. In particular, we use an algorithm by Calinescu, Karloff and Rabani [CKR04]: For a parameter $\Delta > 0$, one can partition a semimetric space into *random clusters* so that the diameter of every cluster is bounded by Δ and each δ -neighborhood around a node is separated, meaning contains jobs assigned to different clusters,

with probability at most $O(\log(n)) \cdot \frac{\delta}{\Delta}$. Setting $\delta := \Theta(\frac{1}{\log(n)})$ and $\Delta := \Theta(1)$ one can then show that a fixed job $j \in J^*$ will be in the same cluster as *all* its ancestors in J^* with probability at least $\frac{1}{2}$, while all clusters have diameter at most $\frac{1}{2}$. By (iii), each cluster will contain at most $2c$ many (unit-length) jobs, and consequently we can schedule all the clusters in parallel, where we drop any job that got separated from any ancestor. Repeating the sampling $O(\log n)$ times then schedules all jobs in J^* . This reasoning results in a $O(\log^2 n)$ -approximation for this problem, which we call $P_\infty \mid \text{prec}, p_j = 1, c\text{--intervals} \mid C_{\max}$. With a bit of care the approximation factor can be improved to $O(\log c \cdot \log m)$.

Finally, the promised $O(\log c \cdot \log m)$ -approximation for the more general problem $P \mid \text{prec}, c \mid C_{\max}$ follows from a reduction to the described special case $P_\infty \mid \text{prec}, p_j = 1, c\text{--intervals} \mid C_{\max}$.

3.1.4 History of the Problem

Precedence-constrained scheduling problems of minimizing the makespan and sum of completion times objectives have been extensively studied for many decades in various settings. We refer the reader to [Pin18, LLRKS93, PST04, AMMS08, Sve09] for more details. Below, we only discuss results directly related to the communication delay problem in the offline setting.

Approximation algorithms. As mentioned earlier, Graham's [Gra66] list scheduling algorithm yields a $(c + 2)$ -approximation for $P \mid \text{prec}, c \mid C_{\max}$, and a $(c + 1)$ -approximation for the P_∞ variant. For unit-size jobs and $c \geq 2$, Giroudeau, König, Moulai and Palaysi [GKMP08] improved the latter ($P_\infty \mid \text{prec}, p_j = 1, c \geq 2 \mid C_{\max}$) to a $\frac{2}{3}(c + 1)$ -approximation. For unit-size jobs and $c = 1$, Munier and König [MK97] obtained a $4/3$ -approximation via LP rounding ($P_\infty \mid \text{prec}, p_j = 1, c = 1 \mid C_{\max}$); for the P variant, Hanen and Munier [HM01] gave an easy reduction from the P_∞ variant that loses an additive term of 1 in the approximation ratio, thus yielding a $7/3$ -approximation. Thurimella and Yesha [THU92] gave a reduction that, given an

α -approximation algorithm for $P_\infty \mid \text{prec}, c, p_j = 1 \mid C_{\max}$, would yield a $(1 + 2\alpha)$ -approximation algorithm for $P \mid \text{prec}, c, p_j = 1 \mid C_{\max}$.

For a constant number of machines, a hierarchy-based approach of Levey and Rothvoss [LR16] for the no-delay setting ($Pm \mid \text{prec}, p_j = 1 \mid C_{\max}$) was generalized by Kulkarni, Li, Tarnawski and Ye [KLTY20] to allow for communication delays that are also bounded by a constant. For any $\varepsilon > 0$ and $\hat{c} \in \mathbb{Z}_{\geq 0}$, they give a nearly quasi-polynomial-time $(1 + \varepsilon)$ -approximation algorithm for $Pm \mid \text{prec}, p_j = 1, c_{jk} \leq \hat{c} \mid C_{\max}$. The result also applies to arbitrary job sizes, under the assumption that preemption of jobs is allowed, but migration is not.

Hardness. Hoogeveen, Lenstra and Veltman [HLV94] showed that even the special case $P_\infty \mid \text{prec}, p_j = 1, c = 1 \mid C_{\max}$ is NP-hard to approximate to a factor better than $7/6$. For the case with bounded number of machines (the P variant) they show $5/4$ -hardness. These two results can be generalized for $c \geq 2$ to $(1 + 1/(c+4))$ -hardness [GKMP08] and $(1 + 1/(c+3))$ -hardness [BGK96], respectively.²

Duplication model. The communication delay problem has also been studied (to a lesser extent) in a setting where jobs can be duplicated (replicated), i.e., executed on more than one machine, in order to avoid communication delays. This assumption seems to significantly simplify the problem, especially when we are also given an unbounded number of machines: already in 1990, Papadimitriou and Yannakakis [PY90] gave a rather simple 2-approximation algorithm for $P_\infty \mid \text{prec}, p_j, c_{jk}, \text{dup} \mid C_{\max}$. Observe that this result holds even when communication delays are unrelated (they depend on the pair of jobs). The only non-trivial approximation algorithm for arbitrary c and a bounded number of machines is due to Lepere and Rapine [LR02], who gave an asymptotic $O(\log c / \log \log c)$ -approximation for $P \mid \text{prec}, p_j = 1, c, \text{dup} \mid C_{\max}$. On the hardness

²Papadimitriou and Yannakakis [PY90] claim a 2-hardness for $P_\infty \mid \text{prec}, p_j = 1, c \mid C_{\max}$, but give no proof. Schuurman and Woeginger [SW99] remark that “it would be nice to have a proof for this claim”.

side, Papadimitriou and Yannakakis [PY90] showed NP-hardness of $P_\infty \mid \text{prec}, p_j = 1, c, \text{dup} \mid C_{\max}$ (using a large delay $c = \Theta(n^{2/3})$).

Many further references can be found in [VLL90, GKMP08, Dro09, GK07, CC91, JKS93, MH97].

3.2 Preliminaries

3.2.1 The Sherali-Adams Hierarchy for LPs with Assignment Constraints

In this section, we review the *Sherali-Adams hierarchy* which provides an automatic strengthening of linear relaxations for 0/1 optimization problems. The authoritative reference is certainly Laurent [Lau03], and we adapt the notation from Friggstad et al. [FKK⁺14]. Consider a set of variable indices $[n] = \{1, \dots, n\}$ and let $U_1, \dots, U_N \subseteq [n]$ be subsets of variable indices. We consider a polytope

$$K = \left\{ x \in \mathbb{R}^n \mid \tilde{A}x \geq \tilde{b}, \sum_{i \in U_k} x_i = 1 \quad \forall k \in [N], \quad 0 \leq x_i \leq 1 \quad \forall i \in [n] \right\},$$

which we also write in a more compact form as $K = \{x \in \mathbb{R}^n \mid Ax \geq b\}$ with $A \in \mathbb{R}^{m \times n}$ and $b \in \mathbb{R}^m$. We note that we included explicitly the “box constraints” $0 \leq x_i \leq 1$ for all variables i . Moreover, the constraint matrix contains *assignment constraints* of the form $\sum_{i \in U_k} x_i = 1$. This is the aspect that is non-standard in our presentation.

The general goal is to obtain a strong relaxation for the integer hull $\text{conv}(K \cap \{0, 1\}^n)$. Observe that any point $x \in \text{conv}(K \cap \{0, 1\}^n)$ can be interpreted as a *probability distribution* X over points $K \cap \{0, 1\}^n$. We know that any distribution can be described by the 2^n many values $y_I = \Pr[\bigwedge_{i \in I} (X_i = 1)]$ for $I \subseteq [n]$ — in fact, the probability of any other event can be reconstructed using the *inclusion-exclusion formula*, for example $\Pr[X_1 = 1 \text{ and } X_2 = 0] =$

$y_{\{1\}} - y_{\{1,2\}}$. While this is an exact approach, it is also an inefficient one. In order to obtain a polynomial-size LP, we only work with variables y_I where $|I| \leq O(1)$. Hence, for $r \geq 0$, we denote $\mathcal{P}_r([n]) := \{S \subseteq [n] \mid |S| \leq r\}$ as all the index sets of size at most r .

Definition 2. Let $SA_r(K)$ be the set of vectors $y \in \mathbb{R}^{\mathcal{P}_{r+1}([n])}$ satisfying $y_\emptyset = 1$ and

$$\sum_{H \subseteq J} (-1)^{|H|} \cdot \left(\sum_{i=1}^n A_{\ell,i} y_{I \cup H \cup \{i\}} - b_\ell y_{I \cup H} \right) \geq 0 \quad \forall \ell \in [m]$$

for all $I, J \subseteq [n]$ with $|I| + |J| \leq r$.

The parameter r in the definition is usually called the *rank* or *number of rounds* of the Sherali-Adams lift. It might be helpful for the reader to verify that for $I = J = \emptyset$, the constraint simplifies to $\sum_{i=1}^n A_{\ell,i} y_{\{i\}} \geq b_\ell y_\emptyset = b_\ell$, which implies that $(y_{\{1\}}, \dots, y_{\{n\}}) \in K$. Moreover it is instructive to verify that for any feasible integral solution $x \in K \cap \{0, 1\}^n$ one can set $y_I := \prod_{i \in I} x_i$ to obtain a vector $y \in SA_r(K)$.

Theorem 22 (Properties of Sherali-Adams). *Let $y \in SA_r(K)$ for some $r \geq 0$. Then the following holds:*

- (a) *For $J \in \mathcal{P}_r([n])$ with $y_J > 0$, the vector $\tilde{y} \in \mathbb{R}^{\mathcal{P}_{r+1-|J|}([n])}$ defined by $\tilde{y}_I := \frac{y_{I \cup J}}{y_J}$ satisfies $\tilde{y} \in SA_{r-|J|}(K)$.*
- (b) *One has $0 \leq y_I \leq y_J \leq 1$ for $J \subseteq I$ and $|I| \leq r + 1$.*
- (c) *If $|J| \leq r + 1$ and $y_i \in \{0, 1\} \forall i \in J$, then $y_I = y_{I \setminus J} \cdot \prod_{i \in I \cap J} y_i$ for all $|I| \leq r + 1$.*
- (d) *For $J \subseteq [n]$ with $|J| \leq r$ there exists a distribution over vectors \tilde{y} such that (i) $\tilde{y} \in SA_{r-|J|}(K)$, (ii) $\tilde{y}_i \in \{0, 1\}$ for $i \in J$, (iii) $y_I = \mathbb{E}[\tilde{y}_I]$ for all $I \subseteq [n]$ with $|I \cup J| \leq r + 1$ (this includes in particular all $I \in \mathcal{P}_{r+1-|J|}([n])$).*

- (e) For $I \subseteq [n]$ with $|I| \leq r$ and $k \in [N]$ one has $y_I = \sum_{i \in U_k} y_{I \cup \{i\}}$.
- (f) Take $H \subseteq [N]$ with $|H| \leq r$ and set $J := \bigcup_{k \in H} U_k$. Then there exists a distribution over vectors \tilde{y} such that (i) $\tilde{y} \in SA_{r-|H|}(K)$, (ii) $\tilde{y}_i \in \{0, 1\}$ for $i \in J$, (iii) $y_I = \mathbb{E}[\tilde{y}_I]$ for all $I \in \mathcal{P}_{r+1-|H|}([n])$.

Proof. For (a)-(d), we refer to the extensive coverage in Laurent [Lau03]. We prove (e) and (f) which are non-standard and custom-tailored to LPs with assignment constraints:

- (e) Fix $I \subseteq [n]$ with $|I| \leq r$. We apply (d) to obtain a distribution over \tilde{y} with $\tilde{y} \in SA_{r-|I|}(K)$ so that $\tilde{y}_i \in \{0, 1\}$ for $i \in I$. Then

$$\sum_{i \in U_k} y_{I \cup \{i\}} \stackrel{\text{linearity}}{=} \mathbb{E} \left[\sum_{i \in U_k} \tilde{y}_{I \cup \{i\}} \right] \stackrel{(c)}{=} \mathbb{E} \left[\tilde{y}_I \cdot \underbrace{\sum_{i \in U_k} \tilde{y}_i}_{=1} \right] = \mathbb{E}[\tilde{y}_I] = y_I.$$

Here we apply (c) for index sets $I \cup \{i\}$ where variables in $J := I$ have been made integral. Note that indeed $|I \cup (I \cup \{i\})| \leq r + 1$ as required.

- (f) By an inductive argument it suffices to consider the case of $|H| = 1$. Let $H = \{k\}$ and set $U := U_k$, i.e. the constraints for polytope P contain the assignment constraint $\sum_{i \in U} x_i = 1$ and we want to make all variables in U integral while only losing a *single* round in the hierarchy. Abbreviate $U^+ := \{i \in U \mid y_{\{i\}} > 0\}$. For $i \in U^+$, define $y^{(i)} \in \mathbb{R}^{\mathcal{P}_r([n])}$ to be the vector with $y_I^{(i)} := \frac{y_{I \cup \{i\}}}{y_i}$. By (a) we know that $y^{(i)} \in SA_{r-1}(K)$. Moreover $y_{\{i\}}^{(i)} = \frac{y_{\{i\}}}{y_{\{i\}}} = 1$. Then the assignment constraint of the LP forces that $y_{\{i'\}}^{(i)} = 0$ for $i' \in U \setminus \{i\}$. Now we define a probability distribution over vectors \tilde{y} as follows: for $i \in U^+$, with probability y_i we set $\tilde{y} := y^{(i)}$. Then (i) and (ii) hold for \tilde{y} as discussed. Property (iii) follows from

$$\mathbb{E}[\tilde{y}_I] = \sum_{i \in U^+} y_i y_I^{(i)} = \sum_{i \in U^+} y_i \frac{y_{I \cup \{i\}}}{y_i} = \sum_{i \in U^+} y_{I \cup \{i\}} \stackrel{(b)}{=} \sum_{i \in U} y_{I \cup \{i\}} \stackrel{(e)}{=} y_I$$

□

It is known that Theorem 22.(f) holds in a stronger form for the SDP-based *Lasserre hierarchy*. Karlin, Mathieu and Nguyen [KMN11] proved a result that can be paraphrased as follows: *if one has any set $J \subseteq [n]$ of variables with the property that there is no LP solution with more than k ones in J , then one can make all variables of J integral while losing only k rounds*. Interestingly, Karlin, Mathieu and Nguyen prove that this is completely false for Sherali-Adams. In particular, for a Knapsack instance with unit size items and capacity $2 - \varepsilon$, the integrality gap is still $2 - 2\varepsilon$ after $\Theta_\varepsilon(n)$ rounds of Sherali-Adams. In a different setting, Friggstad et al. [FKK⁺14] realized that given a “tree constraint”, a Sherali-Adams lift can provide the same guarantees that Rothvoss [Rot11] derived from Lasserre. While Friggstad et al. did not state their insight in the generality that we need here, our Lemma 22.(e)+(f) are inspired by their work.

3.2.2 Semimetric Spaces

A *semimetric space* is a pair (V, d) where V is a finite set (we denote $n := |V|$) and $d : V \times V \rightarrow \mathbb{R}_{\geq 0}$ is a *semimetric*, i.e.

- $d(u, u) = 0$ for all $u \in U$.
- Symmetry: $d(u, v) = d(v, u)$ for all $u, v \in V$.
- Triangle inequality: $d(u, v) + d(v, w) \geq d(u, w)$ for all $u, v, w \in V$.

Recall that the more common notion is that of a *metric*, which additionally requires that $d(u, v) > 0$ for $u \neq v$. For a set $U \subseteq V$ we denote the *diameter* as $\text{diam}(U) := \max_{u, v \in U} d(u, v)$. Our goal is to find a partition $V = V_1 \dot{\cup} \dots \dot{\cup} V_k$ such that the diameter of every cluster V_i is bounded by some parameter Δ . We denote $d(w, U) := \min\{d(w, u) : u \in U\}$ as the distance to the set U . Moreover, for $r \geq 0$ and $U \subseteq V$, let $N(U, r) := \{v \in V \mid d(v, U) \leq r\}$ be the *distance r -neighborhood* of U .

We use a very influential clustering algorithm due to Calinescu, Karloff and Rabani [CKR04], which assigns each $v \in V$ to a random cluster center $c \in V$ such that $d(u, c) \leq \beta\Delta$. Nodes assigned to the same cluster center form one block V_i in the partition. Formally the algorithm is as follows:

CKR CLUSTERING ALGORITHM
Input: Semimetric space (V, d) with $V = \{v_1, \dots, v_n\}$, parameter $\Delta > 0$
Output: Clustering $V = V_1 \dot{\cup} \dots \dot{\cup} V_k$ for some k .
<ol style="list-style-type: none"> (1) Pick a uniform random $\beta \in [\frac{1}{4}, \frac{1}{2}]$ (2) Pick a random ordering $\pi : V \rightarrow \{1, \dots, n\}$ (3) For each $v \in V$ set $\sigma(v) := v_\ell$ so that $d(v, v_\ell) \leq \beta \cdot \Delta$ and $\pi(v_\ell)$ is minimal (4) Denote the points $v \in V$ with $\sigma^{-1}(v) \neq \emptyset$ by $c_1, \dots, c_k \in V$ and return clusters $V_i := \sigma^{-1}(c_i)$ for $i = 1, \dots, k$

Note that the algorithm has two sources of randomness: it picks a random parameter β , and independently it picks a random ordering π . Here the ordering is to be understood such that element v_ℓ with $\pi(v_\ell) = 1$ is the “highest priority” element. The original work of Calinescu, Karloff and Rabani [CKR04] only provided an upper bound on the probability that a short edge (u, v) is separated. Mendel and Naor [MN06] note that the same clustering provides the guarantee of $\Pr[N(u, t) \text{ separated}] \leq 1 - O(\frac{t}{\Delta} \cdot \ln(\frac{|N(u, \Delta)|}{|N(u, \Delta/8)|}))$ for all $u \in V$ and $0 \leq t < \frac{\Delta}{8}$. Mendel and Naor attribute this to Fakcharoenphol, Rao and Talwar [FRT04] (while Fakcharoenphol, Rao and Talwar [FRT04] do not state it explicitly in this form and focus on the “local growth ratio” aspect). Instead of the algorithm by Calinescu, Karloff and Rabani [CKR04], one could also cluster using the techniques of Leighton and Rao [LR99] or those of Garg, Vazirani and Yannakakis [GVY93]. We state the formal claim in a form that will be convenient for us.

Theorem 23 (Analysis of CKR). *Let $V = V_1 \dot{\cup} \dots \dot{\cup} V_k$ be the random partition of the CKR algorithm.*

The following holds:

- (a) The blocks have $\text{diam}(V_i) \leq \Delta$ for $i = 1, \dots, k$.
- (b) Let $U \subseteq V$ be a subset of points. Then

$$\Pr[U \text{ is separated by clustering}] \leq \ln(2|N(U, \Delta/2)|) \cdot \frac{4\text{diam}(U)}{\Delta} \leq \ln(2n) \cdot \frac{4\text{diam}(U)}{\Delta}.$$

In the above, *separated* means that there is more than one index i with $V_i \cap U \neq \emptyset$.

Proof. The claim from (a) is easy to show as

$$\text{diam}(V_i) = \max_{u,v \in V_i} d(u, v) \leq 2 \max_{u \in V_i} \underbrace{d(u, c_i)}_{\leq \beta \Delta} \leq \Delta$$

The tricky part is to show part (b). The following definition and lemma are needed.

Definition 3. Let us say that a node w is a *separator* for U , if

- (A) $\sigma(u) = w$ for at least one $u \in U$
- (B) $\sigma(u) \neq w$ for at least one $u \in U$

Moreover, if the set of separators of U is non-empty, then we call the separator that comes first in the order π the *first separator*.

Next, we show that nodes that are closer to the set U are the most likely to be the first separator:

Lemma 24. Let w_1, \dots, w_n be the nodes sorted so that $d(w_1, U) \leq \dots \leq d(w_n, U)$. Then

$$\Pr[w_s \text{ is the first separator for } U] \leq \frac{4}{s} \cdot \frac{\text{diam}(U)}{\Delta}.$$

Proof. Let $u_{\min} := \text{argmin}\{d(u, w_s) : u \in U\}$ and $u_{\max} := \text{argmax}\{d(u, w_s) : u \in U\}$ be the closest and furthest point from w_s . We claim that in order for w_s to be the first separator, both of the following conditions must hold:

- (i) $d(w_s, u_{\min}) \leq \beta \cdot \Delta < d(w_s, u_{\max})$
- (ii) The order selects w_s as the first node among w_1, \dots, w_s .

We assume that w_s is the first separator, and suppose for the sake of contradiction that either (i) or (ii) (or both) are not satisfied. We verify the cases:

- *Case:* $\beta\Delta < d(w_s, u_{\min})$. Then no point will be assigned to w_s and w_s is not a separator at all.
- *Case:* $\beta\Delta \geq d(w_s, u_{\max})$. As w_s is a separator, there are nodes $u_1, u_2 \in U$ with $\sigma(u_1) = w_s$ and $\sigma(u_2) \neq w_s$. Then $\sigma(u_2)$ has to come earlier in the order π as $d(w_s, u_2) \leq \beta\Delta$. Hence w_s is not the first separator.
- *Case:* w_s is not first among w_1, \dots, w_s with respect to π . By assumption there is an index $1 \leq s_2 < s$ such that $\pi(w_{s_2}) < \pi(w_s)$. As w_s is a separator, there is a $u_1 \in U$ with $\sigma(u_1) = w_s$. Let $u_2 := \operatorname{argmin}\{d(u, w_{s_2}) : u \in U\}$ be the point in the set U that is closest to w_{s_2} . Then $d(u_2, w_{s_2}) = d(w_{s_2}, U) \leq d(w_s, U) \leq d(w_s, u_1) \leq \beta\Delta$. Hence u_2 would be assigned to a point of order at most $\pi(w_{s'}) < \pi(w_s)$, and therefore w_s is not the first separator.

Now we estimate the probability that w_s is the first separator. The parameter β and the permutation are chosen independently, so (i) and (ii) are independent events. Clearly $\Pr[(ii)] = \frac{1}{s}$. Moreover

$$\Pr[(i)] = \frac{|[d(w_s, u_{\min}), d(w_s, u_{\max})] \cap [\frac{\Delta}{4}, \frac{\Delta}{2}]|}{\Delta/4} \leq \frac{4d(u_{\min}, u_{\max})}{\Delta} \leq \frac{4\operatorname{diam}(U)}{\Delta},$$

where we have used the triangle inequality and the notation $|[a, b]| = b - a$ for the length of an interval. □

Now we can finish the proof of Theorem 23:

Proof of Theorem 23. As in Lemma 24, let w_1, \dots, w_n be an order of nodes such that $d(w_1, U) \leq \dots \leq d(w_n, U)$. Note that $L := |N(U, \frac{\Delta}{2})| \leq n$ is the maximal index with $d(w_L, U) \leq \frac{\Delta}{2}$. If U is separated, then there has to be a first separator. Therefore, the following holds:

$$\begin{aligned} \Pr[U \text{ is separated}] &\leq \sum_{s=1}^L \Pr[w_s \text{ is first separator for } U] \\ &\stackrel{\text{Lem 24}}{\leq} \underbrace{\sum_{s=1}^L \frac{1}{s}}_{\leq \ln(2L)} \cdot \frac{4\text{diam}(U)}{\Delta} \leq \ln(2L) \cdot \frac{4\text{diam}(U)}{\Delta} \end{aligned}$$

Here we use that $\Pr[w_s \text{ is first separator for } U] = 0$ for $s > L$ since a node w_s that has a distance bigger than $\frac{\Delta}{2}$ to U will never be a separator. □

□

3.3 Approximation for $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$

In this section, we provide an approximation algorithm for scheduling n unit-length jobs J with communication delay $c \in \mathbb{N}$ on an unbounded number of machines so that precedence constraints given by a partial order \prec are satisfied. Instead of working with $P_\infty \mid \text{prec}, p_j = 1, c \mid C_{\max}$ directly, it will be more convenient to consider a slight variant that we call $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$. This problem variant has the same input but the time horizon is partitioned into time *intervals* of length c , say $I_s = [sc, (s+1)c)$ for $s \in \mathbb{Z}_{\geq 0}$. The goal is to assign jobs to intervals and machines. We require that if $j_1 \prec j_2$ then either j_1 is scheduled in an earlier interval than j_2 or j_1 and j_2 are scheduled in the same interval on the same machine. Other than that, there are no further communication delays. The objective function is to minimize the number of intervals used to process the jobs. In fact we do not need to decide the order of jobs within




Figure 3.1: Left: example of an instance of $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$ with $c = 4$ (where the partial order \prec is the transitive closure of the depicted digraph). Right: a valid schedule in 2 intervals.

intervals as any topological order will work. In a more mathematical notation, the problem asks to find a partition $J = \dot{\bigcup}_{s \in \{0, \dots, S-1\}, i \in \mathbb{N}} J_{s,i}$ with $|J_{s,i}| \leq c$ such that S is minimized and for every $j_1 \prec j_2$ with $j_1 \in J_{s_1, i_1}$ and $j_2 \in J_{s_2, i_2}$ one has either $s_1 < s_2$ or $(s_1, i_1) = (s_2, i_2)$. See Figure 3.1 for an illustration.

It is rather straightforward to give reductions between $P_\infty \mid \text{prec}, p_j = 1, c \mid C_{\max}$ and $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$ that only lose a small constant factor in both directions. The only subtle point to consider here is that when the optimum makespan for $P_\infty \mid \text{prec}, c \mid C_{\max}$ is less than c , the problem admits a PTAS; we refer to Section 3.4 for details.

3.3.1 The Linear Program

Let $m \in \mathbb{N}$ be a parameter defining the number of machines that we admit for the LP. Moreover, let $S \in \mathbb{N}$ be the number of intervals that we allow for the time horizon. To obtain an approximation for $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$ one can set $m := n$ and perform a binary search to find the minimal S for which the LP is feasible. But we prefer to keep the approach general.

We construct the LP in two steps. First consider the variables

$$x_{j,i,s} = \begin{cases} 1 & \text{if } j \text{ is scheduled on machine } i \text{ in interval } I_s \\ 0 & \text{otherwise} \end{cases} \quad \forall j \in J, i \in [m], s \in \{0, \dots, S-1\}$$

Let K be the set of fractional solutions to the following linear system

$$\begin{aligned} \sum_{i \in [m]} \sum_{s \geq 0} x_{j,i,s} &= 1 \quad \forall j \in J \\ \sum_{j \in J} x_{j,i,s} &\leq c \quad \forall i \in [m] \quad \forall s \in \{0, \dots, S-1\} \\ 0 \leq x_{j,i,s} &\leq 1 \quad \forall j \in J, i \in [m], s \in \{0, \dots, S-1\} \end{aligned}$$

So far, K simply assigns jobs (fractionally) to intervals and machines without taking any precedence constraints into account. Next, we will use a lift $x \in \text{SA}_r(K)$ containing variables $x_{(j_1, i_1, s_1), (j_2, i_2, s_2)}$, which provide the probability for the event that j_1 is scheduled in interval s_1 on machine i_1 and j_2 is scheduled in interval s_2 on machine i_2 . We introduce two more types of decision variables:

$$\begin{aligned} y_{j_1, j_2} &= \begin{cases} 1 & j_1 \text{ and } j_2 \text{ are scheduled on the same machine in the same interval} \\ 0 & \text{otherwise} \end{cases} \\ C_j &= \text{index of interval where } j \text{ is processed} \end{aligned}$$

Let $Q(r)$ be the set of vectors (x, y, C) that satisfy

$$\begin{aligned} y_{j_1, j_2} &= \sum_{s \in \{0, \dots, S-1\}} \sum_{i \in [m]} x_{(j_1, i, s), (j_2, i, s)} \\ C_{j_2} &\geq C_{j_1} + (1 - y_{j_1, j_2}) \quad \forall j_1 \prec j_2 \\ C_j &\geq 0 \quad \forall j \in J \end{aligned}$$

$$x \in SA_r(K)$$

The analysis of our algorithm will work for all $r \geq 5$ while solving the LP takes time $n^{O(r)}$. Here we make no attempt at optimizing the constant r . The main technical contribution of this section is the following rounding result:

Theorem 25. *Consider an instance with unit-length jobs J , a partial order \prec , and parameters $c, S, m \in \mathbb{N}$ such that $Q(r)$ is feasible for $r := 5$. Then there is a randomized algorithm with expected polynomial running time that finds a schedule for $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$ using at most $O(\log m \cdot \log c) \cdot S$ intervals.*

We would like to emphasize that we require \prec to be a partial order, which implies that it is transitive. While replacing any acyclic digraph with its transitive closure does not change the set of feasible integral schedules and hence can be done in a preprocessing step, it corresponds to adding constraints to the LP that we rely on in the algorithm and in its analysis.

We will now discuss some properties that are implied by the Sherali-Adams lift:

Lemma 26. *Let $(x, y, C) \in Q(r)$ with $r \geq 2$. Then for any set $\tilde{J} \subseteq J$ of $|\tilde{J}| \leq r - 2$ jobs, there exists a distribution $\mathcal{D}(\tilde{J})$ over pairs (\tilde{x}, \tilde{y}) such that*

- (A) $\tilde{x}_{j,i,s} \in \{0, 1\}$ for all $j \in \tilde{J}$, all $i \in [m]$ and $s \geq 0$.
- (B) $\tilde{y}_{j_1,j_2} = \sum_{s \geq 0} \sum_{i \in [m]} \tilde{x}_{j_1,i,s} \cdot \tilde{x}_{j_2,i,s}$ if $|\{j_1, j_2\} \cap \tilde{J}| \geq 1$.
- (C) $\tilde{x} \in K$, $\tilde{y}_{j_1,j_2} = \sum_{s \in \{0, \dots, S-1\}} \sum_{i \in [m]} \tilde{x}_{(j_1,i,s), (j_2,i,s)}$ for all $j_1, j_2 \in J$.
- (D) $\mathbb{E}[\tilde{x}_{j,i,s}] = x_{j,i,s}$ and $\mathbb{E}[\tilde{y}_{j_1,j_2}] = y_{j_1,j_2}$ for all j, j_1, j_2, i, s .

Proof. By Theorem 22.(f), there is a distribution over $\tilde{x} \in SA_2(K)$ which satisfies (A) and has $\tilde{x} \in K$, $\mathbb{E}[\tilde{x}_{j,i,s}] = x_{j,i,s}$ and $\mathbb{E}[\tilde{x}_{(j_1,i_1,s_1), (j_2,i_2,s_2)}] = x_{(j_1,i_1,s_1), (j_2,i_2,s_2)}$, and additionally is integral on variables involving only jobs from \tilde{J} , where $|\tilde{J}| \leq r - 2$. Here, we crucially use that every job

$j \in \tilde{J}$ is part of an assignment constraint $\sum_{i \in [m]} \sum_{s \geq 0} x_{j,i,s} = 1$, hence making these variables integral results in the loss of only one round per job. Then, the y -variables are just linear functions depending on the x -variables, so we can define

$$\tilde{y}_{j_1, j_2} := \sum_{s \in \{0, \dots, S-1\}} \sum_{i \in [m]} \tilde{x}_{(j_1, i, s), (j_2, i, s)}$$

and the claim follows. \square

From the LP solution, we define a semimetric d . Here the intuitive interpretation is that a small distance $d(j_1, j_2)$ means that the LP schedules j_1 and j_2 mostly on the same machine and in the same interval.

Lemma 27. *Let $(x, y, C) \in Q(r)$ be a solution to the LP with $r \geq 5$. Then $d(j_1, j_2) := 1 - y_{j_1, j_2}$ is a semimetric.*

Proof. The first two properties from the definition of a semimetric (see Section 3.2.2) are clearly satisfied. We verify the triangle inequality. Consider three jobs $j_1, j_2, j_3 \in J$. We apply Lemma 26 with $\tilde{J} := \{j_1, j_2, j_3\}$ and consider the distribution $(\tilde{x}, \tilde{y}) \sim \mathcal{D}(\tilde{J})$. For $j \in \tilde{J}$, define $Z(j) = (\tilde{s}(j), \tilde{i}(j))$ as the random variable that gives the unique pair of indices such that $\tilde{x}_{j, \tilde{i}(j), \tilde{s}(j)} = 1$. Then for $j', j'' \in \tilde{J}$ one has

$$d(j', j'') = \Pr[Z(j') \neq Z(j'')] = \Pr[(\tilde{s}(j'), \tilde{i}(j')) \neq (\tilde{s}(j''), \tilde{i}(j''))]$$

Then indeed

$$\begin{aligned} d(j_1, j_3) &= \Pr[Z(j_1) \neq Z(j_3)] \leq \Pr[Z(j_1) \neq Z(j_2) \vee Z(j_2) \neq Z(j_3)] \\ &\stackrel{\text{union bound}}{\leq} \Pr[Z(j_1) \neq Z(j_2)] + \Pr[Z(j_2) \neq Z(j_3)] = d(j_1, j_2) + d(j_2, j_3). \end{aligned}$$

□

Lemma 28. *For every $j_1 \in J$ one has $\sum_{j_2 \in J} y_{j_1, j_2} \leq c$.*

Proof. Consider the distribution $(\tilde{x}, \tilde{y}) \sim \mathcal{D}(\{j_1\})$. From Lemma 26.(B) we know that $\mathbb{E}[\tilde{y}_{j_1, j_2}] = y_{j_1, j_2}$ and $\tilde{y}_{j_1, j_2} = \sum_{s \in \{0, \dots, S-1\}} \sum_{i \in [m]} \tilde{x}_{j_1, i, s} \cdot \tilde{x}_{j_2, i, s}$. By linearity it suffices to prove that $\sum_{j_2 \in J} \tilde{y}_{j_1, j_2} \leq c$ always. Fix a pair (\tilde{x}, \tilde{y}) . There is a unique pair of indices (i_1, s_1) with $\tilde{x}_{j_1, i_1, s_1} = 1$. Then

$$\sum_{j_2 \in J} \tilde{y}_{j_1, j_2} = \sum_{s \in \{0, \dots, S-1\}} \sum_{j_2 \in J} \sum_{i \in [m]} \underbrace{\tilde{x}_{j_1, i, s}}_{0 \text{ if } i \neq i_1 \text{ or } s \neq s_1} \cdot \tilde{x}_{j_2, i, s} = \sum_{j_2 \in J} \tilde{x}_{j_2, i_1, s_1} \leq c.$$

□

A crucial insight is that for any job j^* , few jobs are very close to j^* with respect to d .

Lemma 29. *Fix $j^* \in J$ and abbreviate $U := \{j \in J \mid d(j, j^*) \leq \beta\}$ for $0 < \beta < 1$. Then $|U| \leq \frac{c}{1-\beta}$.*

Proof. For each $j \in U$ we have $y_{j, j^*} = 1 - d(j, j^*) \geq 1 - \beta$. Combining with the last lemma we have $(1 - \beta)|U| \leq \sum_{j \in J} y_{j, j^*} \leq c$. □

3.3.2 Scheduling a Single Batch of Jobs

We now come to the main building block of our algorithm. We consider a subset J^* of jobs whose LP completion times C_j are very close (within a $\Theta(\frac{1}{\log(c)})$ term of each other) and show we can schedule half of these jobs in a single length- $2c$ interval. The following lemma is the main technical contribution of the paper.

Lemma 30. *Let $(x, y, C) \in Q(r)$ with $r \geq 5$ and let $0 < \delta \leq \frac{1}{64 \log(4c)}$ be a parameter. Let $C^* \geq 0$ and set $J^* \subseteq \{j \in J \mid C^* \leq C_j \leq C^* + \delta\}$. Then there is a randomized rounding procedure that*




Figure 3.2: Visualization of the partition $V = V_1 \dot{\cup} \dots \dot{\cup} V_k$ and the induced sets $V'_\ell \subseteq V_\ell$. Here \prec is the transitive closure of the depicted digraph.

*finds a schedule for a subset $J^{**} \subseteq J^*$ in a single interval of length at most $2c$ such that every job $j \in J^*$ is scheduled with probability at least $1 - 32 \log(4c) \cdot \delta \geq \frac{1}{2}$.*

We denote $\Gamma^-(j)$ as the predecessors of j and $\Gamma^+(j)$ as the successors, and similarly $\Gamma^{-/+}(J') = \{j \in J : \exists j' \in J' \text{ s.t. } j \in \Gamma^{-/+}(j')\}$. Again, recall that we assume \prec to be transitive. The rounding algorithm is the following:

SCHEDULING A SINGLE BATCH

- (1) Run a CKR clustering on the semimetric space (J^*, d) with parameter $\Delta := \frac{1}{4}$ and let V_1, \dots, V_k be the clusters.
- (2) Let $V'_\ell := \{j \in V_\ell \mid \Gamma^-(j) \cap J^* \subseteq V_\ell\}$ for $\ell = 1, \dots, k$.
- (3) Schedule V'_ℓ on one machine for all $\ell = 1, \dots, k$.

We now discuss the analysis. First we show that no cluster is more than a constant factor too large.

Lemma 31. *One has $|V'_\ell| \leq 2c$ for all $\ell = 1, \dots, k$.*

Proof. We know by Theorem 23 that $\text{diam}(V'_\ell) \leq \text{diam}(V_\ell) \leq \Delta < \frac{1}{2}$ where the diameter is with respect to d . Fix a job $j^* \in V'_\ell$. Then we know by Lemma 29 that there are at most $2c$ jobs j with $d(j, j^*) \leq \frac{1}{2}$ and the claim follows. \square

Next, we see that the clusters respect the precedence constraints.

Lemma 32. *The solution V'_1, \dots, V'_k is feasible in the sense that jobs on different machines do not have precedence constraints.*

Proof. Consider jobs processed on different machines, say (after reindexing) $j_1 \in V'_1$ and $j_2 \in V'_2$. If $j_1 \prec j_2$ then we did *not* have $\Gamma^-(j_2) \subseteq V'_2$. This contradicts the definition of the sets V'_ℓ . \square

A crucial property that makes the algorithm work is that predecessors of some job $j \in J^*$ must be very close in d distance.

Lemma 33. *For every $j_1, j_2 \in J^*$ with $j_1 \prec j_2$ one has $d(j_1, j_2) \leq \delta$.*

Proof. We know that

$$C^* \stackrel{j_1 \in J^*}{\leq} C_{j_1} \leq C_{j_1} + \underbrace{(1 - y_{j_1, j_2})}_{=d(j_1, j_2)} \stackrel{LP}{\leq} C_{j_2} \stackrel{j_2 \in J^*}{\leq} C^* + \delta$$

and so $d(j_1, j_2) \leq \delta$. \square

We will use the three statements above together with Theorem 23 to prove Lemma 30.

Proof of Lemma 30. We have already proven that the scheduled blocks have size $|V'_\ell| \leq 2c$ and that there are no dependent jobs in different sets of V'_1, \dots, V'_k . To finish the analysis, we need

to prove that a fixed job $j^* \in J^*$ is scheduled with good probability. Consider the set $U := \{j^*\} \cup (\Gamma^-(j^*) \cap J^*)$ of j^* and its ancestors in J^* .

Since the diameter of U is at most 2δ by Lemma 33, we can use Lemma 29 to see that $|N(U, \Delta/2)| \leq \frac{c}{1-2\delta-\Delta}$. For our choice of $\Delta = 1/4$ and $\delta \leq \frac{1}{64\log(4c)}$, $|N(U, 1/8)| \leq 2c$. From Theorem 23, the cluster is separated with probability at most $\log(4c) \cdot \frac{8\delta}{\Delta} \leq \frac{1}{2}$. \square

To schedule all jobs in J^* , we repeat the clustering procedure $O(\log m)$ times and simply schedule the remaining jobs on one machine.

Lemma 34. *Let $(x, y, C) \in Q(r)$ with $r \geq 5$. Let $C^* \geq 0$ and set $J^* \subseteq \{j \in J \mid C^* \leq C_j < C^* + \delta\}$. Assume that all jobs in $\Gamma^-(J^*) \setminus J^*$ have been scheduled respecting precedence constraints. Then there is an algorithm with expected polynomial running time that schedules all jobs in J^* using at most $O(\log m) + \frac{|J^*|}{mc}$ many intervals.*

Proof. Our algorithm in Lemma 30 schedules each $j \in J^*$ in an interval of length $2c$ with probability at least $1/2$. We run the algorithm for $2 \log m$ iterations, where input to iteration $k+1$ is the subset of jobs that are not scheduled in the first k iterations. For $k \in \{1, 2, \dots, 2 \log m\}$, let J_k^{**} denote the subset of jobs that are scheduled in the k^{th} iteration, and let $J_{k+1}^* := J^* \setminus \{\bigcup_{k'=1}^k J_{k'}^{**}\}$. In this notation, $J_1^* := J^*$. Let $\mathcal{S}(J_k^{**})$ denote the schedule of jobs J_k^{**} given by Lemma 30. We schedule $\mathcal{S}(J_1^{**})$ first, then for $k = 2, \dots, 2 \log m$, we append the schedule $\mathcal{S}(J_k^{**})$ after $\mathcal{S}(J_{k-1}^{**})$. Let $J' := J_{2 \log m+1}^*$ denote the set of jobs that were not scheduled in the $2 \log m$ iterations. We schedule all jobs in J' consecutively on a single machine after the completion of $\mathcal{S}(J_{2 \log m}^{**})$.

From our construction, the length of a schedule for J^* , which is a random variable, is at most $O(\log m) + \lceil \frac{|J'|}{c} \rceil$ many intervals. For $k \in \{1, 2, \dots, 2 \log m\}$, Lemma 30 guarantees that each job $j \in J_k^*$ gets scheduled in the k^{th} iteration with probability at least $1/2$. Therefore, the probability that $j \in J'$, i.e., it does not get scheduled in the first $2 \log m$ iterations, is at most $\frac{1}{2m}$. This implies

that $\mathbb{E}[|J'|] \leq \frac{|J^*|}{2m}$. By Markov's inequality $\Pr[|J'| > \frac{|J^*|}{m}] \leq \Pr[|J'| > 2 \cdot \mathbb{E}[|J'|]] \leq 1/2$. Hence we can repeat the described procedure until indeed we have a successful run with $|J'| \leq \frac{|J^*|}{m}$ which results in the claimed expected polynomial running time.

Let us now argue that the schedule of J^* is feasible. For $k \in \{1, 2, \dots, 2 \log m\}$ and any two jobs $j, j' \in \mathcal{S}(J_k^{**})$, Lemma 30 guarantees that precedence and communication constraints are satisfied. Furthermore, Lemma 30 also ensures that there cannot be jobs j, j' such that $j \in \mathcal{S}(J_k^{**})$, $j' \in \mathcal{S}(J_{k'}^{**})$ and $j' \prec j$ and $k' > k$. Finally note that every length- $2c$ interval can be split into 2 length- c intervals. The claim follows. \square

3.3.3 The Complete Algorithm for $P_\infty \mid \text{prec}, p_j = 1, c\text{-intervals} \mid C_{\max}$

Now we have all the pieces to put the rounding algorithm together and prove its correctness. We partition the jobs into batches, where each batch consists of subset of jobs that have C_j very close to each other in the LP solution. The complete algorithm is given below.

THE COMPLETE ALGORITHM

- (1) Solve the LP and let $(x, y, C) \in Q(r)$ with $r \geq 5$.
- (2) For $\delta = \frac{1}{64 \log(4c)}$ and $k \in \{0, 1, 2, \dots, \frac{S-1}{\delta}\}$, define

$$J_k = \{j \in J : k \cdot \delta \leq C_j < (k+1) \cdot \delta\}$$
- (3) FOR $k = 0$ TO $\frac{S-1}{\delta}$ DO
 - (4) Schedule the jobs in J_k using the algorithm in Subsection 3.3.2.

Now we finish the analysis of the rounding algorithm.

Proof of Theorem 25. Let us quickly verify that the schedule constructed by our algorithm is feasible. For jobs $j_1 \prec j_2$ with $j_1 \in J_{k_1}$ and $j_2 \in J_{k_2}$, the LP implies that $C_{j_1} \leq C_{j_2}$ and so $k_1 \leq k_2$. If $k_1 < k_2$, then j_1 will be scheduled in an earlier interval than j_2 . If $k_1 = k_2 = k$, then Lemma 34 guarantees that precedence constraints are satisfied.

It remains to bound the makespan of our algorithm. Lemma 34 guarantees that for $k \in \{0, 1, 2 \dots \frac{S-1}{\delta}\}$, the jobs in J_k are scheduled using at most $O(\log m) + \frac{|J_k|}{cm}$ many intervals. Then the total number of intervals required by the algorithm is bounded by

$$\frac{S}{\delta} \cdot O(\log m) + \sum_{k=0}^{\frac{S-1}{\delta}} \frac{|J_k|}{cm} = O(\log m \cdot \log c) \cdot S + \frac{|J|}{cm} \leq O(\log m \cdot \log c) \cdot S.$$

Here we use that $|J| \leq S \cdot cm$ is implied by the constraints defining K . □

Remark 1. We note that it is possible to reverse-engineer our solution and write a more compact LP for the problem, enforcing only the necessary constraints such as those given by Lemmas 27 and 28. We present this compact LP and prove an integrality gap of $\Omega(\sqrt{\log n})$ for it in Section 3.5. While the compact LP is simpler and could be solved more efficiently, we feel that the Sherali-Adams hierarchy gives a more principled and intuitive way to tackle the problem and explain how the LP arises, and hence we choose to present it that way.

3.4 Reductions

We now justify our earlier claim: the special case $P_\infty \mid \text{prec}, p_j = 1, c\text{--intervals} \mid C_{\max}$ indeed captures the full computational difficulty of the more general problem $P \mid \text{prec}, c \mid C_{\max}$. The main result for this section will be the following reduction:

Theorem 35. *Suppose there is a polynomial time algorithm that takes a solution for the LP $Q(r)$ with parameters $m, c, S \in \mathbb{N}$ and $r \geq 5$ and transforms it into a schedule for $P_\infty \mid \text{prec}, p_j = 1, c\text{--intervals} \mid C_{\max}$ using at most $\alpha \cdot S$ intervals. Then there is a polynomial time $O(\alpha)$ -approximation for $P \mid \text{prec}, c \mid C_{\max}$.*

For the reduction we will make use of the very well known list scheduling algorithm by

Graham [Gra66] that can be easily extended to the setting with communication delays. Here the notation $\sigma(j) = ([t, t + p_j], i)$ means that the job j is processed in the time interval $[t, t + p_j]$ on machine $i \in [m]$.

GRAHAM'S LIST SCHEDULING


- (1) Set $\sigma(j) := \emptyset$ for all $j \in J$
- (2) FOR $t = 0$ TO ∞ DO FOR $i = 1$ TO m DO
 - (3) Select any job $j \in J$ with $\sigma(j) = \emptyset$ where every $j' \prec j$ satisfies the following:
 - If j' is scheduled on machine i then j' is finished at time $\leq t$
 - If j' is schedule on machine $i' \neq i$ then j' finished at time $\leq t - c$
 - (4) Set $\sigma(j) := ([t, t + p_j], i)$ (if there was such a job)

For example, for the problem $P \mid \text{prec} \mid C_{\max}$, Graham's algorithm gives a 2-approximation. The analysis works by proving that there is a chain of jobs covering all time units where not all machines are busy. Graham's algorithm does *not* give a constant factor approximation for our problem with communication delays, but it will still be useful for our reduction.

Recall that a set of jobs $\{j_1, \dots, j_\ell\} \subseteq J$ with $j_\ell \prec j_{\ell-1} \prec \dots \prec j_1$ is called a *chain*. We denote $\mathcal{Q}(J)$ as the set of all chains in J w.r.t. precedence order \prec .

Lemma 36. *Graham's list scheduling on an instance of $P \mid \text{prec}, c \mid C_{\max}$ results in a schedule with makespan at most $\frac{1}{m} \sum_{j \in J} p_j + \max_{Q \in \mathcal{Q}(J)} \{\sum_{j \in Q} p_j + c \cdot (|Q| - 1)\}$.*

Proof. We will show how to construct the chain Q that makes the inequality hold. Let j_1 be the job which finishes last in the schedule produced by Graham's algorithm and let t_{j_1} be its start time. Let j_2 be the predecessor of j_1 that finishes last. More generally in step i , we denote j_{i+1} as the predecessor of j_i that finishes last. The construction finishes with a job j_ℓ without predecessors.

Figure 3.3: Analysis of Graham's algorithm with communication delay c .

Now let Q be the chain of jobs $j_\ell \prec j_{\ell-1} \prec \dots \prec j_1$. The crucial observation is that for any $i \in \{1, \dots, \ell-1\}$, either all machines are busy in the time interval $[t_{j_{i+1}} + p_{j_{i+1}} + c, t_{j_i})$ or this interval is empty. The reason is that Graham's algorithm does not leave unnecessary idle time and would have otherwise processed j_i earlier. It is also true that all m machines are busy in the time interval $[0, t_{j_\ell})$. The total amount of work processed in these busy time intervals is

$$L := m \cdot \left(t_{j_\ell} + \sum_{i=1}^{\ell-1} \max\{t_{j_i} - (c + p_{j_{i+1}} + t_{j_{i+1}}), 0\} \right) \leq \sum_{j \in J} p_j - \sum_{i=1}^{\ell} p_{j_i}.$$

Then any time between 0 and the makespan falls into at least one of the following categories:

- (a) the busy time periods described above, (b) the times that a job of the chain Q is processed,
- (c) the interval of length c following a job in the chain Q . Thus, we see that the makespan from Graham's list scheduling is at most

$$t_{j_1} + p_{t_{j_1}} \leq \frac{L}{m} + \sum_{j \in Q} p_j + c \cdot (|Q| - 1) \leq \frac{1}{m} \sum_{j \in J} p_j + \left(1 - \frac{1}{m}\right) \sum_{j \in Q} p_j + c \cdot (|Q| - 1).$$

□

It will also be helpful to note that the case of very small optimum makespan can be well

approximated:

Lemma 37. *Any instance for $P \mid \text{prec}, c \mid C_{\max}$ with optimum objective function value at most c admits a PTAS.*

Proof. Let J be the jobs in the instance and let $\text{OPT}_m \leq c$ be the optimum value. Consider the undirected graph $G = (J, E)$ with $\{j_1, j_2\} \in E \Leftrightarrow ((j_1 \prec j_2) \text{ or } (j_2 \prec j_1))$. Let $J = J_1 \dot{\cup} \dots \dot{\cup} J_N$ be the partition of jobs into connected components w.r.t. graph G . We abbreviate $p(J') := \sum_{j \in J'} p_j$. The assumption guarantees that the optimum solution cannot afford to pay the communication delay and hence there is a length- c schedule that assigns all jobs of the same connected component to the same machine. If we think of a connected component J_ℓ as an “item” of size $p(J_\ell)$, then for any fixed $\varepsilon > 0$ we can use a PTAS for $P \parallel C_{\max}$ (i.e. makespan minimization without precedence constraints) to find a partition of “items” as $[N] = I_1 \dot{\cup} \dots \dot{\cup} I_m$ with $\sum_{\ell \in I_i} p(J_\ell) \leq (1 + \varepsilon) \cdot \text{OPT}_m$ in polynomial time [HS87]. Arranging the jobs $\bigcup_{\ell \in I_i} J_\ell$ on machine i in any topological order finishes the argument. \square

Additionally, it is a standard argument to convert an instance with arbitrary p_j to an instance where all $p_j \leq n/\varepsilon$, while only losing a factor of $(1 + 2\varepsilon)$ in the approximation. For $p_{\max} := \max_j p_j$, we simply scale the job lengths and communication delay down by a factor of $\frac{n}{\varepsilon p_{\max}}$ then round them to the nearest larger integer. This results in at most a 2ε fraction of the optimal makespan being rounded up and all job sizes are integral and at most n/ε .

Now we can show the main reduction:

Proof of Theorem 35. Consider an instance of $P \mid \text{prec}, c \mid C_{\max}$ with $p_j, c \in \mathbb{N}$. Let J denote its job set with precedence constraints, and $\text{OPT}_m(J)$ denote its optimal value where m is the number of available machines. By the previous argument, we may assume that $p_j \leq 2n$ for all $j \in J$. Moreover, by Lemma 37 we only need to focus on the case where $\text{OPT}_m(J) > c$. We may guess




Figure 3.4: Splitting jobs into chains of unit-length jobs.

the optimum value of $\text{OPT}_m(J)$ as $\text{OPT}_m(J) \in \{1, \dots, 2n^2\}$.

Let J' denote the job set obtained by splitting each job $j \in J$ into a chain of p_j unit sub-jobs $j^{(1)} \prec \dots \prec j^{(p_j)}$. Moreover, precedence constraints in J are preserved in J' as we set all predecessors of j to be predecessors of $j^{(1)}$ and all successors of j to be successors of $j^{(p_j)}$, see Figure 3.4. We note that $\text{OPT}_m(J') \leq \text{OPT}_m(J)$ as splitting does not increase the value of the optimum. Let $\mathcal{S}_m(J')$ be a schedule achieving the value of $\text{OPT}_m(J')$. Next, observe that $\mathcal{S}_m(J')$ implies an integral solution for $Q(r)$ with parameters m, c, S where $S := \lceil \text{OPT}_m(J)/c \rceil$ and $r := 5$. In particular here we use that if jobs $j_1 \prec j_2$ are scheduled on different machines by $\mathcal{S}_m(J')$, then their starting times differ by at least $c + 1$ and hence they are assigned to different length- c intervals.

Now we execute the assumed α -approximate rounding algorithm and obtain a schedule $\mathcal{S}_{\infty, \text{int}}(J')$ that uses $T \leq \alpha S$ many intervals. We will use this solution $\mathcal{S}_{\infty, \text{int}}(J')$ to construct a schedule $\mathcal{S}_{\infty}(J)$ for $\text{P}\infty \mid \text{prec}, c \mid C_{\max}$ with job set J by running split sub-jobs consecutively on the same processor. This will use $4T$ time intervals in total. Recall that I_s denotes the time interval $[sc, (s+1)c)$. The rescheduling process is as follows:

For a fixed job $j \in J$, let I_{s_1} be the time interval where $j^{(1)}$ is scheduled in $\mathcal{S}_{\infty, \text{int}}(J')$. Then all other sub-jobs of j should be either scheduled in I_{s_1} or the time intervals after I_{s_1} .

- **Case 1: Some sub-job of j is not scheduled in I_{s_1} .**




Figure 3.5: Transformation of the schedule $\mathcal{S}_{\infty, \text{int}}(J')$ (top) to $\mathcal{S}_{\infty}(J)$ (bottom), where $\mathcal{S}_{\infty}(J)$ is compressed by a factor of 4. Here a “private” machine for a job j means the machine never processes any job other than j .

Schedule job j at the beginning of time interval I_{4s_1+1} on a new machine. If j is a short job, then it will finish running by the end of the interval. Otherwise j is a long job. Let I_{s_2} be the last time interval where a sub-job of j is scheduled in $\mathcal{S}_{\infty, \text{int}}(J')$. Then, the length satisfies $p_j \leq c \cdot (s_2 - s_1 + 1)$, which implies that the job finishes by time $c \cdot (4s_1 + 1) + p_j \leq c \cdot (4s_2 - 1)$.

- **Case 2: All sub-jobs of j are scheduled in I_{s_1} .**

Simply schedule job j during time interval I_{4s_1} on the same machine as in $\mathcal{S}_{\infty, \text{int}}(J')$.

See Figure 3.5 for a visualization. Then $\mathcal{S}_{\infty}(J)$ is a valid schedule for $P_{\infty} \mid \text{prec}, c \mid C_{\max}$, with makespan $\leq 4c \cdot T$. Moreover, $\mathcal{S}_{\infty}(J)$ satisfies the following:

- (a) A short job is fully contained in some interval I_s .
- (b) A long job’s start time is at the beginning of some interval I_s .

For $\mathcal{S}_{\infty}(J)$, define a new job set H . Every long job j becomes an element of H with its original running time p_j . Meanwhile, every set of short jobs that are assigned to the same machine in one

time interval becomes an element of H , with running time equal to the sum of running times of the short jobs merged. To summarize, a new job $h \in H$ corresponds to a set $h \subseteq J$ and $p_h = \sum_{j \in h} p_j$.

We define the partial order $\tilde{\prec}$ on H with $h_1 \tilde{\prec} h_2$ if and only if there are $j_1 \in h_1$ and $j_2 \in h_2$ with $j_1 \prec j_2$. One can check that this partial order is well defined. Moreover, by the fact that jobs assigned to the same interval but different machines do not have precedence constraints, the length of the longest chain in $(H, \tilde{\prec})$ in terms of the number of elements is bounded by the number of intervals that are used, which is at most $4T$.

Now run Graham's list scheduling on the new job set H with order $\tilde{\prec}$ and m machines. By Lemma 36, the makespan of the list scheduling is bounded by $\frac{1}{m} \sum_{h \in H} p_h + \max_{Q \in \mathcal{Q}(H)} \{ \sum_{h \in Q} p_h + c \cdot |Q| \}$. As the total sum of the processing times does not change from J to H , we see that $\frac{1}{m} \sum_{h \in H} p_h \leq \text{OPT}_m(J)$. Moreover, for any chain $Q \in \mathcal{Q}(H)$, $\sum_{h \in Q} p_h$ is no greater than the makespan of $\mathcal{S}_\infty(J)$, which is $4cT$. Finally, as argued earlier, the chain has $|Q| \leq 4T$ elements. Above all,

$$\begin{aligned}
\frac{1}{m} \sum_{h \in H} p_h + \max_{Q \in \mathcal{Q}(H)} \left\{ \sum_{h \in Q} p_h + c \cdot |Q| \right\} &\leq \text{OPT}_m(J) + 4cT + 4cT \\
&\leq \text{OPT}_m(J) + 8\alpha \cdot cS \\
&\leq \text{OPT}_m(J) + 16\alpha \cdot \text{OPT}_m(J) \\
&= O(\alpha) \cdot \text{OPT}_m(J).
\end{aligned}$$

□

3.5 Integrality Gap for the Compact LP

We consider the following layered expander graph from Maiti et al. [MRS⁺20]. Let V_1, \dots, V_L be sets of n vertices and for each $j \in V_k$, for $k \in [L - 1]$, let j sample d successors from V_{k+1} uniformly at random without replacement. We choose $n = cm/L$ and $c = d^L$, and set $L = \epsilon_1 \sqrt{\log n}$, $d = 2^{\epsilon_2 \sqrt{\log n}}$, and $m = c$. We will fix ϵ_1 and ϵ_2 as in Maiti et al.; the constants have several specifications, and we will not enumerate all of them, except that $\epsilon_2 \geq 5\epsilon_1$. Call this instance \mathcal{I} . Note that we have purposefully defined \mathcal{I} so that its partial order is not closed under transitivity, as it will be notationally convenient later to consider the successors and predecessors of a node $j \in V_k$ as the nodes j' in $\Gamma^+(j) \cap V_{k+1}$ and $\Gamma^-(j) \cap V_{k-1}$, respectively.

Our compact LP is the following:

$$\begin{aligned}
y_{j_1, j_1} &= 1 & \forall j_1 \in J \\
y_{j_1, j_2} &= y_{j_2, j_1} & \forall j_1, j_2 \in J \\
1 - y_{j_1, j_2} &\leq 1 - y_{j_1, j_3} + 1 - y_{j_2, j_3} & \forall j_1, j_2, j_3 \in J \\
\sum_{j \in J} y_{j_1, j} &\leq c & \forall j_1 \in J \\
C_{j_2} &\geq C_{j_1} + (1 - y_{j_1, j_2}) & \forall j_1 \prec j_2 \\
C_j &\geq 0 & \forall j \in J \\
0 \leq y_{j, j'} &\leq 1 & \forall j, j' \in J
\end{aligned}$$

Note the above completely does away with the x variables from LP. We will prove the following when parameters for the layered expander graph are specified as above:

Lemma 38. *The compact LP has value at most 1 (c -interval) on instance \mathcal{I} .*

Maiti et al. prove that the number of c -intervals required for an integral solution to schedule \mathcal{I} is $\Omega(L)$. Combining this with Lemma 38, we have the following theorem:

Theorem 39. *The compact LP has integrality gap at least $\Omega(\sqrt{\log n})$.*

Before we begin the proof of Lemma 38 we will define some new terms. First let $\phi(j) : J \rightarrow [L]$ be the map sending each job to its corresponding layer; so $j \in V_k$ if and only if $\phi(j) = k$. Next we let $r(j, j') = \min\{L, s(j, j')\}$, where $s(j, j')$ is the distance between j and j' in the undirected graph, i.e. the number of edges in the shortest path between j and j' , when edges of the DAG are made undirected. Note that for $j \in V_k$ with $j' \in V_k \cap \Gamma^+(j)$ and $j'' \in V_{k+2} \cap \Gamma^+(j')$, we have that $r(j, j') = r(j', j'') = 1 - 1/L$ and $r(j, j'') = 1 - 2/L$. The function r is clearly symmetric. We will show that r satisfies the triangle inequality, since normally it is not the case that a minimum of two metrics is still a metric.

Lemma 40. *Let $s(j, j')$ be the distance between j and j' in the undirected graph. Then $r(j, j') = \min\{L, s(j, j')\}$ satisfies the triangle inequality.*

Proof. Fix jobs $j, j', j'' \in J$. Suppose first that $r(j, j') = s(j, j')$. Since $s(j, j') \leq s(j, j'') + s(j'', j')$ and $s(j, j') \leq L$, it follows that $r(j, j') \leq r(j, j'') + r(j'', j')$. On the other hand if $r(j, j') = L$, the triangle inequality would only be violated if j'' was such that $L > s(j, j'') + s(j'', j')$. However, this inequality implies there is a path from j to j'' through j' that has $s(j, j'') + s(j'', j') < L$ edges, which contradicts the definition of $r(j, j')$. \square

Proof of Lemma 38. The set of jobs in instance \mathcal{I} is partitioned into layers V_1, \dots, V_L . For every job j , let j 's completion time be $C_j = \phi(j)/L$. For jobs j, j' , let $y_{j,j'} = 1 - r(j, j')/L$.

Now we check the inequalities of the LP. We begin by seeing that these y values form a semimetric $d(j, j') = 1 - y_{j,j'}$. For $j \in J$, $y_{j,j} = 1 - 0/L = 1$ so that $d(j, j) = 0$, and d is

symmetric since $r(j, j')$ is too. To see that the triangle inequality holds, fix j, j', j'' and observe that by Lemma 40.

$$1 - y_{j,j'} = r(j, j')/L \leq r(j, j'')/L + r(j'', j')/L = 1 - y_{j,j''} + 1 - y_{j'',j'}.$$

Further, we see that for j, j' with $j \prec j'$, $C_j = \phi(j)/L$, $C_{j'} = \phi(j')/L$, and $y_{j,j'} = 1 - (\phi(j') - \phi(j))/L$, so that $C_{j'} - C_j \geq 1 - y_{j,j'}$.

The last inequality to check is that for all $j \in J$, $\sum_{j'} y_{j,j'} \leq c$. Fix a job j , and recall a job j' has $y_{j,j'} = 1 - \ell/L$ exactly when $r(j, j') = \ell$. Every job, except those in V_L , has d successors, and with probability at least $1 - n^{-2}$, every job, except those in V_1 , has at most $d + 10\sqrt{d \log n}$ predecessors. Therefore, at most $2^\ell (d + 10\sqrt{d \log n})^\ell$ jobs j' have $r(j, j') = \ell$. It follows that

$$\begin{aligned} \sum_{j'} y_{j,j'} &\leq \sum_{\ell=1}^{L-1} 2^\ell (d + 10\sqrt{d \log n})^\ell (1 - \ell/L) \\ &= \frac{2(d + 10\sqrt{d \log n})((2(d + 10\sqrt{d \log n}))^L - 2(d + 10\sqrt{d \log n})L + L - 1)}{(2(d + 10\sqrt{d \log n}) - 1)^2 L} \\ &\leq \frac{2^{L+1}(d + 10\sqrt{d \log n})^{L+1}}{(d + 10\sqrt{d \log n})^2} \leq 2^{L+1}(d + 10\sqrt{d \log n})^{L-1} \leq 2^{L+1}(2d)^{L-1} \leq d^L = c, \end{aligned}$$

where the second to last inequality follows from the fact that $d \geq 10\sqrt{\log n}$, and the last inequality follows from the choice of d and L and the fact that $\epsilon_2 \geq 2\epsilon_1$.

□

3.6 Minimizing Weighted Sum of Completion Times

To illustrate the generality of our framework we show that it can be extended to handle different objective functions, in particular we can minimize the *weighted sum of completion times* of the

jobs. Here we restrict to the simplest case where jobs have unit length and an unbounded number of machines are available. In the 3-field notation, this problem is denoted by $P_\infty \mid \text{prec}, p_j = 1, c \mid \sum_j w_j C_j$. The input for this problem is the same as for the makespan minimization problem except that each job j now has a weight $w_j \geq 0$.

The goal is to minimize the objective function $\sum_j w_j C_j$, where C_j is the completion time of j , which is defined as the time slot in which job j is scheduled.

Note that the LP $Q(r)$ has variables C_j that denote the index of the length- c interval where j is being scheduled. A natural approach would be to interpret $c \cdot C_j$ as the completion time of job j and minimize $\sum_{j \in J} w_j \cdot c \cdot C_j$ over $Q(r)$. Then the rounding algorithm from Section 3.3 will indeed schedule each job j so that the completion time is at most $(O(\log c \cdot \log n) \cdot C_j + \Theta(\log n)) \cdot c$. We can observe that if a $O(\log c \cdot \log n)$ approximation is the goal, then this argument suffices for all jobs j where the LP solution has $C_j \geq \Omega(\frac{1}{\log c})$ – but it fails for jobs with $0 \leq C_j \ll 1$.

3.6.1 The linear program

In order to address this case, we first start with a more general LP relaxation compared to the makespan result which tracks the actual time slot where the jobs are processed, rather than just the interval. Again, we use the parameter $m \in \mathbb{N}$ to denote the number of machines that we allow the LP to use (one can set $m := n$) and the parameter $S \in \mathbb{N}$ to denote the number of intervals that we allow for the time horizon. We abbreviate $T := S \cdot c$ as the number of time slots. Note that $T \leq nc$ always suffices for any non-idling schedule. We index time slots as $[T] := \{1, \dots, T\}$ and consider an interval as a discrete set of slots $I_s := \{cs + 1, \dots, c(s + 1)\}$ where $s \in \{0, \dots, S - 1\}$.

Recall that in the makespan result $x_{j,i,s}$ variables indicated if job j got scheduled on machine i in the interval s . Here, we introduce additional variables of the form $z_{j,i,t}$ which indicate if job j is scheduled on machine i at time $t \in [T]$. The variables $x_{j,i,s}$ are fully determined by summing

over appropriate variables $z_{j,i,t}$, but we retain them for notational convenience. Further, similar to our makespan result, we impose an interval structure on the optimal solution and lose an $O(1)$ factor in the approximation ratio.

Let \tilde{K} be the set of fractional solutions to the following LP.

$$\begin{aligned}
\sum_{i \in [m]} \sum_{t \in [T]} z_{j,i,t} &= 1 \quad \forall j \in J \\
\sum_{j \in J} z_{j,i,t} &\leq 1 \quad \forall i \in [m] \quad \forall t \in [T] \\
\sum_{t' < t} \sum_{i \in [m]} z_{j_1,i,t'} &\geq \sum_{t' \leq t} \sum_{i \in [m]} z_{j_2,i,t'} \quad \forall j_1 \prec j_2 \quad \forall t \in [T] \\
\sum_{t \in I_s} z_{j,i,t} &= x_{j,i,s} \quad \forall j \in J \quad \forall s \in \{0, \dots, S-1\} \\
0 \leq z_{j,i,t} &\leq 1 \quad \forall j \in J, i \in [m], t \in [T]
\end{aligned}$$

Similar to the makespan LP, let $\tilde{Q}(r)$ be the set of feasible solutions (x, y, z, C) to the following LP:

$$\begin{aligned}
\text{Minimize} \quad & \sum_{j \in J} w_j \cdot C_j \\
y_{j_1, j_2} &= \sum_{s \in \{0, \dots, S-1\}} \sum_{i \in [m]} x_{(j_1, i, s), (j_2, i, s)} \\
C_{j_2} &\geq C_{j_1} + (1 - y_{j_1, j_2}) \cdot c \quad \forall j_1 \prec j_2 \\
C_j &= \sum_{i \in [m]} \sum_{t \in [T]} z_{j,i,t} \cdot t \quad \forall j \in J \\
(x, z) &\in SA_r(\tilde{K})
\end{aligned}$$

Note that the C_j variables in this LP relaxation denote the actual completion time of j unlike

their role in the makespan result, where they were used to indicate the interval in which j was scheduled. The main technical result for this section is the following:

Theorem 41. *Consider an instance for $P_\infty \mid \text{prec}, p_j = 1, c \mid \sum_j w_j C_j$ and a solution $(x, y, z, C) \in \tilde{Q}(r)$ with $r \geq 5$. Then there is a randomized algorithm with expected polynomial running time that finds a feasible schedule so that (i) $\mathbb{E}[C_j^A] \leq O(\log c \cdot \log n) \cdot C_j$ and (ii) $C_j^A \leq O(\log c \cdot \log n) \cdot C_j + O(\log n) \cdot c$ for all $j \in J$, where C_j^A is the completion time of job j .*

We briefly describe how Theorem 41 implies the approximation algorithm promised in Theorem 21:

Proof of Theorem 21. Note that strictly speaking $\tilde{Q}(r)$ is not actually a relaxation of $P_\infty \mid \text{prec}, p_j = 1, c \mid \sum_j w_j C_j$. However one can take an optimum integral schedule and insert c idle time slots every c time units and obtain a feasible solution for $\tilde{Q}(r)$. This increases the completion time of any job by at most a factor of 2. Then we set $r := 5$ and $m := n$ and solve the LP $\tilde{Q}(r)$ in time polynomial in n . Now consider the randomized schedule from Theorem 41 with completion times C_j^A . Then the expected objective function is $\mathbb{E}[\sum_{j \in J} w_j \cdot C_j^A] \leq O(\log n \cdot \log c) \cdot (\sum_{j \in J} w_j \cdot C_j)$. Markov's inequality guarantees that we can find in expected polynomial time a schedule that satisfies this inequality if we increase the right hand side by a constant factor. This completes the proof. \square

3.6.2 The Rounding Algorithm

Let (x, y, z, C) be an optimal solution to the LP relaxation $\tilde{Q}(r)$ with $r \geq 5$. It remains to show Theorem 41. We partition the jobs based on their fractional completion times. For $\delta = \frac{c}{64 \log(4c)}$ and $k \geq 0$, let $J_k := \{j \in J : k \cdot \delta \leq C_j < (k+1) \cdot \delta\}$.

We give a separate algorithm for scheduling jobs in J_0 within an interval of length at most

$O(\log n) \cdot c$. Now consider the remaining jobs. For $k = 1, 2, \dots$, we schedule jobs in the set J_k using the algorithm from Section 3.3.3,

inserting c empty time slots between the schedule of jobs in the set J_k and J_{k+1} . Let C_j^A denote the completion time of job j in our algorithm.

Lemma 42. *For $k \geq 1$, consider a job $j \in J_k$. Then deterministically $C_j^A \leq O(\log n \cdot \log c) \cdot C_j$.*

Proof. The claim follows from repeating the arguments in Lemma 34, so we only give a sketch here. Fix k and consider scheduling the jobs in the set J_k using the procedure described in Lemma 34, where we repeat the CKR clustering algorithm for $k = \{1, 2, \dots, 2 \log n\}$ iterations. Then the expected number of jobs that did not get scheduled in the first $2 \log n$ iterations is at most $\frac{|J_k|}{n^2} < 1$. Therefore, in expected polynomial time we can find a schedule such that $C_j^A \in [2 \log n \cdot O(c) \cdot k, 2 \log n \cdot O(c) \cdot (k + 1)]$. From the definition of set J_k , the fractional completion time C_j of every job j in J_k is at least $k \cdot \frac{c}{64 \log(4c)}$ in the LP solution. This completes the proof. \square

The only new ingredient for the completion time result is scheduling the jobs in the set J_0 . For $j \in J_0$, let t_j^* denote the earliest time instant t at which the job is scheduled to a fraction of at least $1 - \varepsilon$ in the LP solution. Here $0 < \varepsilon < 1$ is a small constant that we determine later. In scheduling theory this time is also called α -point with $\alpha = 1 - \varepsilon$. Formally

$$t_j^* := \min \left\{ t' \in [T] : \sum_{i=1}^m \sum_{t=1}^{t'} z_{j,i,t} \geq 1 - \varepsilon \right\} \quad (3.1)$$

We use the same semimetric $d(j_1, j_2) := 1 - y_{j_1, j_2}$ as in Section 3.3 and schedule jobs in J_0 using the following procedure.

SCHEDULE FOR J_0

- (1) Run a CKR clustering on the semimetric space (J_0, d) with parameter $\Delta := \frac{1}{12}$ and let V_1, \dots, V_k be the clusters.
 - (2) Let $V'_\ell := \{j \in V_\ell \mid (\Gamma^-(j) \cap J_0) \subseteq V_\ell\}$ for $\ell = 1, \dots, k$.
 - (3) For all $\ell = 1, \dots, k$ assign jobs in V'_ℓ on a single machine and schedule them in the increasing order of t_j^* values breaking ties in an arbitrary manner.
 - (4) Insert a gap of c time slots.
 - (5) Let $J'_0 \subseteq J_0$ be the set of jobs that did not get scheduled in steps (1) - (3). Use Lemma 34 to schedule J'_0 .
-

Lemma 43. *For a job $j_1 \in J_0$, the probability that j_1 gets scheduled in step (5) of the algorithm, i.e., $j_1 \in J'_0$, is at most $O(\log c) \cdot \frac{C_{j_1}}{c}$.*

Proof. The arguments are a slight refinement of Lemma 30. Consider the set $U := \{j_1\} \cup (\Gamma^-(j_1) \cap J_0)$ of j_1 and its ancestors. If $j_0 \prec j_1$, then $0 \leq C_{j_0} + c \cdot d(j_0, j_1) \leq C_{j_1}$ by the LP constraints and so $d(j_0, j_1) \leq \frac{C_{j_1}}{c}$. Then the diameter of U with respect to semimetric d is bounded by $2C_{j_1}$ and hence by Theorem 23.(b) the probability that U is separated is bounded by $\ln(2|N(U, \Delta/2)|) \cdot \frac{4\text{diam}(U)}{\Delta} \leq O(\log c) \cdot \frac{C_{j_1}}{c}$. \square

The next lemma follows from repeating the arguments in Lemma 42.

Lemma 44. *For a job $j \in J_0$, condition on the event that $j \in J'_0$. Then, $C_j^A|_{(j \in J'_0)} \leq O(\log n) \cdot c$.*

We can now prove that every cluster V'_ℓ can be scheduled on one machine so that the completion time of any job is at most twice the LP completion time.

Lemma 45. *For a small enough constant $\varepsilon > 0$ ($\varepsilon = \frac{1}{12}$ suffices) the following holds: Let $U \subseteq J$ be a set of jobs with $\text{diam}(U) \leq \varepsilon$ w.r.t. distance d . Define t_j^* as in Eq (3.1) and schedule the jobs*

in U in increasing order of t_j^* on one machine and denote the completion time of j by C_j^A . Then, in expectation $C_j^A \leq 2t_j^*$ for every $j \in U$.

Proof. Let us index the jobs in $U = \{j_1, \dots, j_{|U|}\}$ so that $t_{j_1}^* \leq \dots \leq t_{j_{|U|}}^*$. Suppose for the sake of contradiction that there is some job j_N with $C_{j_N}^A > 2t_{j_N}^*$. Abbreviate $U^* := \{j_1, \dots, j_N\}$ and $\theta^* := t_{j_N}^*$ so that $1 \leq t_j^* \leq \theta^*$ for $j \in U^*$. We observe that $|U^*| = \sum_{j \in U^*} p_j > 2\theta^*$. Then we have

$$(A) \sum_{j \in U^*} \sum_{i \in [m]} \sum_{t=1}^{\theta^*} z_{j,i,t} \geq (1-\varepsilon)|U^*|, \quad (B) \sum_{j \in U^*} y_{j,j_N} \geq (1-\varepsilon)|U^*|, \quad (C) \sum_{i \in [m]} \sum_{t=1}^{\theta^*} z_{j_N,i,t} \geq 1-\varepsilon$$

where (A) and (C) are by definition of t_j^* and (B) follows from $\text{diam}(U^*) \leq \text{diam}(U) \leq \varepsilon$. Intuitively, this means that we have $|U^*| > 2\theta^*$ many jobs that the LP schedules almost fully on slots $\{1, \dots, \theta^*\}$ while (B) means that the jobs are almost fully scheduled on the same machine.

As before, we will use the properties of the Sherali-Adams hierarchy to formally derive a contradiction. We know by Lemma 26³ that there is a distribution $(\tilde{x}, \tilde{z}, \tilde{y}) \sim \mathcal{D}(j_N)$ so that $\mathbb{E}[\tilde{x}_{j,i,s}] = x_{j,i,s}$, $\mathbb{E}[\tilde{z}_{j,i,t}] = z_{j,i,t}$ and $\mathbb{E}[\tilde{y}_{j_1,j_2}] = y_{j_1,j_2}$ while the variables involving job j_N are integral, i.e.

$\tilde{x}_{j_N,i,s}, \tilde{z}_{j_N,i,t} \in \{0, 1\}$. Consider the three events

$$(A') \sum_{j \in U^*} \sum_{i \in [m]} \sum_{t=1}^{\theta^*} \tilde{z}_{j,i,t} \geq (1-3\varepsilon)|U^*|, \quad (B') \sum_{j \in U^*} \tilde{y}_{j,j_N} \geq (1-3\varepsilon)|U^*|, \quad (C') \sum_{i \in [m]} \sum_{t=1}^{\theta^*} \tilde{z}_{j_N,i,t} = 1.$$

Then by Markov inequality $\Pr[A'] \geq \frac{2}{3}$, $\Pr[B'] \geq \frac{2}{3}$ and $\Pr[C'] \geq 1 - \varepsilon$, and so by the union bound $\Pr[A' \wedge B' \wedge C'] > 0$, assuming $\varepsilon < \frac{1}{3}$. Fix an outcome for $(\tilde{x}, \tilde{y}, \tilde{z})$ where the events A', B', C' happen and let $i_N \in [m], t_N \in \{1, \dots, \theta^*\}$ be the indices with $\tilde{z}_{j_N,i_N,t_N} = 1$. Then the interval

³Strictly speaking, Lemma 26 describes the SA properties for LP $Q(r)$, but an absolutely analogous statement holds for $\tilde{Q}(r)$.

index with $t_N \in I_{s_N}$ satisfies $\tilde{x}_{j_N, i_N, s_N} = 1$. Hence

$$\begin{aligned}
 (1 - 3\varepsilon)|U^*| &\stackrel{(B')}{\leq} \sum_{j \in U^*} \tilde{y}_{j, j_N} \stackrel{LP}{=} \sum_{j \in U^*} \sum_{i \in [m]} \sum_{s \in \{0, \dots, S-1\}} \tilde{x}_{(j, i, s), (j_N, i, s)} \stackrel{\tilde{x}_{j_N, i_N, s_N} = 1}{=} \sum_{j \in U^*} \tilde{x}_{j, i_N, s_N} \\
 &\stackrel{LP}{\leq} \underbrace{\sum_{j \in U^*} \sum_{t=1}^{\theta^*} \tilde{z}_{j, i_N, t}}_{\leq \theta^* \text{ by LP}} + \underbrace{\sum_{j \in U^*} \sum_{t > \theta^*} \tilde{z}_{j, i_N, t}}_{\leq 3\varepsilon|U^*| \text{ by } (A')} \leq \theta^* + 3\varepsilon|U^*|
 \end{aligned}$$

Rearranging gives $|U^*| \leq \frac{1}{1-6\varepsilon}\theta^*$, which is a contradiction for $\varepsilon \leq \frac{1}{12}$. \square

Lemma 46. *For a job $j \in J_0$, condition on the event that it got scheduled in the step (3) of the algorithm. Then, $C_j^A|_{(j \notin J'_0)} \leq O(C_j)$.*

Proof. Consider a job $j \in J_0 \setminus J'_0$. Then $j \in V'_\ell$ and by construction, the set V'_ℓ has diameter at most $\Delta = \frac{1}{12}$ w.r.t. d . Then Lemma 45 guarantees that the completion time is $C_j^A \leq 2t_j^*$ where we set $\varepsilon = \frac{1}{12}$. Finally note that an ε -fraction of j was finished at time t_j^* or later and hence $C_j = \sum_{i \in [m]} \sum_{t \in [T]} z_{j, i, t} \cdot t \geq \varepsilon \cdot t_j^*$. Putting everything together we obtain $C_j^A \leq \frac{2}{\varepsilon} C_j$. \square

We have everything to finish the proof of the completion time result.

Proof of Theorem 41. From Lemma 42, for $k \geq 1$ and $j \in J_k$, we have deterministically $C_j^A \leq O(\log n \cdot \log c) \cdot C_j$. Now consider a job $j \in J_0$. Then,

$$\begin{aligned}
 \mathbb{E}[C_j^A] &= \mathbb{E}[C_j^A | (j \notin J'_0)] \cdot \Pr[(j \notin J'_0)] + \mathbb{E}[C_j^A | (j \in J'_0)] \cdot \Pr[(j \in J'_0)] \\
 &\leq O(C_j) + O(\log c) \cdot \frac{C_j}{c} \cdot O(\log n) \cdot c \quad (\text{from Lemmas 43, 44, 46}) \\
 &\leq O(\log n \cdot \log c) \cdot O(C_j)
 \end{aligned}$$

Finally note that the completion time of a job $j \in J_0$ is always bounded by $C_j^A \leq O(\log n) \cdot c$.

The claim follows. \square

Conclusion

We gave a new framework for scheduling jobs with precedence constraints and communication delays based on metric space clustering. Our results take the first step towards resolving several important problems in this area.

One immediate open question is to understand whether our approach can yield a constant-factor approximation for $P \mid \text{prec}, c \mid C_{\max}$. A more challenging problem is to handle *non-uniform* communication delays in the problem $P \mid \text{prec}, c_{jk} \mid C_{\max}$, where c_{jk} is the communication delay between jobs $j \prec k$.

Chapter 4

Scheduling with Communication Delays on Related Machines

4.1 Introduction

We consider the problem of scheduling jobs with precedence and communication delay constraints on *related machines*. This classic model was first introduced by Rayward-Smith [RS87] and Papadimitriou and Yannakakis [PY90]. In this problem we are given a set J of n jobs, where each job j has a processing length $p_j \in \mathbb{Z}_+$ and a weight $w_j \in \mathbb{Z}_+$. The jobs need to be scheduled on m related machines, where machine $i \in [m]$ has speed $s_i \in \mathbb{Z}_+$. If a job j with processing length p_j is scheduled on the machine i , then it requires p_j/s_i time units to complete. In addition, we are given a *communication delay parameter* $c \in \mathbb{Z}_{\geq 0}$. The jobs have precedence and communication delay constraints, which are given by a partial order \prec . A constraint $j \prec j'$ encodes that job j' can only start after job j is completed. Moreover, if $j \prec j'$ and j, j' are scheduled on different machines, then j' can only start executing at least c time units after j had finished. On the other hand, if j and j' are scheduled on the same machine, then j' can start executing immediately after

j finishes. The goal is to schedule jobs *non-preemptively* so as to minimize a certain objective function. In a non-preemptive schedule, each job j needs to be assigned to a single machine i and executed during a contiguous time interval of length p_j/s_i .

We focus on two widely studied objective functions: (1) minimizing the weighted sum of completion times of jobs, and (2) minimizing the makespan. In the standard 3-field notation¹, these problems are denoted by $Q \mid \text{prec}, c \mid \sum w_j C_j$ and $Q \mid \text{prec}, c \mid C_{\max}$, respectively. In the presence of precedence and communication delay constraints, the weighted completion time objective is more general than the makespan, in the sense that one can use an approximation algorithm for the completion time objective to obtain a comparable result for the makespan. Our main motivation to study these problems is twofold:

- The problems of scheduling jobs with communication delays are some of the most notorious open questions in approximation algorithms and scheduling theory, which have resisted progress for a long time. For this reason, the well-known survey by Schuurman and Woeginger [SW99] and its recent update by Bansal [Ban17] list understanding the approximability of the problems in this model as one of the top-10 open questions in scheduling theory.
- These problems arise in many real-world applications, especially in the context of scheduling in data centers. A precedence constraint $j \prec j'$ typically implies that the input to j' depends on the output of j . Hence, if j and j' are scheduled on different machines, then the *communication delay* due to transferring this output to the other machine often becomes the bottleneck. The problem has received significant attention in applied data center scheduling literature; see [CZM⁺11, GFC⁺12, HCG12, SZA⁺18, ZZC⁺12, ZCB⁺15, LYZ⁺16] for more

¹We adopt the convention of [GLLK79, VLL90], where the respective fields denote: **(1) machine environment:** Q for related machines, P for identical machines, **(2) job properties:** prec for precedence constraints; c for communication delays of length c ; $p_j = 1$ for unit length case, **(3) objective:** C_{\max} for minimizing makespan, $\sum w_j C_j$ for minimizing weighted sum of completion times.

details. Another timely example is in the parallelization of Deep Neural Network training. When DNNs are trained on multiple clusters, the communication costs incurred for synchronizing the weight updates in fact dominate the overall execution time [NHP⁺19]. The resulting *device placement* problem [MPL⁺17, GCL18, JZA19, TPD⁺20] is indeed a variant of scheduling with communication delays.

Scheduling jobs with precedence and communication delays has been studied extensively over many years [RS87, PY90, MK97, HM01, THU92, HLV94, PY90, GKMP08]. Yet, very little was known in terms of the approximation algorithms for problem until the recent work by Maiti et al. [MRS⁺20] and us [DKR⁺20]. Previously, even for the *identical machines* case, where all machines have the same speed, the best algorithm for general c achieved an approximation factor of $2/3 \cdot (c + 1)$ [GKMP08], which only marginally improves on Graham’s list scheduling algorithm that obtains a $(c + 1)$ -approximation, while requiring the assumption that $p_j = 1$. Moreover, no results were known for the related machines case.

In a very recent work, we [DKR⁺20] designed an $O(\log m \log c)$ -approximation algorithm for minimizing makespan in the *identical machines* case. We also showed an $O(\log n \log c)$ -approximation algorithm for the weighted sum of completion times objective, in a special case where $p_j = 1$ and we have an *unlimited* number of machines. In a parallel and independent work, Maiti et al. [MRS⁺20] developed an $O(\log^2 n \log^2 m \log c / \log \log n)$ -approximation algorithm for the makespan objective function on *related machines*. Interestingly, these two results were obtained using rather different techniques.

Maiti et al. [MRS⁺20] obtain a polylogarithmic approximation algorithm for scheduling in the presence of precedence and communication delay constraints in the *job duplication* model. Here, a single job can be scheduled on *multiple machines*, which is known to effectively “hide” the communication delay constraints [PY90]. Quite surprisingly, Maiti et al. then show that one can

convert a schedule with duplication to a feasible schedule without duplication, where every job is processed on a single machine, while increasing the makespan by at most an $O(\log^2 n \log m)$ factor. To create a strict contrast between our techniques and that of Maiti et al., we do not consider the duplication setting or any reduction from it. Additionally, the LP of Maiti et al. does not give rise to a *semimetric* on the set of jobs, which is a crucial part of our scheduling algorithm.

The framework of Maiti et al. extends to the weighted sum of completion times objective, but it obtains an extra additive term in the approximation that may dominate the factor; we discuss this further in Section 4.1.2. Moreover, our previous results [DKR⁺20] also do not imply any approximation guarantees for the related machines case. Specifically, in the presence of communication delay constraints there are no known reductions between the sum of weighted completion time and makespan objective functions similar to [CS99].

Finally, in another recent work, Su et al. [SRVW20] studied the objective of minimizing makespan and sum of weighted completion times on related machines. They showed that a Generalized Earliest Time First (GETF) algorithm achieves a makespan of $O(\log m / \log \log m) \cdot OPT + C$, where C is the total communication delay in the longest chain in the input graph. They show a similar bound on the schedule produced by GETF for the weighted sum of completion times objective. However, both of these bounds do not give any multiplicative approximation guarantees, as the additive terms can be substantially higher than the optimal solution. The additive terms in above bounds are precisely the reason why the scheduling with communication delays problem is significantly more challenging than the case when $c = 0$.

4.1.1 Our Contributions

Let M denote the number of machines types or equivalently speed classes. By standard arguments we can assume that $M = \log(s_{\max}/s_{\min}) \leq O(\log m)$ while only losing a constant factor in the

approximation guarantee, where s_{\max} and s_{\min} denote the fastest and slowest speeds of machines in the input instance.

The main result of this paper is the following:

Theorem 47. *There is a randomized $O(M^2 \cdot \log^2 n)$ -approximation algorithm for $Q \mid \text{prec}, c \mid \sum_j w_j C_j$ with expected polynomial running time. When jobs have unit processing lengths, $Q \mid \text{prec}, c, p_j = 1 \mid \sum_j w_j C_j$, the approximation factor of the algorithm improves to $O(M \cdot \log^2 n)$.*

As $M \leq O(\log n)$, our result gives an $O(\log^4 n)$ -approximation algorithm to the general problem and an $O(\log^3 n)$ -approximation algorithm for the case with unit processing lengths. As a byproduct of our result we also obtain an improved approximation for the makespan objective function.

Theorem 48. *There is a randomized $O(M \cdot \log m \cdot \log n)$ -approximation algorithm for $Q \mid \text{prec}, c \mid C_{\max}$ with expected polynomial running time. When jobs have unit processing lengths, $Q \mid \text{prec}, c, p_j = 1 \mid C_{\max}$, the approximation factor of the algorithm improves to $O(\log m \cdot \log n)$.*

So our result gives an $O(\log^3 n)$ -approximation algorithm for the problem of minimizing makespan (and an $O(\log^2 n)$ -approximation algorithm for the case with unit processing lengths). This improves the $O(\log^5 n / \log \log n)$ -approximation algorithm due to Maiti et al. [MRS⁺20].

4.1.2 Technical Challenges

Absent the communication delay constraints, one can use an approximation for the makespan objective to obtain an approximation algorithm for the weighted sum of completion times objective, with some (negligible) loss in the approximation quality [HSSW97, QS02, Li17]. Namely, a standard approach is to first solve an LP for the weighted sum of completion times problem, and then

geometrically partition jobs according to completion times in the LP solution into buckets of the form $[2^t, 2^{t+1}]$. Next, use an α -approximation algorithm for makespan to schedule these jobs in an interval of length $O(\alpha \cdot 2^t)$. Finally, concatenate the schedules of jobs belonging to different partitions. This approach gives an $O(\alpha)$ -approximation for identical machines, and an extension of this idea to related machines is given by [CS99].

However, this approach fails in the presence of communication delay constraints, even for the identical machines case. In particular, the main technical difficulty arises when in the LP solution a large number of jobs have very small completion times, say $O(c/\log^2 n)$. We need to schedule these jobs so that no precedence and communication delay constraints are violated, while at the same time achieving completion times comparable to the LP. This requires very precise control over job completion times, which cannot be achieved by the existing algorithms for the makespan. The problem becomes further more complex if machines have different speeds and jobs have different weights. This is also the main technical hurdle if one tries to adapt the framework of Maiti et al. [MRS⁺20]. In fact, [MRS⁺20] gives a schedule whose cost can be as large as $O\left((\log^5 n / \log \log n) \cdot OPT + c \cdot \sum_j w_j\right)$, which only implies an approximation factor of $\max\{c, \log^5 n / \log \log n\}$.

4.1.3 Our Techniques and Algorithm Overview

We obtain our results by generalizing the LP hierarchy framework introduced by [DKR⁺20] to handle processors of varying speed and the more general objective of minimizing the weighted sum of completion times. In this section we outline our main techniques and explain our innovations compared to [DKR⁺20]. In particular, as alluded above, the weighted sum of completion times objective requires a finer control over the *time slots where jobs are scheduled* compared to minimizing the makespan. Moreover, allowing different machine types introduces an *assignment*

aspect to the problem, namely assigning jobs to a machine type, which is not present in the identical machines case. Tackling these two challenges simultaneously requires an extended LP as well as more structural insights about the solution produced by the Sherali-Adams hierarchy compared to [DKR⁺20].

To keep the notation simple we will use $\log n$ instead of the potentially smaller quantities $\log m$ and M . We begin our discussion with the problem $Q \mid \text{prec}, p_j = 1, c \mid \sum w_j C_j$, where jobs have unit lengths. We will see later that we can reduce the version with general processing times to the case $p_j = 1$ while losing a factor of $O(\log n)$. A helpful view that already emerged in preceeding work is to partition the time horizon into intervals of length c or larger. Then if one can assign each job j_1 so that any predecessor j_2 is either executed on the same machine or in a earlier interval, then inserting a waiting time of c after each interval will result in a valid schedule while completion times increase by at most a factor of 2.

We start by designing an LP with *time-indexed variables* $x_{j,i,t}$ that specify whether job j is to be scheduled in time slot t on machine i . To accommodate the different speeds, a machine in class k has s_k slots available per unit time interval. Similarly to [DKR⁺20], we take the *Sherali-Adams lift* with a constant number of rounds and use it to extract variables y_{j_1,j_2} , which tell us whether the jobs are scheduled on the same machine within an interval of length c , as well as variables C_j , which denote the *LP completion times*. Our main LP rounding result is that we can generate a schedule at random so that the completion time of every job is at most $O(\log^3 n) \cdot C_j$ in expectation.

Using the Sherali-Adams properties one can prove that the function $d : J \times J \rightarrow \mathbb{R}_{\geq 0}$ with $d(j_1, j_2) := 1 - y_{j_1,j_2}$ is a *semimetric*. One important ingredient of our algorithm is a clustering procedure due to Calinescu, Karloff and Rabani [CKR04] which, for such a semimetric space (J, d) and a parameter $\Delta > 0$, finds a random partition $J = V_1 \dot{\cup} \dots \dot{\cup} V_q$ into blocks of diameter at most Δ such that any set $U \subseteq J$ has a probability of at most $O(\log(n) \cdot \frac{\text{diam}(U)}{\Delta})$ of being separated.

We will use two different lines of arguments for scheduling jobs with very small LP completion time and for the remaining jobs.

Case I: Scheduling jobs with very small LP completion time. Consider the jobs $J_0 := \{j \in J \mid 0 \leq C_j \leq \Theta(\frac{c}{\log^2 n})\}$ whose LP completion time C_j is much smaller than the communication delay. Jobs with tiny C_j have to be scheduled in the first interval with sufficiently high probability, and moreover the position within the first interval has to be proportional to C_j as well. To achieve this, we run the CKR clustering with a rather small parameter $\Delta := \Theta(\frac{1}{\log n})$ and denote $J_0 = V_1 \dot{\cup} \dots \dot{\cup} V_q$ as the obtained partition. Clustering with such small diameter parameter allows us to prove structural insights about clusters that are new and were not required for the setting of [DKR⁺20]. If we consider a job $j \in J_0$, one can prove that all its predecessors $j' \prec j$ have a distance of $d(j, j') \leq \frac{C_j}{c}$. In particular, the probability that j gets separated from any of its ancestors is bounded by $O(\log^2(n) \cdot \frac{C_j}{c})$. Let us denote $V'_\ell \subseteq V_\ell$ as the jobs that are not separated from any ancestor and let $J'_0 := \bigcup_{\ell=1}^q V'_\ell$ be their union. Note that any set V'_ℓ could be processed on a single machine starting at time 0 without the danger of violating any precedence or communication delay constraints. Consequently, we will schedule the jobs in J'_0 right at the beginning of the schedule; after all of them are completed, we will schedule the jobs in $J_0 \setminus J'_0$.

- **Scheduling J'_0 .** This is the most delicate part of the algorithm, as for the jobs in J'_0 even the relative order of the jobs within the sets V'_ℓ has to be decided. Moreover, there is no obvious bijection between the V'_1, \dots, V'_q and the machines, and we do not have a uniform upper bound on the size of the sets V'_ℓ . But we can prove a novel structural lemma that for each V'_ℓ there is a machine type k such that $|V'_\ell| \leq O(s_k c)$ and the LP solution schedules *every* job $j \in V'_\ell$ to an extent of $\Omega(\frac{1}{\log n})$ on machines of class k . Then we assign V'_ℓ to a random machine in that speed class k . Now consider the situation of one machine i of class k and let $J'_0(i, k)$ be the union of jobs assigned to this machine. The order we choose for

sequencing the jobs in $J'_0(i, k)$ is the increasing order of α -points, which for us is the time when the LP has processed a $(1 - \Theta(\frac{1}{\log n}))$ -fraction of job j . Again using properties implied by the Sherali-Adams lift combined with a Chernoff bound we can prove that for every θ , the number of jobs in $J'_0(i, k)$ with α -point at most θ is at most $O(\log^2 n) \cdot s_k \cdot \theta$. We can then conclude that every job in J'_0 is indeed completed by time $O(\log^3 n) \cdot C_j$.

- **Scheduling** $J''_0 := J_0 \setminus J'_0$. As the probability for a job $j \in J_0$ to end up in this case is small enough, it suffices to schedule all the jobs J''_0 in a time horizon of $O(\log n) \cdot c$ without caring about their particular order. In fact, we can simply use the same procedure that we are about to describe for Case II.

Overall we can see that adding up the contributions from both cases, the expected completion time of a job $j \in J_0$ will be $O(\log^3(n) \cdot C_j) + O(\log^2 n \cdot \frac{C_j}{c}) \cdot O(\log n \cdot c) \leq O(\log^3(n) \cdot C_j)$ as required.

Case II: Scheduling jobs with lower-bounded LP completion time. The main result to handle this case is that any subset J_T of jobs with LP completion times $C_j \leq T$ for all $j \in J_T$ can be scheduled with makespan $O(\log^2 n) \cdot T + O(\log n) \cdot c$. The complete algorithm can then be easily obtained by running the argument for all $T \geq \Theta(\frac{c}{\log^2 n})$ that are powers of 2 and concatenating the obtained schedules. We can break the problem further down to scheduling jobs in a narrow band of the form $J^* := \{j \in J_T \mid C^* \leq C_j \leq C^* + \Theta(\frac{c}{\log n})\}$, i.e., jobs that have close LP completion times. Again we run the CKR clustering procedure, but this time with a larger parameter $\Delta := \frac{1}{4}$ (which will result in saving a $\log n$ factor in the approximation guarantee compared to choosing $\Delta = \Theta(\frac{1}{\log n})$). Again we denote the random partition by $J^* = V_1 \dot{\cup} \dots \dot{\cup} V_q$ and define $V'_\ell \subseteq V_\ell$ to be the jobs that were not separated from any ancestor in J^* . Unfortunately in this regime of $\Delta = \Theta(1)$, the mentioned structural claim from Case I does not hold anymore. However we can

prove a weaker result that for every block V'_ℓ there is a speed class k with $|V'_\ell| \leq O(s_k c)$ such that the LP solution schedules jobs in V'_ℓ on *average* to an extend of $\Omega(\frac{1}{\log n})$ on class- k machines. Since in Case II, we do not have to decide an order within the blocks V'_ℓ , this weaker property will be sufficient. Repeating the random clustering $\log n$ times will schedule all jobs in J^* . Finally we use a load argument to conclude the bound on the makespan of J_T .

Conclusion and reduction to general processing times. Overall this line of arguments gives an $O(\log^3 n)$ -approximation for $Q \mid \text{prec}, p_j = 1, c \mid \sum w_j C_j$. Moreover for the problem $Q \mid \text{prec}, p_j = 1, c \mid C_{\max}$ it suffices to consider instances with optimum value at least c and hence we can find an $O(\log^2 n)$ -approximation for that setting.

Finally let us outline how to obtain an approximation to $Q \mid \text{prec}, c \mid \sum w_j C_j$ while losing at most another $O(\log n)$ factor. Consider jobs J that now have arbitrary integer processing times p_j . We perform the natural reduction and split each job into a chain of p_j many unit-length jobs. Then we run the $O(\log^3 n)$ -approximation algorithm for $Q \mid \text{prec}, p_j = 1, c \mid \sum w_j C_j$. The obtained schedule can be interpreted as a schedule for the original length- p_j jobs that respects precedence constraints and communication delays, but uses preemption and migration. By standard arguments it suffices to schedule the jobs with completion time at most T so that the makespan is $O(\log n) \cdot T$. We use the information from the migratory schedule and create a new instance $(\tilde{J}, \tilde{\prec})$ where jobs that are scheduled within a length- c timeframe are merged and we have precedence constraints $j_1 \tilde{\prec} j_2$ whenever a job j_1 was finished before j_2 in the migratory schedule. We can prove that chains in the new partial order $\tilde{\prec}$ contain at most $O(\frac{T}{c})$ many jobs. This implies that a small modification of the SPEED-BASED LIST SCHEDULING by Chudak and Shmoys [CS99] can be used to schedule $(\tilde{J}, \tilde{\prec})$ with makespan $O(\log n) \cdot T$. In particular, the penalty for taking communication delays into account is bounded by $c \cdot (|C| - 1) \leq O(T)$, where $C \subseteq \tilde{J}$ is the chain that defines the bottleneck in the algorithm. That concludes the $O(\log^4 n)$ -approximation for $Q \mid \text{prec}, c \mid \sum w_j C_j$.

4.1.4 Outline

The rest of this paper is organized as follows. In Section 4.2 we give preliminaries on hierarchies, semimetric spaces, and a useful version of the Chernoff bound. In Section 4.3 we discuss our linear program and its properties. In Sections 4.4 and 4.5 we give our algorithm for the weighted sum of completion times objective in the special case of unit-size jobs. Then, in Section 4.6, we show how to reduce the general processing time case to the unit-size case. Together, Sections 4.4–4.6 constitute the proof of Theorem 47. Finally, in Section 4.7 we solve the makespan minimization version while saving one $O(\log n)$ factor, thus proving Theorem 48.

4.2 Preliminaries

4.2.1 The Sherali-Adams Hierarchy for LPs with Assignment Constraints

See Section 3.2 for most of the definitions and properties on the Sherali-Adams hierarchy. Recall that K is a polytope with assignment constraints and $SA_r(K)$ is the *Sherali-Adams* lift of K of rank r .

For $y \in SA_r(K)$, it will be convenient later to use the notation $\tilde{y} \sim \mathcal{D}_y(H)$ with $|H| \leq r$ for the distribution described in Theorem 22.(f). Often we will omit the vector y if the Sherali-Adams lift is clear from the context. If we consider indices $J \subseteq U_h$ for some $h \in [N]$ then we can interpret these as an event $\mathcal{E} = \{\exists i \in J : x_i = 1\}$ and the probability of this event is $\Pr_{\tilde{y} \sim \mathcal{D}_y(h)}[\mathcal{E} \text{ holds in } \tilde{y}] = \sum_{i \in J} y_i$. Assuming that the probability of this event is positive, we can obtain a new Sherali-Adams lift conditioned on that event to happen while losing at most one round in the rank of the solution. We summarize the properties that we require later:

Lemma 49. *Let $y \in SA_r(K)$ for some $r \geq 0$. Fix an index $h \in [N]$ and fix $J \subseteq U_h$ with*

$\sum_{i \in J} y_i > 0$. Define $\bar{y} \in \mathbb{R}^{\mathcal{P}_r([n])}$ with $\bar{y}_I := \frac{\sum_{i \in J} y_{I \cup \{i\}}}{\sum_{i \in J} y_i}$ for $I \in \mathcal{P}_r([n])$. Then $\bar{y} \in SA_{r-1}(K)$ and moreover $\bar{y}_I \leq \frac{y_I}{\sum_{i \in J} y_i}$.

Proof. Abbreviate $\gamma := \sum_{i \in J} y_i > 0$ and $J_+ := \{i \in J \mid y_i > 0\}$. For $i \in J_+$ we denote $y^{(i)} \in \mathcal{P}_r([n])$ as the vector with $y_I^{(i)} := \frac{y_{I \cup \{i\}}}{y_i}$. By Theorem 22.(a) we know that $y^{(i)} \in SA_{r-1}(K)$. Then y is a convex combination of the vectors $y^{(i)}$ as one can see from $\bar{y}_I = \sum_{i \in J_+} \frac{y_i}{\gamma} \cdot y_I^{(i)}$. By convexity of $SA_{r-1}(K)$, this implies that $\bar{y} \in SA_{r-1}(K)$. The moreover part follows from $\sum_{i \in J} y_{I \cup \{i\}} \leq \sum_{i \in U_h} y_{I \cup \{i\}} = y_I$, using Theorem 22.(e). \square

Suppose that $\bar{y} \in SA_{r-1}(K)$ is the vector obtained by conditioning on the event $\mathcal{E} = \{\exists i \in J : x_i = 1\}$ according to Lemma 49 with $J \subseteq U_h$. Then we will also abbreviate the distribution $\mathcal{D}_{\bar{y}}(h')$ more conveniently as the conditional distribution $\mathcal{D}_y(h' \mid \mathcal{E})$.

4.2.2 Semimetric Spaces

See Section 3.2.2 for the definition of a semimetric space, the statement of the CKR algorithm, and Theorem 23 on the analysis of the CKR algorithm.

4.2.3 Concentration

We will also make use of the following version of the Chernoff bound. See for example the textbook of Dubhashi and Panconesi [DP09].

Lemma 50. *There is a universal constant $C > 0$ such that the following holds. Let $X := X_1 + \dots + X_n$ be a sum of independent random variables with $0 \leq X_\ell \leq \alpha$ for all $\ell \in [n]$ and $\mathbb{E}[X] \leq \alpha$ for some $\alpha > 0$. Then for any $N \geq 4$ one has $\Pr[X > C \frac{\log N}{\log \log N} \cdot \alpha] \leq \frac{1}{N}$.*

4.3 The Linear Program and Its Properties

Let J be a set of $|J| = n$ jobs, each with a *processing time* $p_j = 1$ and a weight $w_j \geq 0$. Let $\{1, \dots, m\}$ be the indices of the available machines where we assume to have m_k machines of speed $s_k \in \mathbb{N}$. We let M be the number of different machine *types*, also referred to as speed classes. By the standard argument of geometrically grouping machines with speeds within a constant factor of each other [Li17, MRS⁺20], we can assume that $M \leq O(\log(\frac{s_{\max}}{s_{\min}})) \leq O(\log m)$ while we lose a constant factor in the approximation. The goal is to find a schedule such that jobs $j_1 \prec j_2$ are either scheduled on the same machine (with j_1 finishing before j_2 is started) or they are scheduled on different machines and j_2 starts at least c time units after j_1 is finished. By a slight abuse of notation we denote $[m_k]$ as the indices of the machines of speed s_k .

4.3.1 The Linear Program

It will be convenient to partition the time horizon into *intervals* I_0, I_1, \dots, I_{L-1} of c time units each. As machines have different speeds and hence can handle different numbers of jobs per interval, we abbreviate the discrete set of time slots that a speed- s_k machine has in I_ℓ as $I_{\ell,k} := \{\ell \cdot c + \frac{t}{s_k} : t = 1, \dots, s_k c\}$. We note that indeed $|I_{\ell,k}| = s_k c$. Moreover we set $I_{*,k} := \bigcup_{\ell=0}^{L-1} I_{\ell,k}$ as all time slots where a speed- s_k machine can schedule jobs.

We construct the LP in two steps. First consider the variables

$$x_{j,i,t} = \begin{cases} 1 & \text{if } j \text{ is scheduled on machine } i \text{ in time slot } t \in I_{*,k}, \\ 0 & \text{otherwise} \end{cases}$$

for all $j \in J$, $k \in [M]$, $i \in [m_k]$, and $t \in I_{*,k}$. Let K be the set of fractional solutions to the




Figure 4.1: Slots for different machine types.

following linear system:

$$\begin{aligned}
 \sum_{k \in [M], i \in [m_k]} \sum_{t \in I_{*,k}} x_{j,i,t} &= 1 \quad \forall j \in J \\
 \sum_{j \in J} x_{j,i,t} &\leq 1 \quad \forall k \in [M] \quad \forall i \in [m_k] \quad \forall t \in I_{*,k} \\
 0 \leq x_{j,i,t} &\leq 1 \quad \forall j \in J \quad \forall k \in [M] \quad \forall i \in [m_k] \quad \forall t \in I_{*,k}
 \end{aligned}$$

We note that the assignment constraints are the set of equations $\sum_{k \in [M], i \in [m_k]} \sum_{t \in I_{*,k}} x_{j,i,t} = 1$ for every $j \in J$.

Next, we will use a lift $x \in \text{SA}_r(K)$, which contains in particular the variables $x_{(j_1, i_1, t_1), (j_2, i_2, t_2)}$ that provide the probability for the event that j_1 is scheduled at time t_1 on machine i_1 and j_2 is scheduled at time t_2 on machine i_2 . We introduce more types of decision variables:

$$y_{j_1, j_2, k} = \begin{cases} 1 & j_1 \text{ and } j_2 \text{ are scheduled on the same machine of type } k \text{ in the same interval,} \\ 0 & \text{otherwise,} \end{cases}$$

$$y_{j_1, j_2} = \begin{cases} 1 & j_1 \text{ and } j_2 \text{ are scheduled on the same machine in the same interval,} \\ 0 & \text{otherwise,} \end{cases}$$

C_j = completion time of job j .

The LP relaxation is then as follows:

$$\begin{aligned}
\text{Minimize} \quad & \sum_{j \in J} w_j \cdot C_j \quad (\text{LP}) \\
y_{j_1, j_2, k} = & \sum_{\ell \in \{0, \dots, L-1\}} \sum_{i \in [m_k]} \sum_{t_1 \in I_{\ell, k}} \sum_{t_2 \in I_{\ell, k}} x_{(j_1, i, t_1), (j_2, i, t_2)} \quad \forall j_1, j_2 \in J \quad \forall k \in [M] \\
y_{j_1, j_2} = & \sum_k y_{j_1, j_2, k} \quad \forall j_1, j_2 \in J \\
C_{j_2} \geq & C_{j_1} + (1 - y_{j_1, j_2}) \cdot c \quad \forall j_1 \prec j_2 \\
C_j = & \sum_{k \in [M]} \sum_{i \in [m_k]} \sum_{t \in I_{*, k}} x_{j, i, t} \cdot t \quad \forall j \in J \\
x \in & SA_r(K)
\end{aligned}$$

Following the discussion in Section 4.2.1, we know that for every job j , there is a distribution that we denote as $(\tilde{x}, \tilde{y}) \sim \mathcal{D}(j^*)$ such that $\mathbb{E}[\tilde{x}_{j, i, t}] = x_{j, i, t}$ and $\mathbb{E}[\tilde{y}_{j_1, j_2}] = y_{j_1, j_2}$ with $\tilde{x} \in SA_{r-1}(K)$ where job j^* is integrally assigned and (\tilde{x}, \tilde{y}) satisfies the first two constraints in (LP). This is immediate for the x -part as $x \in SA_r(K)$, and follows for the y -variables as these linear in the x -variables. If \mathcal{E} is an event, then we write $(\tilde{x}, \tilde{y}) \sim \mathcal{D}(j^* \mid \mathcal{E})$ as the conditional distribution (conditioning on the event \mathcal{E} occurring), see again Section 4.2.1 for details.

4.3.2 Properties of the LP

We will now discuss some properties that are implied by the Sherali-Adams lift. The properties proved in this section, which we use crucially in our rounding algorithms, are the main technical contributions of this paper.

Lemma 51. *Let (x, y, C) be a solution to (LP) with $r \geq 5$. Then $d(j_1, j_2) := 1 - y_{j_1, j_2}$ is a*

semimetric.

This property was proven in [DKR⁺20] for a special case of $s_k = 1$ and is not hard to extend to our more general LP. From now on, the symbol d as well as the quantity $\text{diam}(\cdot)$ will always refer to this particular semimetric. For the special case of $s_k = 1$ for all k , one can also find a proof in [DKR⁺20] that any set $U \subseteq J$ with $\text{diam}(U) \leq \frac{1}{2}$ has size $|U| \leq 2c$. While this is obviously false for arbitrary speeds s_k , we can prove a similar claim:

Lemma 52. *For $k \in [M]$ and $j_1 \in J$ one has $\sum_{j_2 \in J} y_{j_1, j_2, k} \leq s_k \cdot c \cdot \sum_{i \in [m_k]} \sum_t x_{j_1, i, t} \leq s_k \cdot c$.*

Proof. The second inequality is trivial as $\sum_{i \in [m_k]} \sum_t x_{j_1, i, t} \leq 1$, so we only justify the first inequality. As we are proving a linear inequality, it suffices to show this for a fixed outcome $(\tilde{x}, \tilde{y}) \sim \mathcal{D}(j_1)$ where job j_1 is assigned integrally. If in (\tilde{x}, \tilde{y}) the job j_1 is not assigned to a machine in $[m_k]$, then both sides are 0. So suppose that j_1 is assigned to a machine $i_1 \in [m_k]$ and to interval ℓ_1 . Then indeed

$$\sum_{j_2 \in J} \tilde{y}_{j_1, j_2, k} = \sum_{j_2 \in J} \sum_{t \in I_{\ell_1, k}} \tilde{x}_{j_2, i_1, t} \leq s_k c = s_k c \underbrace{\sum_{i \in [m_k]} \sum_t \tilde{x}_{j_1, i, t}}_{=1}.$$

□

Lemmas 53 and 54 are key technical insights behind the algorithms for the related machines setting. They say that if one considers a set of jobs which are close to each other with respect to d , then there exists a type k^* such that we can schedule all the jobs in an interval of length $O(c)$ on a single machine of type k^* . Moreover, the LP also schedules a good fraction of these jobs on the same machine type. These lemmas are important in assigning jobs to machine types in our algorithms.

Lemma 53. *Let $U \subseteq J$ be a non-empty subset of jobs with $\text{diam}(U) \leq \frac{1}{4}$. Then there exists a $k^* \in [M]$ such that*

- (i) $|U| \leq O(1) \cdot s_{k^*} \cdot c$, and
- (ii) $\sum_{j \in U} \sum_{i \in [m_{k^*}]} \sum_t x_{j,i,t} \geq \Omega(\frac{1}{M}) \cdot |U|$.

Proof. Let us sort the speed classes $[M]$ so that $s_1 \geq \dots \geq s_M$ and abbreviate $z_{j,k} := \sum_{i \in [m_k]} \sum_t x_{j,i,t}$ as the fraction of job j scheduled on class k . Moreover let $\rho_k := \sum_{j \in U} z_{j,k}$ be the load of cluster U on class k . We fix an arbitrary job $j^* \in U$. Let $k_{\text{median}} \in [M]$ be the median speed class for job j^* , meaning that $\sum_{k \leq k_{\text{median}}} z_{j^*,k} \geq \frac{1}{2}$ and $\sum_{k \geq k_{\text{median}}} z_{j^*,k} \geq \frac{1}{2}$. We split the remaining proof into two separate claims.

Claim I. *One has $|U| \leq 4 \cdot s_{k_{\text{median}}} \cdot c$.*

Proof of Claim I. First we observe that for every $j \in U$ we have $\sum_{k \geq k_{\text{median}}} y_{j^*,j,k} \geq \frac{1}{2} - d(j^*, j) \geq \frac{1}{4}$. We use this to estimate

$$\begin{aligned} \frac{|U|}{4} &\leq \sum_{k \geq k_{\text{median}}} \sum_{j \in U} y_{j^*,j,k} \stackrel{\text{Lem 52}}{\leq} \sum_{k \geq k_{\text{median}}} s_k c \sum_{i \in [m_k]} \sum_t x_{j^*,i,t} \\ &\leq s_{k_{\text{median}}} c \underbrace{\sum_{k \geq k_{\text{median}}} \sum_{i \in [m_k]} \sum_t x_{j^*,i,t}}_{\leq 1} \leq s_{k_{\text{median}}} c \end{aligned}$$

Rearranging gives $|U| \leq 4s_{k_{\text{median}}} c$ as claimed. \square

Claim II. *There is a class $k^* \in \{1, \dots, k_{\text{median}}\}$ where $\rho_{k^*} \geq \frac{|U|}{4M}$.*

Proof of Claim II. For any job $j \in U$ we have $\sum_{k \leq k_{\text{median}}} z_{j,k} \geq (\sum_{k \leq k_{\text{median}}} z_{j^*,k}) - d(j, j^*) \geq \frac{1}{2} - \frac{1}{4} = \frac{1}{4}$. Hence $\sum_{k \leq k_{\text{median}}} \rho_k \geq \frac{|U|}{4}$. Then at least one index $k^* \in \{1, \dots, k_{\text{median}}\}$ will have $\rho_{k^*} \geq \frac{|U|}{4M}$. \square

While the above lemma is enough to obtain our approximation algorithm for makespan on related machines, for the weighted completion time objective we need the following strengthening:

if one considers a set of jobs which are *very* close to each other, then there exists a type k^* such that we can schedule all the jobs in an interval of length $O(c)$ on a single machine of type k^* . Moreover, the LP solution also schedules at least $\Omega(1/M)$ fraction of *every one of these jobs* on the same machine type. Recall that Lemma 53 only satisfied this condition on average.

Lemma 54. *Let $U \subseteq J$ be a non-empty set with $\text{diam}(U) \leq \frac{1}{8M}$ and abbreviate $z_{j,k} := \sum_{i \in [m_k]} \sum_t x_{j,i,t}$. Then*

$$(i) \quad |z_{j_1,k} - z_{j_2,k}| \leq \text{diam}(U) \leq \frac{1}{8M} \text{ for all } j_1, j_2 \in U \text{ and } k \in [M].$$

Moreover there are indices $\pi(U) \subseteq [M]$ such that

$$(ii) \quad z_{j,k} \geq \frac{1}{4M} \text{ for all } j \in U \text{ and } k \in \pi(U)$$

$$(iii) \quad \sum_{k \in \pi(U)} \min\{z_{j,k} : j \in U\} \geq \frac{1}{2}$$

$$(iv) \quad |U| \leq 2 \cdot s_k \cdot c \text{ for all } k \in \pi(U).$$

Proof. We prove the points in order.

- (i) Actually we claim the stronger property of $|z_{j_1,k} - z_{j_2,k}| \leq d(j_1, j_2)$. Sample a distribution $(\tilde{x}, \tilde{y}) \sim \mathcal{D}(j_1, j_2)$ and let $\sigma(j_1) \in [m]$ be the random variable that denotes the machine index with $\sum_t \tilde{x}_{j_1, \sigma(j_1), t} = 1$ (similarly we define $\sigma(j_2)$). Then we can see that

$$|z_{j_1,k} - z_{j_2,k}| = |\Pr[\sigma(j_1) \in [m_k]] - \Pr[\sigma(j_2) \in [m_k]]| \leq |\Pr[\sigma(j_1) \neq \sigma(j_2)]| \leq d(j_1, j_2).$$

Now we fix any job $j^* \in U$ and set $\pi(U) := \{k \in [M] : z_{j^*,k} \geq \frac{1}{4M}\}$. Then we continue the proof:

$$(ii) \quad \text{By (i) we know that for each } j \in U \text{ and } k \in \pi(U) \text{ we have } z_{j,k} \geq z_{j^*,k} - \frac{1}{8M} \geq \frac{1}{4M}.$$

(iii) We have

$$\begin{aligned}
\sum_{k \in \pi(U)} \min\{z_{j,k} : j \in U\} &\geq \sum_{k \in \pi(U)} \left(z_{j^*,k} - \frac{1}{8M} \right) \\
&\geq \underbrace{\sum_{k \in [M]} z_{j^*,k}}_{=1} - \sum_{k \in [M] \setminus \pi(U)} \underbrace{z_{j^*,k}}_{\leq 1/(4M)} - \frac{1}{8M} \underbrace{|\pi(U)|}_{\leq M} \geq \frac{1}{2}
\end{aligned}$$

(iv) Fix a machine type $k \in \pi(U)$ and consider the event $\mathcal{E} := “\sum_{i \in [m_k]} \sum_t \tilde{x}_{j^*,i,t} = 1”$ (meaning the event that j^* is assigned to a machine in class k). Note that $\Pr_{(\bar{x}, \bar{y}) \sim \mathcal{D}(j^*)}[\mathcal{E}] = z_{j^*,k} \geq \frac{1}{4M}$. Now, let (\bar{x}, \bar{y}) be the Sherali-Adams lift (see Lemma 49) conditioned on the event \mathcal{E} . Then trivially $\sum_{i \in [m_k]} \sum_t \bar{x}_{j^*,i,t} = 1$. As we have conditioned on an event with probability $z_{j^*,k} \geq \frac{1}{4M}$, the chance of other events cannot increase by more than a factor of $4M$ (see again Lemma 49). In particular for $j \in U$ we have $1 - \bar{y}_{j^*,j} \leq 4M \cdot (1 - y_{j^*,j}) = 4M \cdot d(j^*, j) \leq \frac{1}{2}$. As $\bar{y}_{j^*,j,k'} = 0$ for all $k' \neq k$ and $j \in U$ we can conclude that $\bar{y}_{j^*,j,k} \geq \frac{1}{2}$ for all $j \in U$. Finally by Lemma 52 we know that $\sum_{j \in J} \bar{y}_{j^*,j,k} \leq s_k c$ which then gives $|U| \leq 2s_k c$.

□

As it is always the case when using the Sherali-Adams hierarchy in the context of approximation algorithms, it is a valid question how much of the power of the hierarchy is really needed. Some properties such as the triangle inequality from Lemma 51 or the constraints from Lemma 52 could be enforced by simply adding them to the original LP. However at various places such as Lemmas 54, 57 and 58 our analysis makes heavily use of the local consistency provided by the hierarchy. It remains open whether there is a more efficient linear program, say only with the original variables $x_{i,j,t}, y_{j_1,j_2}, C_j$ that suffices.

4.4 The Rounding Algorithm for Q | prec, $p_j = 1, c \mid \sum w_j C_j$

In this section, we describe the rounding algorithm which proves our main technical result, Theorem 55. We let C_j^A denote the algorithm's completion time of job j . We denote $\Gamma^-(j)$ as the predecessors of j and $\Gamma^+(j)$ as the successors, and similarly $\Gamma^{-/+}(J') = \{j \in J : \exists j' \in J' \text{ s.t. } j \in \Gamma^{-/+}(j')\}$.

Theorem 55. *There is a polynomial-time randomized algorithm that, given a solution (x, y, C) to (LP) for $r \geq 5$, produces a schedule with completion times $\{C_j^A\}_{j \in J}$ such that*

$$\mathbb{E}[C_j^A] \leq O(M \cdot \log^2 n) \cdot C_j \leq O(\log^3 n) \cdot C_j.$$

Before we describe our rounding algorithm, we set up some notation. Let (x, y, C) be an optimal solution to the LP with $r \geq 5$. We partition the jobs based on their fractional completion times in the LP solution. For any constant $a > 128$ set $\delta = \frac{1}{a \cdot M \cdot \log 2n}$. Consider

$$J_0 := \{j : C_j < \delta \cdot c\} \tag{4.1}$$

and for $\gamma \geq 1$ define

$$J_\gamma := \{j : 2^{\gamma-1} \cdot \delta \cdot c \leq C_j < 2^\gamma \cdot \delta \cdot c\} \tag{4.2}$$

Our rounding algorithm has the following steps:

- Schedule J_0 via the first batch scheduling algorithm given in Section 4.4.1
- For $\gamma = 1, 2, \dots$, schedule J_γ via the intermediate scheduling algorithm given in Section 4.4.2.

- Concatenate the schedules of J_0, J_1, J_2, \dots

We show that the expected completion time of every job in the above schedule satisfies Theorem 55. We need the following scheduling subroutine for the rounding steps above.

Theorem 56. *Let (x, y, C) be a solution to the LP with $r \geq 5$ and let $T^* \in \mathbb{N}$. For some subset $J' \subset J$, suppose in the LP solution completion time $C_j \leq T^*$ for every job $j \in J'$. Then there is a randomized polynomial time algorithm that schedules all jobs such that $C_j^A \leq O(\log m \cdot \log n \cdot T^*) + O(\log m \cdot c)$ for every job $j \in J'$.*

Note that the above theorem immediately implies our result for the makespan minimization problem on related machines. We prove this theorem in Section 4.5.

4.4.1 First Batch Scheduling

Here we give an algorithm for scheduling jobs in J_0 . We define the α -point of job j as the earliest time t_j^* when the LP solution has completed an α -fraction of j . Formally,

$$t_j^* := \min \left\{ t' \in [T] : \sum_{i=1}^m \sum_{t \leq t'} x_{j,i,t} \geq \alpha \right\}. \quad (4.3)$$

For $\beta \leq 1/4M$, consider any subset $U \subseteq J$ with $\text{diam}(U) \leq \beta$. Let $\pi(U) \subseteq [M]$ denote the indices of machine types that satisfy the conditions of Lemma 54. We use the same semimetric $d(j_1, j_2) := 1 - y_{j_1, j_2}$ and schedule jobs in J_0 using the following algorithm.

FIRST BATCH SCHEDULING

- (1) Run a CKR clustering on the semimetric space (J_0, d) with parameter $\Delta := \frac{1}{100M}$ and let V_1, \dots, V_q be the clusters.
- (2) Let $V'_\ell := \{j \in V_\ell \mid (\Gamma^-(j) \cap J_0) \subseteq V_\ell\}$ for $\ell = 1, \dots, q$, and $J'_0 := \cup_{\ell=1}^q V'_\ell$.
- (3) FOR $\ell = 1$ TO q DO
 - (4) Sample a machine type k^* from the set $\pi(V'_\ell)$ with probability $\frac{\min\{z_{jk^*}: j \in V'_\ell\}}{\sum_{k \in \pi(V'_\ell)} \min\{z_{jk}: j \in V'_\ell\}}$.
 - (5) Assign V'_ℓ to a machine $i \in [m_{k^*}]$ with probability $\frac{1}{m_{k^*}}$.
- (6) For a machine $i \in [m_k]$ of type k , let $J'_0(i, k)$ denote the set of jobs assigned to machine i .
- (7) For all k and for all $i \in [m_k]$, schedule $J'_0(i, k)$ in the increasing order of their α -points for $\alpha = (1 - \frac{1}{100M})$ as defined in Eq. (4.3).
- (8) Insert a gap of c time slots.
- (9) Let $J''_0 := J_0 \setminus J'_0$ be the set of jobs that did not get scheduled in steps (1) - (5). Use Theorem 56 to schedule J''_0 .

We now argue that expected completion time of a job j in above algorithm is comparable to its LP cost. First we focus on bounding the completion time of jobs in J'_0 . We need the following crucial lemma.

Lemma 57. *Let $U \subseteq J$ be a set of jobs with $\text{diam}(U) \leq \frac{1}{100M}$ w.r.t. distance d . For every job $j \in U$, define t_j^* as in Eq (4.3). For $\theta > 0$, consider the set of jobs $U^* := \{j \in U : t_j^* < \theta\}$ with α -point less than θ . Then $|U^*| \leq 2 \cdot s_{k^*} \cdot \theta$ for any $k^* \in \pi(U)$.*

The lemma formalizes the intuition that as the jobs in U^* are very close to each other, it must be the case that the LP schedules all of them on the same machine. As a good fraction of each of

these jobs are scheduled in the interval of length θ there cannot be too many of them.

Proof. We prove the lemma by contradiction. Consider a set U^* that violates the conditions in the lemma and fix a job $j^* \in U^*$. Let $k^* \in \pi(U)$. Then we have

$$(A) \sum_{j \in U^*} \sum_{i \in [m]} \sum_{t > \theta} x_{j,i,t} < \frac{1}{100M} \cdot |U^*|, \quad (B) \sum_{j \in U^*} y_{j,j^*} \geq \left(1 - \frac{1}{100M}\right) |U^*|,$$

$$(C) \sum_{i \in [m_{k^*}]} \sum_{t \leq \theta} x_{j^*,i,t} \geq \frac{1}{5M}.$$

where (A) follows from the definition of α -point of jobs, (C) follows from $\sum_{i \in [m_{k^*}]} \sum_{t=1}^{\theta} x_{j^*,i,t} \geq 1/4M - 1/100M > 1/5M$, and (B) follows from $\text{diam}(U^*) \leq \text{diam}(U) \leq \frac{1}{100M}$. Now we appeal to the properties of the Sherali-Adams hierarchy. Consider the event $\mathcal{E} := \sum_{i \in [m_{k^*}]} \sum_{t \leq \theta} x_{j^*,i,t} = 1$ and note that from (C) we know that the probability of the event \mathcal{E} is at least $\frac{1}{5M}$. Now let (\bar{x}, \bar{y}) be the Sherali-Adams lift with $\bar{x} \in SA_{r-1}(K)$ conditioned on this event, see Lemma 49 for details. In particular one has $\sum_{i \in [m_{k^*}]} \sum_{t \leq \theta} \bar{x}_{j^*,i,t} = 1$, meaning that the LP solution (\bar{x}, \bar{y}) schedules j^* fully on machines of class k^* and until time θ .

As we have conditioned on an event of probability at least $\frac{1}{5M}$, the probabilities of other events cannot increase by more than a factor of $5M$ (see the “moreover” part in Lemma 49). In particular for $j \in U$ we have $1 - \bar{y}_{j^*,j} \leq 5M \cdot (1 - y_{j^*,j}) = 5M \cdot d(j^*, j) \leq \frac{1}{10}$. Similarly, $\sum_{j \in U^*} \sum_{i \in [m]} \sum_{t > \theta} \bar{x}_{j,i,t} \leq 5M \cdot (\sum_{j \in U^*} \sum_{i \in [m]} \sum_{t > \theta} x_{j,i,t}) \leq 5M \cdot \frac{1}{100M} \cdot |U^*| \leq \frac{|U^*|}{10}$. Therefore we have

$$(A') \sum_{j \in U^*} \sum_{i \in [m_{k^*}]} \sum_{t > \theta} \bar{x}_{j,i,t} \leq \frac{|U^*|}{10}, \quad (B') \sum_{j \in U^*} \bar{y}_{j,j^*} \geq \frac{9}{10} \cdot |U^*|, \quad (C') \sum_{i \in [m_{k^*}]} \sum_{t \leq \theta} \bar{x}_{j^*,i,t} = 1.$$

Then by Markov's inequality there exists an outcome $(\tilde{x}, \tilde{y}) \sim \mathcal{D}_{\bar{y}}(j^*)$ where the job j^* is integrally

assigned and

$$(A'') \sum_{j \in U^*} \sum_{i \in [m_{k^*}]} \sum_{t > \theta} \tilde{x}_{j,i,t} \leq \frac{|U^*|}{5}, \quad (B'') \sum_{j \in U^*} \tilde{y}_{j,j^*} \geq \frac{4}{5} \cdot |U^*|, \quad (C'') \sum_{i \in [m_{k^*}]} \sum_{t \leq \theta} \tilde{x}_{j^*,i,t} = 1.$$

We fix that outcome and let $i^* \in [m_{k^*}]$, $t^* \in \{1, \dots, \theta\}$ be the indices with $\tilde{x}_{j^*,i^*,t^*} = 1$. Let ℓ^* be the interval index such that $t^* \in I_{\ell^*,k^*}$. Then

$$\begin{aligned} \frac{4}{5} \cdot |U^*| &\stackrel{(B'')}{\leq} \sum_{j \in U^*} \tilde{y}_{j,j^*} \stackrel{LP}{=} \sum_{j \in U^*} \sum_{\ell \in \{0, \dots, L-1\}} \sum_k \sum_{i \in [m_k]} \sum_{t_1 \in I_{\ell,k}} \sum_{t_2 \in I_{\ell,k}} \tilde{x}_{(j,i,t_1),(j^*,i,t_2)} \\ &\leq \sum_{j \in U^*} \sum_{t \in I_{\ell^*,k^*}} \tilde{x}_{j,i^*,t} = \underbrace{\sum_{j \in U^*} \sum_{t \in I_{\ell^*,k^*}: t \leq \theta} \tilde{x}_{j,i^*,t}}_{\leq s_{k^*} \cdot \theta \text{ by LP}} + \underbrace{\sum_{j \in U^*} \sum_{t \in I_{\ell^*,k^*}: t > \theta} \tilde{x}_{j,i^*,t}}_{\leq \frac{1}{5} \cdot |U^*| \text{ by } (A'')} \\ &\leq s_{k^*} \cdot \theta + \frac{|U^*|}{5} \end{aligned}$$

Rearranging gives $|U^*| \leq \frac{5}{3} \cdot s_{k^*} \cdot \theta^*$, which is a contradiction. \square

Lemma 58. *Let $\theta > 0$ and i be any machine of type k . Then with high probability $1 - 1/n^{100}$ it holds that*

$$|\{j \in J'_0(i, k) : t_j^* \leq \theta\}| \leq O\left(\frac{\log n}{\log \log n} \cdot s_k \cdot \theta\right)$$

Proof. Fix $\theta^* > 0$ and any machine i^* of type k . Consider any outcome of the clustering itself.

The randomness needed for the lemma is only over the assignment in step (3)-(5) of the algorithm.

We define the following random variables. Let

$$X_\ell = \begin{cases} |\{j \in V'_\ell : t_j^* \leq \theta^*\}| & \text{if } V'_\ell \text{ is assigned to machine } i^* \text{ of type } k \text{ by our algorithm,} \\ 0 & \text{otherwise} \end{cases}$$

and abbreviate $X = \sum_{\ell=1}^q X_\ell$. Note the random variables X_1, \dots, X_q are independent and $X = |\{j \in J'_0(i^*, k) : t_j^* \leq \theta^*\}|$. From Lemma 57 we know that $X_\ell \leq 2s_k \cdot \theta^*$ for all ℓ .

Next, we estimate $\mathbb{E}[X]$. Recall that the algorithm uses the indices $\pi(V'_\ell) \subseteq [M]$ from Lemma 54 which in particular satisfy $z_{j,k} \geq \frac{1}{4M}$ and $|z_{j,k} - z_{j',k}| \leq \frac{1}{8M}$ whenever $j, j' \in V'_\ell$ with $k \in \pi(V'_\ell)$.

Abbreviate

$$\mu_{\ell,k} := \begin{cases} \min\{z_{j,k} : j \in V'_\ell\} & \text{if } k \in \pi(V'_\ell) \\ 0 & \text{otherwise.} \end{cases}$$

Recall that for every ℓ we have $\sum_{k \in [M]} \mu_{\ell,k} \geq 1/2$. We prove two technical claims that we need for our analysis:

Claim I. For any ℓ and k and $j \in V'_\ell$ one has $\mu_{\ell,k} \leq 2z_{j,k}$.

Proof of Claim I. If $k \notin \pi(V'_\ell)$ then the left hand side is 0, so suppose $k \in \pi(V'_\ell)$. Then $\mu_{\ell,k} = \min\{z_{j',k} : j' \in V'_\ell\} \leq z_{j,k} + \frac{1}{8M} \leq 2z_{j,k}$. \square

Claim II. For $j \in V'_\ell$ with $t_j^* \leq \theta^*$ and $k \in \pi(V'_\ell)$ one has $z_{j,k} \leq 2 \sum_{i \in [m_k]} \sum_{t \in I_{*,k} : t \leq \theta^*} x_{j,i,t}$.

Proof of Claim II. Since $k \in \pi(V'_\ell)$ we know that $z_{j,k} \geq \frac{1}{4M}$. We consider the distribution $(\tilde{x}, \tilde{y}) \sim \mathcal{D}(j)$ and denote \tilde{i} and \tilde{t} as the random indices so that $\tilde{x}_{j,\tilde{i},\tilde{t}} = 1$. Then $\Pr[\tilde{i} \in [m_k]] = z_{j,k}$ and $\Pr[\tilde{t} > \theta^*] = \sum_{k \in [M]} \sum_{i \in [m_k]} \sum_{t \in I_{*,k} : t > \theta^*} x_{j,i,t} \leq \frac{1}{100M}$ by the definition of α -points and the assumption $t_j^* \leq \theta^*$. Then

$$\begin{aligned} \sum_{i \in [m_k]} \sum_{t \in I_{*,k} : t \leq \theta^*} x_{j,i,t} &= \Pr[\tilde{i} \in [m_k] \text{ and } \tilde{t} \leq \theta^*] \\ &\geq \Pr[\tilde{i} \in [m_k]] - \Pr[\tilde{t} > \theta^*] \\ &\geq z_{j,k} - \frac{1}{100M} \geq \frac{1}{2} z_{j,k} \quad \square \end{aligned}$$

Now we can bound the expectation of our random variable as

$$\begin{aligned}
\mathbb{E}[X] &= \frac{1}{m_k} \sum_{\ell=1}^q \Pr[V'_\ell \text{ assigned to class } k] \cdot |\{j \in V'_\ell : t_j^* \leq \theta^*\}| \\
&= \frac{1}{m_k} \sum_{\ell=1}^q \frac{\mu_{\ell,k}}{\sum_{k'} \mu_{\ell,k'}} \cdot |\{j \in V'_\ell : t_j^* \leq \theta^*\}| \\
&\stackrel{\text{Claim I}}{\leq} \frac{4}{m_k} \sum_{\ell: k \in \pi(V'_\ell)} \sum_{j \in V'_\ell : t_j^* \leq \theta^*} z_{j,k} \\
&\stackrel{\text{Claim II}}{\leq} \frac{8}{m_k} \sum_{\ell: k \in \pi(V'_\ell)} \sum_{j \in V'_\ell : t_j^* \leq \theta^*} \sum_{i \in [m_k]} \sum_{t \in I_{*,k} : t \leq \theta^*} x_{j,i,t} \\
&\leq \frac{8}{m_k} \sum_{i \in [m_k]} \sum_{t \in I_{*,k} : t \leq \theta^*} \underbrace{\sum_{j \in J'_0} x_{j,i,t}}_{\leq 1 \text{ by (LP)}} \leq 8 \cdot |\{t \in I_{*,k} : t \leq \theta^*\}| \leq 8s_k \cdot \theta^*
\end{aligned}$$

Here we also have used that $\sum_{k' \in [M]} \mu_{\ell,k'} \geq \frac{1}{2}$. As X is sum of independent random variables X_ℓ each of which is bounded by $O(s_k \cdot \theta^*)$, we apply the Chernoff bound from Lemma 50 and obtain that $\Pr[X > C' \cdot \frac{\log n}{\log \log n} \cdot s_k \theta^*] \leq 1/n^{1000}$ for some constant $C' > 0$. To complete the lemma, we simply do a union bound over all possible values of θ^* and i . The number of machines is $m \leq n$ and the number of relevant θ^* 's in the time horizon is bounded by $\sum_{k \in [M]} |I_{*,k}| \leq M \cdot 2n$. \square

Lemma 59. *For every job $j \in J'_0$, with high probability, the completion time of j in our schedule*

$$C_j^A \leq O\left(M \cdot \frac{\log n}{\log \log n} \cdot C_j\right).$$

Proof. Fix a job $j^* \in J'_0$, and suppose job j^* is scheduled on machine i of type k in our schedule. From Lemma 58, there are at most $O(\frac{\log n}{\log \log n} \cdot s_k \cdot t_{j^*}^*)$ jobs whose $t_j^* \leq t_{j^*}^*$. Therefore, the completion time is $C_{j^*}^A \leq O(\frac{\log n}{\log \log n} \cdot t_{j^*}^*)$. Note that in the LP solution at least an $\frac{1}{100M}$ -fraction of j^* was scheduled at or after $t_{j^*}^*$. Hence, $C_j \geq \frac{1}{100M} \cdot t_{j^*}^*$. Putting things together we conclude $C_j^A \leq O(\frac{\log n}{\log \log n} \cdot t_{j^*}^*) \leq O(M \cdot \frac{\log n}{\log \log n} \cdot C_j)$. \square

Now we focus on bounding the completion time of jobs in J_0'' .

Lemma 60. *For every job $j \in J_0''$, the completion time of j in our schedule $C_j^A \leq O(\log n \cdot c)$.*

Proof. From the definition of set J_0 , for all $j \in J_0''$ the completion time $C_j \leq c \cdot \delta$ in the LP solution. This implies that at least a 1/2-fraction of every job $j \in J_0''$ is scheduled in the interval $[0, 2c \cdot \delta]$ in the LP solution. We take the LP solution restricted to J_0'' in the interval $[0, c \cdot \delta]$, and invoke Theorem 56 to construct a schedule of jobs J_0'' . Theorem 56 guarantees that every job $j \in J_0''$ is scheduled in an interval of length at most $2c \cdot \delta \cdot (\log m \cdot \log n + M^2) + O(\log m \cdot c) = O(\log m \cdot c)$ from our choice of δ . To complete the proof, it remains to account for the increase in completion time caused by jobs in J_0' as our algorithm schedules J_0'' after scheduling jobs in J_0' . However, from Lemma 59 every job in J_0' has completion time at most $O(\log n \cdot c)$. Thus the lemma follows. \square

Lemma 61. *For any job $j_1 \in J_0$, the probability that $j_1 \in J_0''$ is at most $O(M \cdot \log n \cdot \frac{C_{j_1}}{c})$.*

Proof. Consider the set $U := \{j_1\} \cup (\Gamma^-(j_1) \cap J_0)$ of j_1 and its ancestors. If $j_0 \prec j_1$, then $0 \leq C_{j_0} + c \cdot d(j_0, j_1) \leq C_{j_1}$ by the LP constraints and so $d(j_0, j_1) \leq \frac{C_{j_1}}{c}$. Then the diameter of U with respect to semimetric d is bounded by $2C_{j_1}/c$ and hence by Theorem 23.(b) the probability that U is separated is bounded by $\ln(2n) \cdot \frac{4\text{diam}(U)}{\Delta} \leq O(M \cdot \log n \cdot \frac{C_{j_1}}{c})$. \square

We can now bound the expected completion time of every job in J_0 .

Lemma 62. *For every job $j \in J_0$, $\mathbb{E}[C_j^A] \leq O(M \cdot \log^2 n) \cdot C_j$.*

Proof. Fix a job $j_1 \in J_0$ and consider

$$\begin{aligned} \mathbb{E}[C_{j_1}^A] &= \mathbb{E}[C_{j_1}^A | (j_1 \in J_0')] \cdot \Pr[(j_1 \in J_0')] + \mathbb{E}[C_{j_1}^A | (j_1 \in J_0'')] \cdot \Pr[(j_1 \in J_0'')] \\ &\stackrel{(*)}{\leq} O\left(M \cdot \frac{\log n}{\log \log n} \cdot C_{j_1}\right) + O\left(M \cdot \log n \cdot \frac{C_{j_1}}{c}\right) \cdot O(\log n \cdot c) \end{aligned}$$

$$\leq O(M \cdot \log^2 n) \cdot C_{j_1},$$

where (*) follows from Lemmas 59, 60, 61. □

We also note the following simple consequence of previous lemmas:

Lemma 63. *For every job $j \in J_0$, $C_j \leq O(\log n \cdot c)$. In other words, the makespan of our schedule for J_0 is at most $O(\log n \cdot c)$.*

4.4.2 Intermediate Batch Scheduling

Now we describe our algorithm to schedule jobs in the set J_γ for $\gamma > 0$, which is similar to that of scheduling jobs in J_0'' . Recall that from the definition of set J_γ in Eq. (4.2), $2^{\gamma-1}\delta \cdot c < C_j \leq 2^\gamma\delta \cdot c$ in the LP solution. We take the LP solution restricted to J_γ in the interval $[0, 2^\gamma\delta \cdot c]$, and invoke Theorem 56 to construct a schedule of jobs J_γ . Theorem 56 ensures that every job $j \in J_\gamma$ is scheduled in an interval of length at most $O(2^\gamma\delta \cdot c \cdot \log m \cdot \log n) + O(\log n \cdot c)$. As a consequence, we obtain the following lemma.

Lemma 64. *Fix any $\gamma^* \geq 0$. For every $j \in J_{\gamma^*}$, $C_j^A \leq O(2^{\gamma^*} \cdot \delta \cdot \log m \cdot \log n) + O(\gamma^* \cdot \log n \cdot c)$.*

Proof. For $\gamma = 0, 1, \dots$, let $\mathcal{S}(J_\gamma)$ denote the schedule of jobs in the set J_γ . Our final schedule \mathcal{S} is simply a concatenation of schedules $\mathcal{S}(J_0), \mathcal{S}(J_1), \dots$, while inserting c empty time slots between $\mathcal{S}(J_\gamma)$ and $\mathcal{S}(J_{\gamma+1})$. Let T_γ denote the makespan of $\mathcal{S}(J_\gamma)$. From Lemma 63 and Theorem 56, $T_\gamma \leq O(2^\gamma \cdot \delta \cdot c \cdot \log m \cdot \log n) + O(\log n \cdot c)$. Fix the job $j^* \in J_{\gamma^*}$ with the highest completion time in \mathcal{S} . We can bound the completion time of j^* as

$$\begin{aligned}
C_{j^*}^A &\leq \sum_{\gamma=0}^{\gamma^*} T_\gamma \leq O(\gamma^* \cdot \log n \cdot c) + \sum_{\gamma=1}^{\gamma^*} O(2^\gamma \cdot \delta \cdot c \cdot \log m \cdot \log n) \\
&\leq O(\gamma^* \cdot \log n \cdot c) + O(2^{\gamma^*} \cdot \delta \cdot c \cdot \log m \cdot \log n)
\end{aligned}$$

which completes the proof. \square

We finish the proof of main theorem for the weighted sum of completion times of jobs.

Proof of Theorem 55. For every job $j \in J_0$, from Lemma 62, we have $\mathbb{E}[C_j^A] \leq O(M \cdot \log^2 n) \cdot C_j$. From the previous Lemma 64, for $\gamma > 0$ and $j \in J_\gamma$ we have, $C_j^A \leq O(\gamma \cdot \log n \cdot c) + O(2^\gamma \cdot \delta \cdot c \cdot \log m \cdot \log n)$. On the other hand, for every $j \in J_\gamma$, by definition, $C_j > 2^\gamma \cdot \delta \cdot c$. Thus, $C_j^A \leq O(M \cdot \log^2 n) \cdot C_j$. \square

4.5 Proof of Theorem 56

Recall that we are given a subset of jobs $J' \subseteq J$, and a solution (x, y, C) to the LP with $r \geq 5$ and $T \in \mathbb{N}$. It is promised that in (x, y, C) , the completion times satisfy $C_j \leq T$ for every job $j \in J'$. Then there is a randomized polynomial time algorithm that schedules all jobs such that $C_j^A \leq O(\log m \cdot \log n \cdot T) + O(\log m \cdot c)$ for every job $j \in J'$.

4.5.1 Main Scheduling Subroutine

The following lemma is the main scheduling subroutine towards proving Theorem 56.

Lemma 65. *Let $C^* \geq 0$ and $\delta = 1/64 \log(2n)$. Consider the set $J^* \subseteq \{j \in J' \mid C^* \leq C_j \leq C^* + \delta \cdot c\}$. Then there is a randomized rounding procedure that finds a subset $J^{**} \subseteq J^*$, a partition*

of J^{**} into $J^{**} = \cup_{k=1}^M J^{**(k)}$ where $J^{**(k)}$ are jobs assigned to machines of type k , and a schedule for J^{**} in an interval of length at most $5c + \sum_k \frac{|J^{**(k)}|}{s_k \cdot m_k}$ such that every job $j \in J^*$ is scheduled with probability at least $\frac{1}{2}$.

The rounding algorithm to prove the above lemma is the following:

SCHEDULING A SINGLE BATCH ON RELATED MACHINES

- (1) Run a CKR clustering on the semimetric space (J^*, d) with parameter $\Delta := \frac{1}{4}$ and let V_1, \dots, V_q be the clusters.
- (2) Let $V'_\ell := \{j \in V_\ell \mid \Gamma^-(j) \cap J^* \subseteq V_\ell\}$ for $\ell = 1, \dots, q$.
- (3) For all $\ell = 1, \dots, q$, assign the jobs in V'_ℓ to type k^* satisfying the conditions in Lemma 53.
- (4) For all $\ell = 1, \dots, q$, if V'_ℓ is assigned to type k^* , assign all jobs in V'_ℓ to the machine of type k^* which has the least load (breaking ties arbitrarily).

We prove few lemmas that will be helpful in proving the desired properties of the algorithm. The first lemma shows that the clusters produced by the algorithm are not too large.

Lemma 66. *For all $\ell = 1, \dots, q$, one has $|V'_\ell| \leq O(s_{k^*}c)$. Here, k^* is the machine type that V'_ℓ is assigned in our algorithm. Thus, it takes $O(c)$ time to process all jobs in V'_ℓ .*

Proof. We know by Theorem 23 that $\text{diam}(V'_\ell) \leq \text{diam}(V_\ell) \leq \Delta \leq \frac{1}{4}$. The lemma follows by Lemma 53 and by our choice of k^* . □

Further it is easy to see that the clusters respect precedence constraints.

Lemma 67. *The solution V'_1, \dots, V'_q is feasible in the sense that jobs on different machines do not have precedence constraints.*

Proof. Consider jobs processed on different machines, say (after reindexing) $j_1 \in V'_1$ and $j_2 \in V'_2$. If $j_1 \prec j_2$ then we did *not* have $\Gamma^-(j_2) \subseteq V'_2$. This contradicts the definition of the sets V'_ℓ . \square

We will use the two statements above together with Theorem 23 to prove Lemma 65.

Proof of Lemma 65. Call the schedule produced by the above algorithm $\mathcal{S}(J^*)$. By Lemma 66, for all $\ell = 1, \dots, q$, the processing time of V'_ℓ is at most $O(c)$. By Lemma 67, there are no dependent jobs in different sets of V'_1, \dots, V'_q .

The interval in which the jobs in J^* are scheduled can be divided into c -length time intervals I_1, \dots, I_ℓ . To prove the lemma, it suffices to consider the case where $\ell \geq 6$. We say that a machine of type k is *full* in a c -length interval if it is processing $c \cdot s_k$ jobs in $\mathcal{S}(J^*)$. From our greedy packing and the bound on processing time of V'_ℓ , it follows that there exists a machine type $k \in [M]$ such that every machine $i \in [m_k]$ is full in $I_1, \dots, I_{\ell-5}$. Therefore,

$$c \cdot (\ell - 5) \leq \max_k \frac{|J^{**}(k)|}{s_k \cdot m_k} \leq \sum_k \frac{|J^{**}(k)|}{s_k \cdot m_k}.$$

Thus, the total length of the interval used in a single batch scheduling is bounded by

$$c \cdot \ell \leq 5c + \sum_k \frac{|J^{**}(k)|}{s_k \cdot m_k}.$$

It remains to prove that a fixed job $j^* \in J^*$ is scheduled with good probability. Consider the set $U := \{j^*\} \cup (\Gamma^-(j^*) \cap J^*)$ of j^* and its ancestors in J^* . By the LP constraint $C_{j_1} + d(j_1, j_2) \cdot c \leq C_{j_2}$, rearranging we see that $d(j_1, j_2) \leq \delta$ for every $j_1, j_2 \in U$. So the diameter of U is at most 2δ . Now we appeal to Lemma 53. For our choice of $\Delta = 1/4$ and $\delta \leq \frac{1}{64 \log(2n)}$, we apply Theorem 23 to see that the cluster is separated with probability at most $\log(2n) \cdot \frac{8\delta}{\Delta} \leq \frac{1}{2}$. \square

To schedule all jobs in J^* , we repeat the clustering procedure $O(\log m)$ times and simply

schedule the remaining jobs on the fastest machine.

Lemma 68. *Let $C^* \geq 0$ and $\delta = 1/64 \log(2n)$. Consider the set $J^* \subseteq \{j \in J' \mid C^* \leq C_j \leq C^* + \delta \cdot c\}$. Then there is an algorithm with expected polynomial running time that finds disjoint subsets $J^{*(k)} \subset J^*$, where $J^{*(k)}$ are jobs assigned to machines of type k , and schedules all jobs in J^* using at most $O(\log m \cdot c) + \frac{|J^*|}{m \cdot s_{\max}} + \sum_k \frac{|J^{*(k)}|}{s_k \cdot m_k}$ many time slots.*

Proof. We run the above algorithm for $2 \log m$ iterations. Input to iteration $h + 1$ is the set of jobs that are not scheduled in the first h iterations. For $h \in \{1, 2, \dots, 2 \log m\}$, let J_h^{**} be the set of jobs scheduled in iteration h , and let $J_{h+1}^* := J^* \setminus \{\bigcup_{i=1}^h J_i^{**}\}$. In this notation, $J_1^* := J^*$.

Let $\mathcal{S}(J_h^{**})$ be the schedule of jobs J_h^{**} obtained from by Lemma 65. We schedule $\mathcal{S}(J_1^{**})$ first, then append schedule $\mathcal{S}(J_h^{**})$ after $\mathcal{S}(J_{h-1}^{**})$ while inserting c empty time slots between them, for $h = 2, \dots, 2 \log m$. Let $\hat{J} := J_{2 \log m+1}^*$ be the set of jobs that were not scheduled in these $2 \log m$ iterations. We schedule all jobs in \hat{J} consecutively on a single machine with the *fastest speed* after the completion of $\mathcal{S}(J_{2 \log m}^{**})$.

From our construction, the length of a schedule for J^* , which is a random variable, is at most $O(\log m \cdot c) + \lceil \frac{|\hat{J}|}{s_{\max}} + \sum_k \frac{|J^{*(k)}|}{s_k \cdot m_k} \rceil$. For $h \in \{1, 2, \dots, 2 \log m\}$, Lemma 65 guarantees that each job $j \in J_h^*$ gets scheduled in the h^{th} iteration with probability at least $1/2$. Therefore, the probability that $j \in \hat{J}$, i.e. it does not get scheduled in the first $2 \log m$ iterations, is at most $\frac{1}{2^m}$. This implies that $\mathbb{E}[|\hat{J}|] \leq \frac{|J^*|}{2^m}$. The claimed expected polynomial running time bound now simply follows from appealing to Markov's inequality.

Finally, by our construction precedence and communication delay constraints are satisfied. \square

4.5.2 The Complete Algorithm

To schedule all jobs in J' we do the following.

THE COMPLETE ALGORITHM

- (1) Let (x, y, C) be a solution to (LP) with $r \geq 5$ for J' .
 - (2) For $\delta = \frac{1}{64 \log(2n)}$ and $h \in \{0, 1, 2 \dots \frac{T-1}{c \cdot \delta}\}$, define

$$J_h = \{j \in J : h \cdot \delta \cdot c \leq C_j < (h+1) \cdot \delta \cdot c\}$$
 - (3) FOR $h = 0$ TO $\frac{T-1}{c \cdot \delta}$ DO
 - (4) Schedule the jobs in J_h using the algorithm in Subsection 4.5.1.
 - (5) Insert c new empty idle slots.
-

Proof of Theorem 56. First consider the case when $T > \delta \cdot c$. The total number of time slots used by our algorithm can be upper bounded by

$$\frac{T}{c \cdot \delta} \cdot O(\log m \cdot c) + \sum_{h=0}^{\frac{T-1}{c \cdot \delta}} \frac{|J_h|}{m \cdot s_{\max}} + \sum_{h,k} \frac{|J_h^{(k)}|}{s_k \cdot m_k} = O(\log m \cdot \log n) \cdot T + \frac{|J'|}{m \cdot s_{\max}} + \sum_k \frac{|J^{(k)}|}{s_k \cdot m_k}.$$

Here, $J_h^{(k)}$ is the set of jobs assigned on machines of type k when scheduling J_h , introduced in lemma 68, and $J^{(k)} = \bigcup_{h=0}^{\frac{T-1}{c \cdot \delta}} J_h^{(k)}$. From lemma 53 and the choice of machine type in the algorithm, $\sum_{j \in J^{(k)}} \sum_t \sum_{i \in [m_k]} x_{j,i,t} \geq \Omega(1/M) \cdot |J^{(k)}|$. On the other hand, by the constraints of the LP, we have

$$\sum_{j \in J^{(k)}} \sum_{t=0}^T \sum_{i \in [m_k]} x_{j,i,t} \leq \sum_{j \in J} \sum_t \sum_{i \in [m_k]} x_{j,i,t} \leq m_k \cdot s_k \cdot T.$$

Thus, $\frac{|J^{(k)}|}{s_k \cdot m_k} \leq O(M \cdot T)$, and $\sum_k \frac{|J^{(k)}|}{s_k \cdot m_k} \leq O(M^2 \cdot T)$. It is also implied from the same constraints of the LP that $|J'| \leq T \cdot m \cdot s_{\max}$. As $M^2 \leq O(\log m \cdot \log n)$, the proof follows assuming $T > \delta \cdot c$.

Now if $T \leq \delta \cdot c$, then $J' = J^*$ and our algorithm does only one iteration. Thus from Lemma

68 the total number of time slots used by our algorithm is at most

$$O(\log m \cdot c) + \frac{|J'|}{m \cdot s_{\max}} + \sum_k \frac{|J'^{(k)}|}{s_k \cdot m_k} \leq O(\log m \cdot c) + O(M^2 \cdot T).$$

This completes the proof. □

Finally, we conclude this section by noting that above proof in fact gives an $O(\log^2 n)$ approximation for the makespan objective on related machines for the unit length case.

4.6 Reductions

An extension of the classic *list scheduling* algorithm of Graham [Gra66] is the *speed based list scheduling* algorithm by Chudak and Shmoys [CS99]. In this setting, one has a set of jobs J with precedence constraints and an assignment that determines on what machine type a job will be executed. Then we process the jobs greedily in the sense that any available job is scheduled as early as possible on one of the machines belonging to its assigned speed class. The makespan can then be upper bounded by the maximum chain length plus the sum of the loads (if summed over all speed classes).

As in previous sections, we will consider a set J of jobs with processing time p_j and m machines, partitioned into machine types $(m_k, s_k)_{k \in [M]}$, meaning that there are m_k machines of speed s_k available. We design a slight extension of speed based list scheduling that allows us to take communication delays into account.

SPEED-BASED LIST SCHEDULING (WITH COMMUNICATION DELAYS)

Input: Jobs J with $p_j \geq 0$, machine types $(m_k, s_k)_{k \in [M]}$; assignment $\pi : J \rightarrow [M]$; communication delay c

Output: Feasible non-preemptive schedule

- (1) Set $\sigma(j) := \emptyset$ for all $j \in J$
- (2) FOR $t = 0$ TO ∞ DO FOR $i = 1$ TO m DO
 - (3) IF i is idle at time t THEN select any job $j \in J$ satisfying the following
 - $\sigma(j) = \emptyset$
 - Every $j' \prec j$ has been completed. Moreover if $j' \prec j$ was scheduled on machine $i' \neq i$, then j' must have finished by time $t - c$
 - Machine i is of class $\pi(j)$
 - (4) Set $\sigma(j) := ([t, t + \frac{p_j}{s_{\pi(j)}}], i)$ (if there was such a job)

We provide an analysis which follows closely Chudak and Shmoys [CS99] as well as Graham [Gra66].

Lemma 69. *Suppose we are given jobs J with processing time p_j and a partial order \prec as well as machine types $(m_k, s_k)_{k \in [M]}$ and an assignment $\pi : J \rightarrow [M]$. Denote $D_k := \sum_{j \in J: \pi(j)=k} \frac{p_j}{m_k s_k}$ as the load on type k . Then the SPEED-BASED LIST SCHEDULING algorithm produces a schedule of length*

$$\sum_{k=1}^M D_k + \max \left\{ \left(\sum_{j \in C} \frac{p_j}{s_{\pi(j)}} \right) + (|C| - 1) \cdot c \mid C \subseteq J \text{ is a chain} \right\} \quad (4.4)$$

Proof. Consider the schedule produced by SPEED-BASED LIST SCHEDULING. Let j_1 denote the last job that finishes. For $\ell \geq 1$ suppose we have already constructed the sequence j_1, \dots, j_ℓ . Then we choose $j_{\ell+1}$ as that job with $j_{\ell+1} \prec j_\ell$ which is finished last in the schedule. We terminate the procedure when we reach a chain $j_q \prec \dots \prec j_2 \prec j_1$ where j_q does not have any predecessors. Let

$i_{j_\ell} \in [m]$ be the machine index where j_ℓ is scheduled and let $[t_{j_\ell}, t_{j_\ell} + \frac{p_{j_\ell}}{s_{\pi(j_\ell)}})$ be the time interval in which j_ℓ is scheduled. We can make the observation that by the time $t_{j_{\ell+1}} + c$, all predecessors of j_ℓ have been completed and moreover a time of c has passed. Hence if the period between $t_{j_{\ell+1}} + c$ and t_{j_ℓ} is non-empty, then all machines in class $\pi(j_\ell)$ had to be busy during that period. We can make the more general conclusion that for the constructed chain $C := \{j_q, \dots, j_1\}$ it is true that for every time unit in $[0, t_{j_1} + \frac{p_{j_1}}{s_{\pi(j_1)}})$ one of the following holds:

- (a) a job from C is processed at time t
- (b) a job from C has finished less than c time units ago
- (c) in at least one class k , all machines are busy at t

We note that the duration of time that can fall into one of the categories (a), (b), (c) is indeed upper bounded by (4.4). The claim is then proven. \square

This can be combined with an assignment argument:

Lemma 70. *Let J be a set of jobs with processing times p_j , a partial order \prec and machine types $(m_k, s_k)_{k \in [M]}$. Suppose there is a pre-emptive migratory schedule with makespan T^* . Then there is an assignment $\pi : J \rightarrow [M]$ such that each load $D_k := \sum_{j \in J: \pi(j)=k} \frac{p_j}{m_k s_k}$ satisfies $D_k \leq 4T^*$ and the maximum chain length is $\Delta \leq 2T^*$.*

Proof. After scaling the time horizon we may assume that the job lengths p_j are multiples of 2 and the individual parts in the schedule have unit length. Let T^* be the makespan of the schedule. Let $x_{jk} := \frac{\text{\#job parts of } j \text{ assigned to class } k}{p_j}$ we see that we have a fractional solution to the assignment LP

$$\begin{aligned} \sum_{k=1}^M x_{jk} &= 1 \quad \forall j \in J && \text{(Assignment-LP-I)} \\ \sum_{j \in J} x_{jk} \frac{p_j}{s_k m_k} &\leq T^* \quad \forall k \in [M] \end{aligned}$$

$$\begin{aligned} \sum_{j \in C} \sum_{k=1}^M x_{jk} \frac{p_j}{s_k} &\leq T^* \quad \forall \text{ chain } C \subseteq J \\ 0 \leq x_{jk} &\leq 1 \quad \forall j \in J \quad \forall k \in [M] \end{aligned}$$

We now perform a standard procedure in approximation algorithms that is usually called *filtering*. Consider a job j and delete the 1/2 of its parts that are on the slowest machines and scale the fractional solution by a factor of 2 on the remaining parts. Then we obtain a fractional solution satisfying

$$\begin{aligned} \sum_{k=1}^M y_{jk} &= 1 \quad \forall j \in J \quad (\text{Assignment-LP-II}) \\ \sum_{j \in J} y_{jk} \frac{p_j}{s_k m_k} &\leq 2T^* \quad \forall k \in [M] \\ 0 \leq y_{jk} &\leq 1 \quad \forall j \in J \quad \forall k \in [M] \end{aligned}$$

Moreover due to the filtering, any assignment $\pi : J \rightarrow [M]$ with $y_{j,\pi(j)} > 0$ for all $j \in J$, satisfies the following two properties:

(A) One has $\frac{p_j}{s_{\pi(j)}} \leq 2T^*$.

(B) The maximum chain length w.r.t. π is $\Delta \leq 2T^*$.

If we imagine $q_{jk} := \frac{p_j}{s_k m_k}$ as item sizes, then the seminal rounding argument of Shmoys-Tardos shows that any fractional solution to (Assignment-LP-II) can be rounded to an integral solution where the right hand side is exceeded by at most $\max\{q_{jk} : y_{jk} > 0\} \leq 2T^*$. The obtained integral solution is then the desired assignment π . \square

We use this for the main reduction:

Theorem 71. *Suppose there is an α -approximation to $Q \mid \text{prec}, p_j = 1, c \mid \sum w_j C_j$. Then there is an $O(M\alpha)$ -approximation for $Q \mid \text{prec}, c \mid \sum w_j C_j$.*

Proof. Let J be the set of jobs with processing times $p_j \in \mathbb{N}$ and weights w_j with original precedence constraints \prec . Let OPT be the minimum weighted sum of completion times over all feasible schedules. We run the α -approximation algorithm for $Q \mid \text{prec}, p_j = 1, c \mid \sum w_j C_j$ treating each job as p_j many unit length tasks that form a chain where the last unit length task receives a weight of w_j and the other tasks receive weight 0. The algorithm will provide us with a schedule σ which satisfies communication delays but it spreads tasks of the same job over different machines (and even different machine types). Let C_j be the completion time of the last task of job j . Then $\sum_{j \in J} w_j C_j \leq \alpha OPT$ by construction.

We abbreviate $J^*(0) := \{j \in J \mid C_j \leq c\}$ and $J^*(r) := \{j \in J \mid c2^{r-1} < C_j \leq c2^r\}$ for $r \in \mathbb{Z}_{\geq 1}$. We will create a schedule that schedules first $J^*(0)$, then $J^*(1), J^*(2), \dots$ where we start the schedule for $J^*(r)$ exactly c time units after the last job of $J^*(r-1)$ is finished. Note that we can re-sort the tasks of jobs in $J^*(0)$ so that (i) the tasks of the same job are scheduled consecutively; (ii) for every job j the last task is still completed by time C_j , (iii) the precedence constraints and communication constraints (within $J^*(0)$) are still satisfied. Note that this re-sorting can be done by sorting the jobs in non-increasing order of C_j 's.

So we fix a value of $r \geq 1$ and set $T^* := c2^{r-1}$ and $J^* := J^*(r)$. It remains to show how to schedule the jobs J^* within a time interval of $O(T^*)$.

We partition the time horizon into intervals of length c each. Consider the jobs J^* and let σ^* be the schedule with makespan T^* which satisfies precedence constraints and communication delays but is potentially preemptive and migratory. We construct a modified set \tilde{J} by merging jobs whose tasks are scheduled on the same machine in one such interval $[\beta c, (\beta + 1)c[$ with $\beta \in \mathbb{Z}_{\geq 0}$ (we call these jobs *compact*); all other jobs are inherited without a change (we call such

jobs *spread out*).

We denote the inherited migratory schedule for \tilde{J} by $\tilde{\sigma}$; the completion time of a job is denoted by \tilde{C}_j and the start time is denoted by \tilde{S}_j . Note that $\tilde{C}_j - \tilde{S}_j \geq p_j$ but because of the preemption this inequality might be strict. We create a new partial order $\tilde{\prec}$ such that one has $j_1 \tilde{\prec} j_2$ for $j_1, j_2 \in \tilde{J}$ if one of the following conditions is satisfied:

1. j_1, j_2 are assigned to the same machine in $\tilde{\sigma}$ and $\tilde{C}_{j_1} \leq \tilde{S}_{j_2}$.
2. j_1, j_2 are assigned to different machines in $\tilde{\sigma}$ and j_1 finishes in an interval before j_2 is started.

Note that in particular this precedence order $\tilde{\prec}$ implies \prec . The maximum chain length is trivially bounded by T^* . But it is important to note that the maximum *number* of jobs in any chain of $\tilde{\prec}$ is bounded by $O(T^*/c)$ as every job is either spread out and hence crosses an interval boundary or it is the unique compact job included in an interval (on that machine). Then Lemma 70 provides us with an assignment $\pi : \tilde{J} \rightarrow [M]$ such that the loads are $D_k \leq 4T^*$ and the maximum chain length is $\Delta \leq 2T^*$. We use the SPEED-BASED LIST SCHEDULING algorithm to find a non-preemptive schedule respecting precedence constraints $\tilde{\prec}$ and communication delays while the makespan is bounded by

$$\sum_{k=1}^M D_k + \sum_{j \in C} \frac{p_j}{s_{\pi(j)}} + (|C| - 1) \cdot c \leq 4MT^* + 2T^* + O\left(\frac{T^*}{c}\right) \cdot c \leq O(MT^*)$$

for some chain $C \subseteq \tilde{J}$. □

4.7 Makespan Minimization on Related Machines

We now prove our result on makespan minimization on related machines, Theorem 48. The proof essentially follows from first using Theorem 56 to obtain a schedule of cost $O(\log^2 n)$ for the unit

length case and then appealing to the reduction in Theorem 71, which increases the makespan at most by a factor of $O(\log n)$. These two steps together give an $O(\log^3 n)$ -approximation for $Q \mid \text{prec}, c \mid C_{\max}$.

Note a small technicality here. Both Theorem 56 and Theorem 71 were proved for the weighted completion time objective and not the makespan objective function. However, it is not hard to see that proofs of both the theorems directly extend to the makespan problem.

For technical completeness, we give a formal proof of our result for makespan. Our proof also serves the purpose of explaining how an algorithm for the sum of weighted completion times can be used to obtain an algorithm for makespan. In our proof below, we first reduce the problem of minimizing makespan to a special case of the problem of minimizing the weighted sum of completion times. But we *do not* use the result on weighted completion time as a black box – that would only imply an $O(\log^4 n)$ -approximation. Instead, we tweak our rounding a bit, and use Theorem 56 and Theorem 71 directly to obtain an $O(\log^3 n)$ -approximation for makespan.

Proof. Consider an input instance $\mathcal{I} := (J, \prec)$ of $Q \mid \text{prec}, c \mid C_{\max}$, and let $T(*)$ denote the optimal makespan for any instance $(*)$ of the problem. We consider two cases:

- Case $T(\mathcal{I}) < c$: In this case, we find all connected components, G_1, G_2, \dots, G_q , in the underlying (undirected) graph of the input instance. Since $T(\mathcal{I}) < c$, for $v \in [q]$ all jobs in G_v are scheduled on the same machine by an optimal solution. For every $v \in [q]$, we create a meta-job j_v with processing time equal to the sum of processing times of jobs in G_v . Then we appeal to the $O(1)$ -approximation algorithm for makespan minimization on related machines without precedence or communication delay constraints.
- Case $T(\mathcal{I}) \geq c$. Here we reduce \mathcal{I} to an instance of $\mathcal{I}' := (J', \prec')$ of $Q \mid \text{prec}, c \mid \sum w_j C_j$. We create a new dummy job j_D with $p_{j_D} = 1$ and weight $w_{j_D} = 1$. We define $J' := J \cup \{j_D\}$

and set weights $w_j = 0$ for every job $j \in J' \setminus \{j_D\}$. Moreover, we define the new precedence constraints \prec' by augmenting the precedence constraints of the instance \mathcal{I} with precedence constraints $j \prec j_D$ from every job $j \in J' \setminus \{j_D\}$.

Suppose \mathcal{S} is a feasible schedule for the instance \mathcal{I}' . From our construction, the makespan of \mathcal{S} is (exactly) equal to the total weighted completion times of jobs in \mathcal{I}' . This follows because only the dummy job j_D has non-zero weight and it has to be scheduled at the last.

We now give an algorithm for solving the instance \mathcal{I}' . Similar to our weighted completion time result, we solve the problem via a two step procedure. We treat each job j with processing time p_j as a chain of p_j unit length jobs. Using standard arguments [DKR⁺20], we can assume that this reduction is both of polynomial size and time. After this reduction, we get a new input instance $\mathcal{I}'' := (J'', \prec')$ of $\mathbf{Q} \mid \text{prec}, p_j = 1, c \mid \sum w_j C_j$.

We observe that $T(\mathcal{I}'') \in [T(\mathcal{I}), 2T(\mathcal{I})]$. This follows because i) the weight of all jobs except j_D is zero in our instance; ii) j_D can be scheduled only after completing all the other jobs, which requires at least $T(\mathcal{I})$ time steps; iii) $c \leq T(\mathcal{I})$.

We solve our LP for the weighted completion time problem on the instance \mathcal{I}'' . Let (x, y, C) be the optimal solution for the LP. We note that all jobs in the LP solution are scheduled in the interval $[0, 2T(\mathcal{I})]$. Here we used the facts that $c \leq T(\mathcal{I})$ and that the LP can indeed schedule all jobs in \mathcal{I}'' within an interval of length $c + T(\mathcal{I})$. We invoke Theorem 56 to obtain a schedule \mathcal{S} with cost at most $O(\log m \cdot \log n) \cdot 2T(\mathcal{I})$. Finally we use Theorem 71 to convert this schedule \mathcal{S} to a schedule \mathcal{S}' where every job is scheduled on one machine in consecutive time steps. From the guarantees of Theorem 71, we know the cost of \mathcal{S}' is at most $O(M \cdot \log m \cdot \log n) \cdot T(\mathcal{I})$.

As noted earlier, the makespan of the schedule \mathcal{S}' is (exactly) equal to the total weighted completion time of jobs in \mathcal{I}' . Thus

$$\text{Makespan of } (\mathcal{S}') \leq O(M \cdot \log m \cdot \log n) \cdot T(\mathcal{I})$$

which completes the proof.

□

4.8 Conclusion

Building off our previous framework for scheduling jobs with precedence constraints and communication delays, we provide polylogarithmic approximation algorithms for scheduling jobs with communication delays on related machines given the objectives of minimizing makespan or minimizing the weighted sum of completion times.

Perhaps the most interesting open question is whether our approach can yield a constant-factor approximation for $P \mid \text{prec}, c \mid C_{\max}$. While a compact form of our linear program is sufficient for the identical machines setting, it is not clear whether it is sufficient on related machines; this is because our analysis relied on the local consistency properties of the Sherali-Adams hierarchy. The compact form does have an integrality gap, but we do not know of any integrality gap for the full lifted LP. Another problem is to handle *non-uniform* communication delays in the problem $P \mid \text{prec}, c_{jk} \mid C_{\max}$; here c_{jk} is the communication delay between jobs $j \prec k$.

Bibliography

- [AFS08] Arash Asadpour, Uriel Feige, and Amin Saberi. Santa claus meets hypergraph matchings. In *Approximation, Randomization and Combinatorial Optimization. Algorithms and Techniques, 11th International Workshop, APPROX 2008, and 12th International Workshop, RANDOM 2008, Boston, MA, USA, August 25-27, 2008. Proceedings*, pages 10–20, 2008.
- [AFS12] Arash Asadpour, Uriel Feige, and Amin Saberi. Santa claus meets hypergraph matchings. *ACM Trans. Algorithms*, 8(3):24:1–24:9, July 2012.
- [AKS15] Chidambaram Annamalai, Christos Kalaitzis, and Ola Svensson. Combinatorial algorithm for restricted max-min fair allocation. In *Proceedings of the Twenty-Sixth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2015, San Diego, CA, USA, January 4-6, 2015*, pages 1357–1372, 2015.
- [AL12] Miguel F. Anjos and Jean B" Lasserre. *Handbook on Semidefinite, Conic and Polynomial Optimization*. Springer, 2012.
- [AMMS08] C. Ambühl, M. Mastrolilli, N. Mutsanas, and O. Svensson. Precedence constraint scheduling and connections to dimension theory of partial orders. In *Bulletin of the European Association for Theoretical Computer Science (EATCS. Citeseer*, 2008.
- [Ann16] Chidambaram Annamalai. Finding perfect matchings in bipartite hypergraphs. In *Proceedings of the Twenty-Seventh Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2016, Arlington, VA, USA, January 10-12, 2016*, pages 1814–1823, 2016.
- [AS10] Arash Asadpour and Amin Saberi. An approximation algorithm for max-min fair allocation of indivisible goods. *SIAM J. Comput.*, 39(7):2970–2989, 2010.
- [Ban17] N. Bansal. Scheduling open problems: Old and new. MAPSP 2017. <http://www.mapsp2017.ma.tum.de/MAPSP2017-Bansal.pdf>, 2017.
- [BD05] Ivona Bezáková and Varsha Dani. Allocating indivisible goods. *SIGecom Exchanges*, 5(3):11–18, 2005.

- [BGK96] E. Bampis, A. Giannakos, and J.C. König. On the complexity of scheduling with large communication delays. *European Journal of Operational Research*, 94(2):252 – 260, 1996.
- [BK10] N. Bansal and S. Khot. Inapproximability of hypergraph vertex cover and applications to scheduling problems. In *Automata, Languages and Programming, 37th International Colloquium, ICALP 2010, Bordeaux, France, July 6-10, 2010, Proceedings, Part I*, volume 6198 of *Lecture Notes in Computer Science*, pages 250–261. Springer, 2010.
- [BS06] Nikhil Bansal and Maxim Sviridenko. The santa claus problem. In *Proceedings of the Thirty-eighth Annual ACM Symposium on Theory of Computing*, STOC '06, pages 31–40, New York, NY, USA, 2006. ACM.
- [CC91] J. Y. Colin and P. Chrétienne. C.p.m. scheduling with small communication delays and task duplication. *Operations Research*, 39(4):680–684, 1991.
- [CCK09] Deeparnab Chakrabarty, Julia Chuzhoy, and Sanjeev Khanna. On allocating goods to maximize fairness. In *50th Annual IEEE Symposium on Foundations of Computer Science, FOCS 2009, October 25-27, 2009, Atlanta, Georgia, USA*, pages 107–116, 2009.
- [CKR04] G. Călinescu, H. Karloff, and Y. Rabani. Approximation algorithms for the 0-extension problem. *SIAM J. Comput.*, 34(2):358–372, 2004.
- [CM18] Siu-Wing Cheng and Yuchen Mao. Restricted max-min fair allocation. In *45th International Colloquium on Automata, Languages, and Programming, ICALP 2018, July 9-13, 2018, Prague, Czech Republic*, pages 37:1–37:13, 2018.
- [CM19] Siu-Wing Cheng and Yuchen Mao. Restricted max-min allocation: Approximation and integrality gap. In *46th International Colloquium on Automata, Languages, and Programming, ICALP 2019, July 9-12, 2019, Patras, Greece.*, pages 38:1–38:13, 2019.
- [CS99] Fabián A. Chudak and David B. Shmoys. Approximation algorithms for precedence-constrained scheduling problems on parallel machines that run at different speeds. *J. Algorithms*, 30(2):323–343, 1999.
- [CS12] M. Chowdhury and I. Stoica. Coflow: A networking abstraction for cluster applications. In *Proceedings of the 11th ACM Workshop on Hot Topics in Networks*, HotNets-XI, page 31?36, New York, NY, USA, 2012. Association for Computing Machinery.
- [CZM⁺11] M. Chowdhury, M. Zaharia, J. Ma, M. I. Jordan, and I. Stoica. Managing data transfers in computer clusters with orchestra. In *Proceedings of the ACM SIGCOMM 2011 Conference*, SIGCOMM '11, page 98?109. Association for Computing Machinery, 2011.

- [DKR⁺20] Sami Davies, Janardhan Kulkarni, Thomas Rothvoss, Jakub Tarnawski, and Yihao Zhang. Scheduling with communication delays via LP hierarchies and clustering. In *Proceedings of the 61st Annual IEEE Symposium on Foundations of Computer Science (to appear), FOCS 2020, November 16-19, 2020, Virtual Event, 2020*.
- [DKR⁺21] Sami Davies, Janardhan Kulkarni, Thomas Rothvoss, Jakub Tarnawski, and Yihao Zhang. Scheduling with communication delays via LP hierarchies and clustering II: weighted completion times on related machines. In Dániel Marx, editor, *Proceedings of the 2021 ACM-SIAM Symposium on Discrete Algorithms, SODA 2021, Virtual Conference, January 10 - 13, 2021*, pages 2958–2977. SIAM, 2021.
- [DP09] Devdatt P. Dubhashi and Alessandro Panconesi. *Concentration of Measure for the Analysis of Randomized Algorithms*. Cambridge University Press, 2009.
- [Dro09] M. Drozdowski. *Scheduling with Communication Delays*, pages 209–299. Springer London, London, 2009.
- [DRZ20] Sami Davies, Thomas Rothvoss, and Yihao Zhang. A tale of santa claus, hypergraphs and matroids. In *Proceedings of the 2020 ACM-SIAM Symposium on Discrete Algorithms, SODA 2020, Salt Lake City, UT, USA, January 5-8, 2020*, pages 2748–2757. SIAM, 2020.
- [EKS08] Tomáš Ebenlendr, Marek Krcál, and Jirí Sgall. Graph balancing: a special case of scheduling unrelated parallel machines. In *Proceedings of the Nineteenth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2008, San Francisco, California, USA, January 20-22, 2008*, pages 483–490, 2008.
- [Fei08] Uriel Feige. On allocations that maximize fairness. In *Proceedings of the Nineteenth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA '08*, pages 287–293, Philadelphia, PA, USA, 2008. Society for Industrial and Applied Mathematics.
- [FKK⁺14] Z. Friggstad, J. Könemann, Y. Kun-Ko, A. Louis, M. Shadravan, and M. Tulsiani. Linear programming hierarchies suffice for directed steiner tree. In *Integer Programming and Combinatorial Optimization - 17th International Conference, IPCO 2014, Bonn, Germany, June 23-25, 2014. Proceedings*, pages 285–296, 2014.
- [FRT04] J. Fakcharoenphol, S. Rao, and K. Talwar. A tight bound on approximating arbitrary metrics by tree metrics. *J. Comput. Syst. Sci.*, 69(3):485–497, 2004.
- [GCL18] Yuanxiang Gao, Li Chen, and Baochun Li. Spotlight: Optimizing device placement for training deep neural networks. In Jennifer Dy and Andreas Krause, editors, *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 1676–1684, StockholmsmÅCâCňssan, Stockholm Sweden, 10–15 Jul 2018. PMLR.

- [GFC⁺12] Z. Guo, X. Fan, R. Chen, J. Zhang, H. Zhou, S. McDirmid, C. Liu, W. Lin, J. Zhou, and L. Zhou. Spotting code optimizations in data-parallel pipelines through periscope. In *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation (OSDI 12)*, pages 121–133, 2012.
- [GK07] R. Giroudeau and J.C. König. Scheduling with communication delays. In Eugene Levner, editor, *Multiprocessor Scheduling*, chapter 4. IntechOpen, Rijeka, 2007.
- [GKMP08] R. Giroudeau, J.C. König, F. K. Moulai, and J. Palaysi. Complexity and approximation for precedence constrained scheduling problems with large communication delays. *Theoretical Computer Science*, 401(1):107 – 119, 2008.
- [GLLK79] R. L. Graham, E. L. Lawler, J. K. Lenstra, and A. H. G. Rinnooy Kan. Optimization and approximation in deterministic sequencing and scheduling: a survey. *Ann. Discrete Math.*, 4:287–326, 1979.
- [Gra66] R. L. Graham. Bounds for certain multiprocessing anomalies. *Bell System Technical Journal*, 45(9):1563–1581, 1966.
- [GVY93] N. Garg, V. V. Vazirani, and M. Yannakakis. Approximate max-flow min-(multi)cut theorems and their applications. *SIAM Journal on Computing*, 25:698–707, 1993.
- [Hax95] Penny E. Haxell. A condition for matchability in hypergraphs. *Graphs and Combinatorics*, 11(3):245–248, 1995.
- [HCB⁺19] Y. Huang, Y. Cheng, A. Bapna, O. Firat, D. Chen, M. Chen, H. Lee, J. Ngiam, Q. V. Le, Y. Wu, and Z. Chen. Gpipe: Efficient training of giant neural networks using pipeline parallelism. In *Advances in Neural Information Processing Systems*, 2019.
- [HCG12] C. Hong, M. Caesar, and P.B. Godfrey. Finishing flows quickly with preemptive scheduling. *ACM SIGCOMM Computer Communication Review*, 42(4):127–138, 2012.
- [HLV94] J.A. Hoogeveen, J.K. Lenstra, and B. Veltman. Three, four, five, six, or the complexity of scheduling with communication delays. *Operations Research Letters*, 16(3):129–137, 1994.
- [HM01] C. Hanen and A. Munier. An approximation algorithm for scheduling dependent tasks on m processors with small communication delays. *Discrete Applied Mathematics*, 108(3):239 – 257, 2001.
- [HS87] D. S. Hochbaum and D. B. Shmoys. Using dual approximation algorithms for scheduling problems theoretical and practical results. *J. ACM*, 34(1):144–162, 1987.

- [HSSW97] Leslie A Hall, Andreas S Schulz, David B Shmoys, and Joel Wein. Scheduling to minimize average completion time: Off-line and on-line approximation algorithms. *Mathematics of operations research*, 22(3):513–544, 1997.
- [JKS93] H. Jung, L.M. Kirousis, and P. Spirakis. Lower bounds and efficient algorithms for multiprocessor scheduling of directed acyclic graphs with communication delays. *Information and Computation*, 105(1):94 – 104, 1993.
- [JR17] Klaus Jansen and Lars Rohwedder. On the configuration-lp of the restricted assignment problem. In *Proceedings of the Twenty-Eighth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2017, Barcelona, Spain, Hotel Porta Fira, January 16-19*, pages 2670–2678, 2017.
- [JR18a] Klaus Jansen and Lars Rohwedder. Compact LP relaxations for allocation problems. In *1st Symposium on Simplicity in Algorithms, SOSA 2018, January 7-10, 2018, New Orleans, LA, USA*, pages 11:1–11:19, 2018.
- [JR18b] Klaus Jansen and Lars Rohwedder. Local search breaks 1.75 for graph balancing. *CoRR*, abs/1811.00955, 2018.
- [JR18c] Klaus Jansen and Lars Rohwedder. A note on the integrality gap of the configuration LP for restricted santa claus. *CoRR*, abs/1807.03626, 2018.
- [JZA19] Zhihao Jia, Matei Zaharia, and Alex Aiken. Beyond data and model parallelism for deep neural networks. In *Proceedings of the 2nd SysML Conference, SysML ’19*, Palo Alto, CA, USA, 2019.
- [KLTY20] J. Kulkarni, S. Li, J. Tarnawski, and M. Ye. *Hierarchy-Based Algorithms for Minimizing Makespan under Precedence and Communication Constraints*, pages 2770–2789. 2020.
- [KMN11] A. Karlin, C. Mathieu, and C. Thach Nguyen. Integrality gaps of linear and semi-definite programming relaxations for knapsack. In *Integer Programming and Combinatorial Optimization - 15th International Conference, IPCO 2011, New York, NY, USA, June 15-17, 2011. Proceedings*, pages 301–314, 2011.
- [Lau03] M. Laurent. A comparison of the sherali-adams, lovász-schrijver, and lasserre relaxations for 0-1 programming. *Math. Oper. Res.*, 28(3):470–496, 2003.
- [Li17] Shi Li. Scheduling to minimize total weighted completion time via time-indexed linear programming relaxations. *SIAM Journal on Computing*, (0):FOCS17–409, 2017.
- [LLRKS93] E. Lawler, J.K. Lenstra, A. Rinnooy Kan, and D. Shmoys. Sequencing and scheduling: Algorithms and complexity. *Handbooks in operations research and management science*, 4:445–522, 1993.

- [LR99] T. Leighton and S. Rao. Multicommodity max-flow min-cut theorems and their use in designing approximation algorithms. *J. ACM*, 46(6):787–832, November 1999.
- [LR02] R. Lepere and C. Rapine. An asymptotic $o(\ln \rho / \ln \ln \rho)$ -approximation algorithm for the scheduling problem with duplication on large communication delay graphs. In Helmut Alt and Afonso Ferreira, editors, *STACS 2002*, pages 154–165, Berlin, Heidelberg, 2002. Springer Berlin Heidelberg.
- [LR16] E. Levey and T. Rothvoss. A $(1+\epsilon)$ -approximation for makespan scheduling with precedence constraints using LP hierarchies. In *Proceedings of the 48th Annual ACM SIGACT Symposium on Theory of Computing, STOC 2016, Cambridge, MA, USA, June 18-21, 2016*, pages 168–177, 2016.
- [LRK78] J. K. Lenstra and A. H. G. Rinnooy Kan. Complexity of scheduling under precedence constraints. *Oper. Res.*, 26(1):22–35, February 1978.
- [LST87] Jan Karel Lenstra, David B. Shmoys, and Éva Tardos. Approximation algorithms for scheduling unrelated parallel machines. In *28th Annual Symposium on Foundations of Computer Science, Los Angeles, California, USA, 27-29 October 1987*, pages 217–224, 1987.
- [LYZ⁺16] S. Luo, H. Yu, Y. Zhao, S. Wang, S. Yu, and L. Li. Towards practical and near-optimal coflow scheduling for data center networks. *IEEE Transactions on Parallel and Distributed Systems*, 27(11):3366–3380, 2016.
- [MH97] A. Munier and C. Hanen. Using duplication for scheduling unitary tasks on m processors with unit communication delays. *Theoretical Computer Science*, 178(1):119 – 127, 1997.
- [MK97] A. Munier and J.C. König. A heuristic for a scheduling problem with communication delays. *Operations Research*, 45(1):145–147, 1997.
- [MN06] M. Mendel and A. Naor. Ramsey partitions and proximity data structures. In *47th Annual IEEE Symposium on Foundations of Computer Science (FOCS 2006), 21-24 October 2006, Berkeley, California, USA, Proceedings*, pages 109–118, 2006.
- [MPL⁺17] Azalia Mirhoseini, Hieu Pham, Quoc V Le, Benoit Steiner, Rasmus Larsen, Yuefeng Zhou, Naveen Kumar, Mohammad Norouzi, Samy Bengio, and Jeff Dean. Device placement optimization with reinforcement learning. In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pages 2430–2439. JMLR. org, 2017.
- [MRS⁺20] Biswaroop Maiti, Rajmohan Rajaraman, David Stalfa, Zoya Svitkina, and Aravindan Vijayaraghavan. Scheduling precedence-constrained jobs on related machines with

- communication delay. *To appear at 61st Annual IEEE Symposium on Foundations of Computer Science (FOCS 2020)*, 2020.
- [NHP⁺19] D. Narayanan, A. Harlap, A. Phanishayee, V. Seshadri, N. Devanur, G. Ganger, P. Gibbons, and M. Zaharia. Pipedream: Generalized pipeline parallelism for dnn training. In *Proc. 27th ACM Symposium on Operating Systems Principles (SOSP)*, Huntsville, ON, Canada, October 2019.
- [Pin18] M. Pinedo. *Scheduling: theory, algorithms, and systems*. Springer, 2018.
- [PST04] K. Pruhs, J. Sgall, and E. Torng. Online scheduling. In Joseph Y.-T. Leung, editor, *Handbook of Scheduling - Algorithms, Models, and Performance Analysis*. Chapman and Hall/CRC, 2004.
- [PY90] C. H. Papadimitriou and M. Yannakakis. Towards an architecture-independent analysis of parallel algorithms. *SIAM J. Comput.*, 19(2):322–328, April 1990.
- [QS02] Maurice Queyranne and Maxim Sviridenko. Approximation algorithms for shop scheduling problems with minsum objective. *Journal of Scheduling*, 5(4):287–305, 2002.
- [Rot11] T. Rothvoß. Directed steiner tree and the lasserre hierarchy. *CoRR*, abs/1111.5473, 2011.
- [RS87] V.J. Rayward-Smith. Uet scheduling with unit interprocessor communication delays. *Discrete Applied Mathematics*, 18(1):55 – 71, 1987.
- [SRVW20] Yu Su, Xiaoqi Ren, Shai Vardi, and Adam Wierman. Communication-aware scheduling of precedence-constrained tasks on related machines. *CoRR*, abs/2004.14639, 2020.
- [Sve09] O. Svensson. *Approximability of some classical graph and scheduling problems*. PhD thesis, Università della Svizzera italiana, 2009.
- [Sve10] O. Svensson. Conditional hardness of precedence constrained scheduling on identical machines. In *Proceedings of the Forty-Second ACM Symposium on Theory of Computing, STOC '10*, pages 745–754, New York, NY, USA, 2010. Association for Computing Machinery.
- [Sve11] Ola Svensson. Santa claus schedules jobs on unrelated machines. In *Proceedings of the 43rd ACM Symposium on Theory of Computing, STOC 2011, San Jose, CA, USA, 6-8 June 2011*, pages 617–626, 2011.
- [SW99] P. Schuurman and G. J. Woeginger. Polynomial time approximation algorithms for machine scheduling: Ten open problems, 1999.

- [SZA⁺18] A. Shymyrbay, A. Zhanbolatov, A. Amankhan, A. Bakambekova, and I. Ukaegbu. Meeting deadlines in datacenter networks: An analysis on deadline-aware transport layer protocols. In *2018 International Conference on Computing and Network Communications (CoCoNet)*, pages 152–158. IEEE, 2018.
- [THU92] R. THURIMELLA. A scheduling principle for precedence graphs with communication delay. *International Conference on Parallel Processing*, 3:229–236, 1992.
- [TPD⁺20] Jakub Tarnawski, Amar Phanishayee, Nikhil R. Devanur, Divya Mahajan, and Fanny Nina Paravecino. Efficient algorithms for device placement of dnn graph operators, 2020.
- [VLL90] B. Veltman, B.J. Lageweg, and J.K. Lenstra. Multiprocessor scheduling with communication delays. *Parallel Computing*, 16(2):173 – 182, 1990.
- [VW11] José Verschae and Andreas Wiese. On the configuration-lp for scheduling on unrelated machines. In *Algorithms - ESA 2011 - 19th Annual European Symposium, Saarbrücken, Germany, September 5-9, 2011. Proceedings*, pages 530–542, 2011.
- [ZCB⁺15] Y. Zhao, K. Chen, W. Bai, M. Yu, C. Tian, Y. Geng, Y. Zhang, D. Li, and S. Wang. Rapier: Integrating routing and scheduling for coflow-aware data center networks. In *2015 IEEE Conference on Computer Communications (INFOCOM)*, pages 424–432. IEEE, 2015.
- [ZZC⁺12] J. Zhang, H. Zhou, R. Chen, X. Fan, Z. Guo, H. Lin, J.Y. Li, W. Lin, J. Zhou, and L. Zhou. Optimizing data shuffling in data-parallel computation by understanding user-defined functions. In *Presented as part of the 9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12)*, pages 295–308, 2012.